

ADB Lecture 4

Query Optimization

- Goal is to Select the best available Strategy For executing the query
- Evaluation Plan = Query tree + algorithm for each operator + access methods for each relation
- * Involves Cost-based & heuristic-based approaches
- Heuristic-based Query Optimization
 - apply heuristic rules to modify internal representation of query - to improve its expected Performance.
 - High-level Steps
 - Query Parser generates an initial internal representation.
 - Apply heuristic rules to optimize the internal representation
 - Based on the access paths (e.g. index) available on the files involved in the query, generate a query evaluation Plan.
 - Main Heuristic Rule
 - apply first the the operations that reduce the size of intermediate results.
(Push down Selections, Projections to delay Cartesian Products)

• General Transformation Rules

$$1. \sigma_{c_1 \wedge c_2 \wedge \dots \wedge c_n}(R) = \sigma_{c_1}(\sigma_{c_2}(\dots(\sigma_{c_n}(R))))$$

• Cascading Selections

$$2. \sigma_{c_1}(Q_2(R)) = Q_2(\sigma_{c_1}(R))$$

• Commutativity of Selections

$$3. \pi_{a_1}(R) = \pi_{a_1}(\pi_{a_2}(\dots(\pi_{a_n}(R))))$$

• Cascading Projections

• Requires $a_i \subseteq a_{i+1} \forall i$



$$4. \pi_a(\sigma_c(R)) = \sigma_c(\pi_a(R))$$

• Selection Projection Comm.

• Requires that every attribute in c is in a .

$$5. R \bowtie_c S = S \bowtie_c R$$

• Join and \bowtie are Comm.
(also, \cup and \cap)

$$6. \sigma_c(R \bowtie_c S) = (\sigma_c(R)) \bowtie_c S$$

• Requires that c uses only attributes from R , else

$$\sigma_c(R \bowtie_c S) = \sigma_{c_1}(R) \bowtie_c \sigma_{c_2}(S)$$

• Computing σ with \bowtie

• c_1 uses attrib. only from R , c_2 only from S (not always possible)

$$7. \text{TC}_L(R \bowtie_c S) = \text{TC}_{A_1, \dots, A_n}(R) \bowtie_c \text{TC}_{B_1, \dots, B_n}(S)$$

. If $L = \{A_1, \dots, A_n, B_1, \dots, B_n\}$ and A_i only in R , B_i only in S

. if the Join Condition Contains additional attributes } no
 A_j, B_j not in L then they must be added to each } Problem
 Projection. } if x
 instead of \bowtie_c

$$8. R \bowtie (S \bowtie T) = (R \bowtie S) \bowtie T$$

. associativity
 of $\bowtie, \times, \cup, \cap$

$$9. \sigma_c(R \times S) = R \bowtie_c S$$

. Combine (x, c)
 into \bowtie

• Heuristic Algebraic Optimization Using Equiv.

1. Use Rule 1 to break up Selects

2. Use Rules 2, 4, 6 to move each Select operation as far down

3. Use Rules 5, 8 to rearrange tree leaf nodes so that more restrictive Selects are executed first.

4. Use Rule 9 to form joins

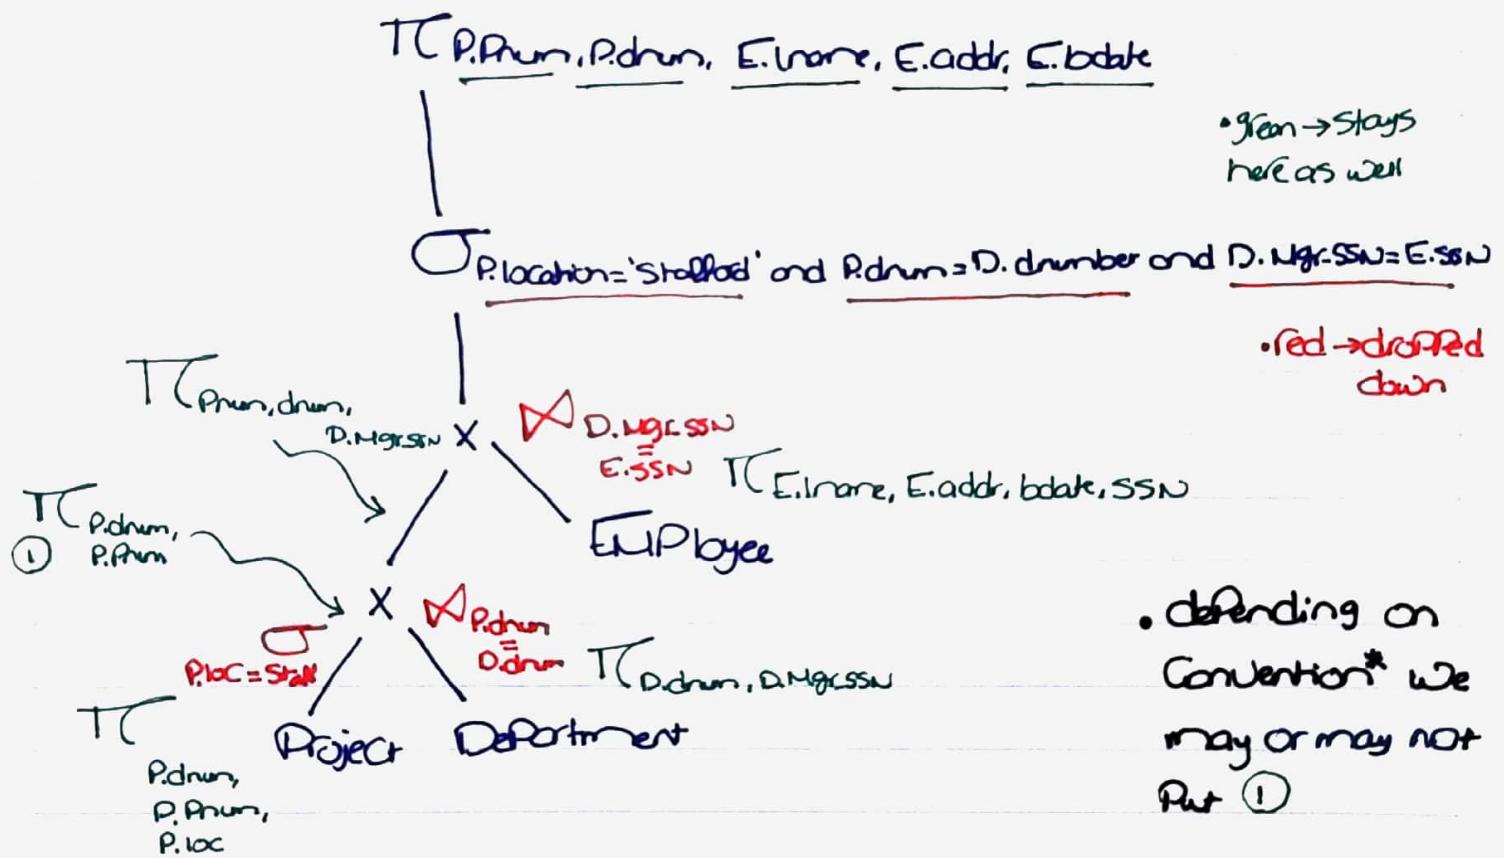
5. Use Rules 3, 4, 7 to break up and make down Projections (may need to create new TC in the process) as far as possible

6. Identify Subtrees that represent groups of operations that can be executed by a single algo.

Example

Select P.Pnum, P.dnum, E.lname, E.addr, E.bdate
 From Project P, department D, employee E
 Where P.dnum = D.dnumber and D.Mgr-ssn = E.SSN
 and P.location = 'Stafford'

. initial (canonical query tree) → heuristic Optimization



* Query Execution Plans

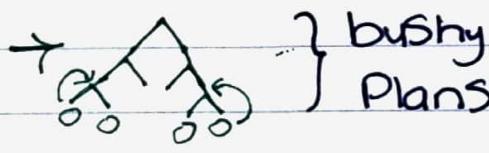
(Execution = Evaluation)

Materialized Evaluation

- the result of an operation is stored as a temporary relation on disk

Dirkted Evaluation

- one result of an operation is produced its forward to the next operator.



better & faster

deep { → requires a left/right most tree ↗ only write to disk here

Plans { }

- A query Optimizer doesn't rely only on heuristics, it also uses
- Cost-based Query Optimization

1. Estimate the Costs of executing a query using different execution Strategies

2. Compare and Choose Strategy with lowest Cost estimate

more Suitable For Compiled Queries (Store the Strategy along Side the query rather than having to figure it out during runtime)

→ Issues

- designing a Cost Function to estimate the actual cost of executing the query.

- No. of execution Strategies to be considered.
(usually non-Polynomial)

• Estimating the Cost of a Plan Query tree +..

1. For each node estimate the Cost of Performing the Corresponding operation

2. For each node estimate the Size of the result and its Properties (eg, Sorted)

→ important as result is input to next operator

• ① is Possible because For each access method / operation we know the cost given input size (Cardinality) // depends on ②

→ An estimate of the entire Pbn can hence be found by Combining estimates at each node

But how do we estimate output size (Cardinality)?

Selectivity Factor

→ Given a query $\underbrace{\quad\quad\quad}_{\# \text{tuples}}$

- Maximum Cardinality of result = Product of Cardinalities of Participating relations

(in the worst case where clause always true and it's just a cross product)

- But the where clause matters

→ Suppose relations A, B, C have cardinalities P_A, P_B, P_C (maximum cardinality of result is $P_A P_B P_C$)

→ then if the selectivity of each condition in where clause is S_1, S_2, S_3 , the total result size can be estimated at

$$(P_A P_B P_C)(S_1 S_2 S_3)$$

• assumes independence*

• Recall, the selectivity of a predicate (condition) in the where clause is the ratio of expected result size to maximum result size

Selectivity Estimates

⇒ Condition is Column = Value

$$S_c = \frac{1}{\#\text{Keys}(\text{Column})}$$

assumes uniform dist.

- i.e., if column takes one of 10 unique values and they're all equally likely then $S_c = P(\text{Col.} = \text{Value}) = \frac{1}{10}$

\Rightarrow Condition is Column 1 = Column 2.

$$S_c = \frac{1}{\max(\#\text{Keys}(C1), \#\text{Keys}(C12))}$$

→ Assume that Col1 has more unique values and that each value in Col2 will match with Col1. In this case, each value in Col2 matches with a probability $\frac{1}{\#\text{Keys}(C1)}$

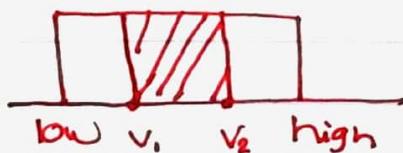
\Rightarrow Condition is Column > Value

$$S_c = \frac{\text{high(Column)} - \text{Value}}{\text{high(Column)} - \text{low(Column)}}$$

\Rightarrow Condition is Value1 < Column < Value2

$$S_c = \frac{\text{Value2} - \text{Value1}}{\text{high(Column)} - \text{low(Column)}}$$

, these two also follow from uniform dist. assumption (area under curve)



\Rightarrow Condition is Column in list

$$S_c = \frac{\text{len(list)}}{\#\text{Keys}(\text{Column})} \quad \left. \begin{array}{l} \text{For value in list} \\ \text{check } \text{Column} = \text{value} \\ (\text{P}(C=v_1) + \text{P}(C=v_2) + \dots) \end{array} \right\}$$

\Rightarrow Condition is Column in Subquery

$$S_c = \frac{|Q_1|}{\# \text{Keys}(\text{Column})}$$

- Let $|Q_1|$ be estimated size of subquery
- Treat like list

$$\Rightarrow S_{\bar{c}} = 1 - S_c \quad (c \text{ is the predicate, } \bar{c} \text{ is its negation})$$

$$\Rightarrow S_{C_1 \text{ or } C_2} = S_{C_1} + S_{C_2} - S_{C_1} \cdot S_{C_2}$$

- Assumes indep. of C_1, C_2
- addition rule of Prob.

• We have So Far assumed

① Values in a Single Column Follow a uniform distribution

- rarely holds
- Can be relaxed with better Statistical methods; namely, histograms.

② Values across Columns are uncorrelated.

- (e.g. For independence of C_1, C_2)
- rarely holds and had to relax

* Histograms

\rightarrow Data structure to Capture the actual Probability distribution of column values

\rightarrow Trade off between size and accuracy

- larger size \rightarrow more accuracy but more hard to maintain

$\xrightarrow{\hspace{1cm}}$ Two classes: equi-width & equi-depth

- Desirable Properties: Small, accurate, fast access

\downarrow e.g. single block access \downarrow \downarrow Single look up access

• Equi-Width Histogram

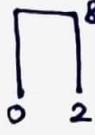
- total range of column value is divided into subranges of equal width. (buckets)
- histogram associates each bucket with the no. of tuples in its range

e.g., if column value ranges from 0 to 14 then a possible equi-width histogram is

min	max	Count
0	2	8
3	5	4
6	8	15
9	11	3
12	14	15
$\Sigma = 45$		

$$\text{width} = 3$$

- we will assume that all values in same bucket follow a uniform dist.

i.e.,  $P(0 \leq v \leq 2) = \frac{8}{45}$
 $P(v=0) = P(v=1) = P(v=2) = \frac{8}{45} \cdot \frac{1}{3}$

$$S_{6 \leq v \leq 10} = P(6 \leq v \leq 10)$$

$$= \frac{3}{3} \cdot \frac{15}{45} + \frac{2}{3} \cdot \frac{3}{45} = \frac{17}{45}$$

$6 \leq v < 8$ $9 \leq v \leq 10$

- if we were to assume a uniform dist.

$$S_{6 \leq v \leq 10} = \frac{10-6}{14-0} = \frac{12}{42}$$

it gets much worse if there's an outlier at 30.

errors

* Note that from this point onwards, S_C will denote the Output Size and \hat{S}_C will denote Selectivity (S_C in slides).

• Equi-depth Histogram

→ The total range of the Column Value is divided into Subranges (buckets) Such that the no. of tuples in each is approx. equal.

min	max	Count
0	3	8
4	7	10
8	9	10
10	13	7
14	14	9
$\Sigma = 44$		

→ to construct, set a threshold and look for smallest subrange satisfying it.

Some estimation algo.

$$\left\{ \begin{aligned} S_{6 \leq v \leq 10} &= \frac{2}{4} \cdot \frac{10}{44} + \frac{2}{2} \cdot \frac{10}{44} + \frac{1}{4} \cdot \frac{7}{44} = \frac{17}{44} \\ &\quad 4 \leq v \leq 7 \qquad \qquad \qquad 8 \leq v \leq 9 \qquad \qquad \qquad 10 \leq v \leq 10 \end{aligned} \right.$$

→ Can be proven to be more accurate than equi-width. (esp. Continuous data with no outliers)

• Catalog Information Used in Cost Functions

→ (to estimate cost of various exec. stat.)

- # Info about File Size
 - no. of records (r)
 - Record Size (R)
 - no. of blocks (b)
 - blocking factor (bfr)

- # Info about indexes
 - No. of levels (x)
 - No. of B+ level index blocks (blt)
 - No. of distinct values of attribute (d)
 - Selectivity of attribute Given Condition (S)
 - Selection cardinality of attr. (S * r = S')

$C_{S1}, S1 = \text{method}$

- Examples of Cost Functions for Select (in terms of # IOs of blocks)

S1. Linear Search by brute force

→ Search all file blocks (non-Key or not found)

$$C_{S1} = b$$

→ if it's a Key (no duplicates)

$$C_{S1} = b/2$$

S2. Binary Search

→ Non-Key

↓ to reach 1st block

$$C_{S2} = \log_2 b + \lceil \frac{S}{bfr} \rceil - 1$$

to read the remaining blocks

• if block holds 10 records and output cardinality S is 50 records then **15 more blocks to read**

→ Key

$$C_{S2} = \log_2 b$$

(even if unique, $S=1$ and formula above holds)

S_{3a}. Use a Primary index to retrieve a single record

$$C_{S3a} = X + 1$$

↓
to reach 1st
index level
(leaves)

Final I/O to
get needed data block

S_{3b}. Use a hash key to retrieve a single record

$$C_{S3b} = 1 \quad (\text{linear hashing})$$

- read the right bucket

$$C_{S3b} = 2 \quad (\text{extendible hashing})$$

- directory is also a block

clustered or Prim.
↓

S₄. Using an Ordering index to retrieve multiple records

$$C_{S4} = X + \frac{b}{2}$$

↓
to reach
leaves

↓
assuming an inequality
condition, roughly, on avg.
half records qualify.

S₅. Using a Clustered index to retrieve mult. records

$$C_{S5} = X + \lceil S/bfr \rceil$$

- like about but
S not necessarily
 $\frac{b}{2} r$ ($S_c = \frac{1}{2}$ not forced).

S_{6a}. Using a Secondary index for equality search

$$C_{S6a} = X + S$$

← block access for
each qualifying record
in index.

- replace S with S+1 if block of pointers
(unique index)

S6b. Secondary Index for multiple records $\langle, \langle, \rangle, \rangle$

$$C_{\text{sec}} = X + \frac{b_{12}}{2} + \frac{r}{2}$$

↓ ↓ ↓
 to reach 1/2 of block's and they correspond to
 leaves quality (leaves) 1/2 the records
(reach blocks
of Pointers)
• i.e., $r/2$ block Pointer on avg.

S7. Conjunctive Selection

- Use S1 (linear Search)
- or → use one of the conditions to retrieve the records } S2
 - (which condition depends on sort/hash options) then } to S6
 - Check the rest of the conditions in memory.

S8. Conjunctive Selection using Composite index

→ Special Case of S3a, S5, S6a

Example:

$$r = 10000 \quad b = 2000 \quad bfr = 5$$

Access Paths:

- Clustering index on salary with 3 levels and $S_{\text{sal}} = 20$
- Secondary index on key attrib. SSN with 4 levels and $S_{\text{SSN}} = 12$

Query: $\sigma_{\text{SSN}=1234}(\text{Employee})$

→ S1: $C_{\text{S1}} = b/2 = 1000$

→ S6a: $C_{\text{S6a}} = X + 1 = 4 + 1$

* The cost model for Cost-based Optimization in our case is #TIO

→ depends on cardinalities of inputs hence its a Cardinality-based cost model

G
Trivial

. Cost-based Query Optim. = Plan enumeration + Selection (min cost)

The Combinations of these two define the Search Space.

access method Selection

join algorithm Selection (enum.)

• Search Space is huge

• Plan enumeration = Search Space exploration

→ exhaustive exploration is hence a bad idea

• Often takes more than just evaluating the query

• Query Optimization method = Way to explore Search Space

Dynamic Programming

Randomised exploration
Rule-based Opt.

→ Why not just brute force?

• Consider

5 relations, 2 access methods for each and 3 possible join algorithms with left-deep plan

• In this case there are $5! \times 2^5 \times 3^4$ Possible Plans



* Dynamic Programming

1. Identify Cheapest way to access each relation given local predicates (Conditions)

- if one of the ways isn't cheapest overall but cheapest w.r.t an interesting order (helpful in future) then keep it as well. ①

2. For every relation and for every join condition, find the Cheapest way to join in a second

- ①

3. Keep applying Principle 2 until costs for the whole plan is found

Example:

• Want to optimally join A,B,C,D (3 available join algos)

1st iteration:

→ Find best way to access A, B, C, D

	Cache best	Cost
A	index	100
B	season	50
C	...	
D		

• 2nd iteration

→ 4 relations $\Rightarrow 4C_2$ Possible Pairs (joins)

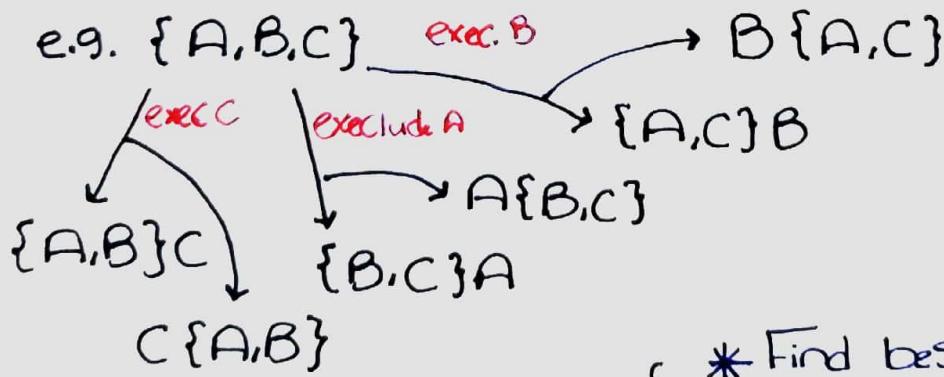
- $\{A, B\} \leftarrow A \bowtie B \text{ or } B \bowtie A$
- $\{A, C\}$
- $\{A, D\}$
- $\{B, C\}$
- $\{B, D\}$
- $\{C, D\}$

} Given the best access methods estimate best way to join each (order of relations and join method)
e.g.
 $\{A, B\} \rightarrow A \bowtie B$ with Sort merge
... record all with the cost in Cache.

- if 3 join methods, we need $3 \times 2 \times 4C_2$ Computations and 7C_2 results (4 is no. of rel.)

• 3rd iteration

→ $4C_3$ Possible trios and each ordered in $2 \times 3 = 6$ Possible ways (with 3 join possible for each)



The SubProblem in {} has already been Solved!

e.g. $(A \bowtie B) \bowtie C$

Sort merge block nested

} * Find best way to compute $\{A, B, C\}$
 $\rightarrow 6 \times 3$ Possibilities
 + join methods

- In total, $4C_3 \times 6 \times 3$ Computations and $4C_3$ results

• 4th iteration

→ $4C_4 \times 4 \times 2 \times 3$ Computations and $4C_4$ outputs
 ↓
 Only $\{A, B, C, D\}$ Excluded
 A, B, C, D before or after {}

• Clearly, final result gives the best plan.

→ Called Selinger Optimization

$R \leftarrow$ Set of relations to join

For ℓ in $1, 2, \dots, |R|$

For each subset S of length ℓ of R

① $\left\{ \begin{array}{l} \text{OPTjoin}(S) = a.\text{join}(S-a) \\ \text{where } a \text{ is the relation that minimizes} \\ \text{Cost(optjoin}(S-a)) + \min \text{Cost to} \\ \text{join } (S-a) \text{ and } a + \min \text{Cost to} \\ \text{access } a \end{array} \right.$

① e.g. For the subset $\{A, B, C\}$ we choose to join
one of the three with optimal result of the other two.
a S-a
while taking into account the cost of the new join
and cost to access a.

• Complexity

→ $\sum_{i=1}^n {}^n C_i$; Subsets are considered in total $= 2^n - 1$

• the 'powerset' mathematically includes all subsets
even $\{\}$ and has 2^n elements

• if we take it to be 2^n , this equivalent to the
no. of ways to write a binary of length n .

- For each subset length (K) we have K Possibilities
(removing each of the relations and using Previous subprob)

$\rightarrow K \leq n$

\longrightarrow Suppose we have m ways to evaluate
the join (6 in Prev. example)

\downarrow
 y
Comm. 3 methods

$$\text{Total Cost} = O(nm2^n)$$

\rightarrow Interesting Orders

- non-optimal joins Producing sorted data can help in future (e.g. if final result will be sorted)
- In Cache maintain best overall Plan and best interesting Order Plan

\Rightarrow Complexity becomes $O(2nm2^n)$

- double evaluations due to interesting order

* Suppose there were I interesting orders to consider

$$O((I+1)nm2^n)$$

- each time $I+1$ costs are estimated.

* Rule-based Optimization

- An issue of if-then rules

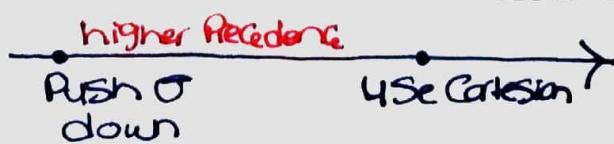
if (Condition-list)

, //Not where clause

apply transformation on Plan so far.

\rightarrow Keep the new Plan only if cheaper than original

- Order of rules matter (Precedence)



* Randomised Exploration

- Useful in big queries with huge search spaces
- to explore a bigger portion of it; let optimizer jump to a far state with some probability (Periodically).
 - . Otherwise can get stuck in a smaller region of search space.

. Oracle DBMS V8

USES

- Rule-based optimization (through heuristics)
- Cost-based optimization (optimizer is CPU & memory)
- allow developer hints to optimizer (they might know more)

Semantic Query Optimization

- uses constraints from DB schema (e.g. unique attr.)
- to modify query so its more efficient to execute (Generally).

e.g. query with EMP. Salary > Manager. Salary

- , There is an insertion constraint that employee can't have higher salary
- hence semantic query optimizer knows this has $S=0$