

Quiz

Hashing , B+trees, and Extendible Hashing


Advanced Database

Consider an extendible hashing structure such that:

- Each bucket can hold up to two records
- The hashing function uses the **lowest g bits** (i.e **least significant bits** of the number), where g is the global depth

- a. Starting from an empty table, insert 6, 15, 34, 18
 - i. What is the total depth of the resulting table?
 - ii. What is the local depth of the bucket containing 34?
- b. Starting from the result in (a), you insert keys 16, 7, 10, 20, 9
 - i. Which key will first cause a split without doubling the size of the table?
 - ii. Which key will first make the table double the size?

- a. Insert 24 to the following B+tree
- b. Delete 23 from the B+tree obtained in part a.



- Consider the relations **R(A, B)**, **S(C, D)** and **T(D, B, E)** with keys underlined, i.e., attribute A is a key for R, C is a key for S, and D is a key for T.
- Tuples from S and T take up 50 bytes**, while **tuples from R comprise 40 bytes**. Values for attribute **S.D** take up **10 bytes**.
- The **page size is set to 4000 bytes** and there are **8 pages in the buffer**.
- The statistics show that **S contains 36000 tuples**; that **T contains 10000 tuples**; that **R contains 12000 tuples**.
- Furthermore, values in **S.D** are **uniformly distributed in the range [0,100]**, values in **T.E** are **uniformly distributed in the range [10000 , 30000]**, and **T** has **5000 distinct values for attribute B**. In addition, the following **histogram information is available for R.B**:

range	number of tuples
[0,19]	5000
[20,39]	200
[40,59]	800
[60,79]	2000
[80,100]	4000

- Moreover, relation S has a clustered B+-tree index on attribute D;
relation T has unclustered hash indexes on attributes B and E (separately).
- Assume that we only consider **left-deep plans**, and two join algorithms (**block nested-loops**, and **sort-merge join**) only. Use **Selinger Optimizer (dynamic programming algorithm)** to find the optimal cost-based query-plan to the following SQL query (i.e. find the best query tree and the algorithm associated with each operator in the tree).
- SELECT R.A, S.C, T.D FROM R,S, T
WHERE R.B==30 AND S.D <=60 AND T.E != 10000 AND R.B == T.B AND S.D == T.D**


Answers

Consider an extendible hashing structure such that:


- Each bucket can hold up to two records
- The hashing function uses the **lowest g bits** (i.e **least significant bits** of the number), where g is the global depth

- a. Starting from an empty table, insert 6, 15, 34, 18
 - i. What is the total depth of the resulting table? $\Rightarrow 3$
 - ii. What is the local depth of the bucket containing 34? $\Rightarrow 3$
- b. Starting from the result in (a), you insert keys 16, 7, 10, 20, 9
 - i. Which key will first cause a split without doubling the size of the table? $\Rightarrow 9$
 - ii. Which key will first make the table double the size? $\Rightarrow 10$

- a. Insert 24 to the following B+tree
- b. Delete 23 from the B+tree obtained in part a.



- a. Insert 24 to the following B+tree
- b. Delete 23 from the B+tree obtained in part a.



- Consider the relations **R(A, B)**, **S(C, D)** and **T(D, B, E)** with keys underlined, i.e., attribute A is a key for R, C is a key for S, and D is a key for T.
- Tuples from S and T take up 50 bytes**, while **tuples from R comprise 40 bytes**. Values for attribute **S.D** take up **10 bytes**.
- The **page size is set to 4000 bytes** and there are **8 pages in the buffer**.
- The statistics show that **S contains 36000 tuples**; that **T contains 10000 tuples**; that **R contains 12000 tuples**.
- Furthermore, values in **S.D** are **uniformly distributed in the range [0,100]**, values in **T.E** are **uniformly distributed in the range [10000 , 30000]**, and **T** has **5000 distinct values for attribute B**. In addition, the following **histogram information is available for R.B**:

range	number of tuples
[0,19]	5000
[20,39]	200
[40,59]	800
[60,79]	2000
[80,100]	4000

- Moreover, relation S has a clustered B+-tree index on attribute D;
relation T has unclustered hash indexes on attributes B and E (separately).
- Assume that we only consider **left-deep plans**, and two join algorithms (**block nested-loops**, and **sort-merge join**) only. Use **Selinger Optimizer (dynamic programming algorithm)** to find the optimal cost-based query-plan to the following SQL query (i.e. find the best query tree and the algorithm associated with each operator in the tree).
- SELECT R.A, S.C, T.D FROM R,S, T
WHERE R.B==30 AND S.D <=60 AND T.E != 10000 AND R.B == T.B AND S.D == T.D**

$$\pi_{R.A, S.C, T.D} ((\sigma_{R.B=30}(R)) \bowtie_{R.B=T.B} (\sigma_{T.E \neq 10000}(T)) \bowtie_{S.D=T.D} (\sigma_{S.D \leq 60}(S)))$$

$\pi_{R,A,S,C,T,D}((\sigma_{R,B=30}(R)) \bowtie_{R,B=T,B} (\sigma_{T,E \neq 10000}(T)) \bowtie_{T,D=S,D} (\sigma_{S,D \leq 60}(S)))$

The best way to select one relation

- $\sigma_{R,B=30}(R)$
 - Full table scan to search the values (no index nor unique values)
= (number of tuples) * (record size) / (page size)
= 12000 * 40 / 4000 = 120 IOs
 - 200 tuples [20,39] **(from givens)**
 - Cardinality = 1/20 * 200 = 10 tuples = 400 bytes = 1 block
- $\sigma_{T,E \neq 10000}(T)$
 - Hash indexes are not useful for inequality predicates.
 - Full table scan = (number of tuples) * (record size) / (page size)
= 10000 * 50 / 4000 = 125 IOs
 - T.E are uniformly distributed in the range [10000 , 30000] **(from givens)**
 - Cardinality = 20000/20001 * 10000 = 10000 tuples = 500000 bytes = 125 blocks
- $\sigma_{S,D \leq 60}(S)$
 - #IOs = index access (using the B+tree index) + fetch blocks
 - S.D are uniformly distributed in the range [0,100] **(from givens)**
 - Cardinality = 61/101 * 36000 = 21743 tuple = 1087150 bytes = 272 blocks
 - Index record \Rightarrow D key (10 bytes) + pointer (assume 10 bytes) \Rightarrow 20 bytes
 \Rightarrow 4000 / 20 = 200 index entry / page
 - The relation can fit into 36000*50/4000 = 450 block
 - First level = ceil(450 / 200) = 3 pages
 - Second level have only 3 pointers to the first level \Rightarrow 2 level index is needed
 - #IOs = 2 + 272 = 274 IOs

$$\pi_{R,A,S,C,T,D} ((\sigma_{R.B=30}(R)) \bowtie_{R.B==T.B} (\sigma_{T.E != 10000}(T)) \bowtie_{T.D==S.D} (\sigma_{S.D \leq 60}(S)))$$

The best way to select two relations

- $(\sigma_{R.B=30}(R)) \bowtie_{R.B==T.B} (\sigma_{T.E != 10000}(T))$
 - Nested loop join
 - Cost = $M + \text{ceil}(M / (B-2)) * N \Rightarrow (M \text{ is the smaller relation})$
 $= 1 + \text{ceil}(1/6) * 125 = 126 \text{ IOs}$
 - Sort merge join
 - Cost = $2M(1 + \text{ceil}(\log_{B-1}(\text{ceil}(M/B)))) + 2N(1 + \text{ceil}(\log_{B-1}(\text{ceil}(N/B)))) + M + N$ (higher cost)
 - Total cost = current cost + sub-problems cost = $126 + 120 + 125 = 371 \text{ IOs}$
 - Cardinality = $10 * 10000 / \max(1,5000(\text{number of distinct values in T.B})) = 20 \text{ tuples} = 1 \text{ block}$
- $(\sigma_{T.E != 10000}(T)) \bowtie_{T.D==S.D} (\sigma_{S.D \leq 60}(S))$
 - Nested loop join
 - Cost = $M + \text{ceil}(M / (B-2)) * N \Rightarrow (M \text{ is the smaller relation})$
 $= 125 + \text{ceil}((125)/(8-2)) * 272 = 5837 \text{ IOs}$
 - Sort merge join
 - Note that the relation S is sorted on the attribute D \Rightarrow Ignore sorting S
 - Cost = $2M(1 + \text{ceil}(\log_{B-1}(\text{ceil}(M/B)))) + 2N(1 + \text{ceil}(\log_{B-1}(\text{ceil}(N/B)))) + M + N$
 $= 2 * 125 * 3 + 2 * 272 * 3 + 125 + 272 = 1147 \text{ IOs}$
 - Total cost = current cost + sub-problems cost = $1147 + 274 + 125 = 1546 \text{ IOs}$
 - Cardinality = $21743 * 10000 / \max(61,10000) = 21743 \text{ tuples}$
 $\Rightarrow 21743 * 100 / 4000 = 544 \text{ blocks}$

$$\pi_{R,A,S,C,T,D} ((\sigma_{R,B=30}(R)) \bowtie_{R,B==T.B} (\sigma_{T,E !=10000}(T)) \bowtie_{T,D==S.D} (\sigma_{S,D \leq 60}(S)))$$

The best way to select three relations

- $((\sigma_{R,B=30}(R)) \bowtie_{R,B==T.B} (\sigma_{T,E !=10000}(T))) \bowtie_{T,D==S.D} (\sigma_{S,D \leq 60}(S))$

Continue by yourself

- Nested loop join

- $\text{Cost} = M + \text{ceil}(M / (B-2)) * N \Rightarrow (M \text{ is the smaller relation})$
 - =

- Sort merge join

- $\text{Cost} = 2M(1 + \text{ceil}(\log_{B-1}(\text{ceil}(M/B)))) + 2N(1 + \text{ceil}(\log_{B-1}(\text{ceil}(N/B)))) + M + N$
 - =

- Total cost = current cost + sub-problems cost =

- $((\sigma_{T,E !=10000}(T)) \bowtie_{T,D==S.D} (\sigma_{S,D \leq 60}(S))) \bowtie_{R,B==T.B} (\sigma_{R,B=30}(R))$

- Nested loop join

- $\text{Cost} = M + \text{ceil}(M / (B-2)) * N \Rightarrow (M \text{ is the smaller relation})$
 - =

- Sort merge join

- $\text{Cost} = 2M(1 + \text{ceil}(\log_{B-1}(\text{ceil}(M/B)))) + 2N(1 + \text{ceil}(\log_{B-1}(\text{ceil}(N/B)))) + M + N$
 - =

- Total cost = current cost + sub-problems cost =