

## ADB Sheet 3 Sol.

15.13.

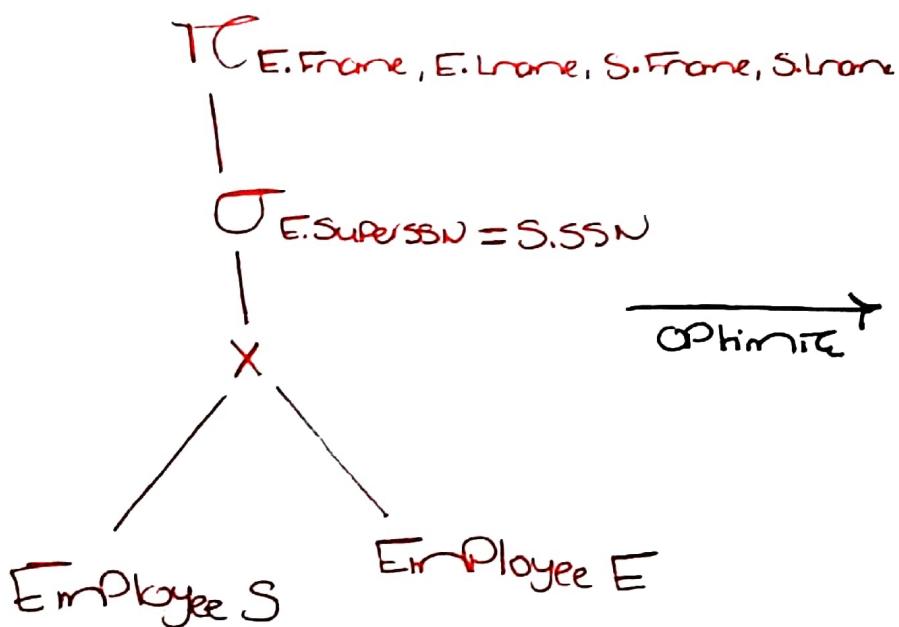
→ Consider SQL queries Q1, Q8, Q13, Q4, Q27 from the book

- For each, draw an initial and optimized query tree

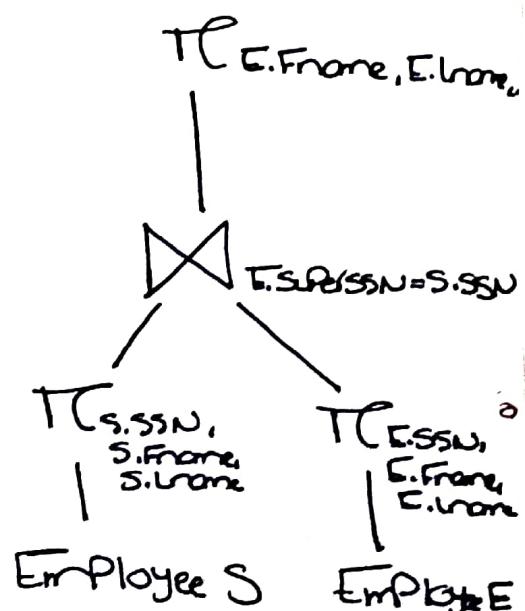
Q  
Circ.

Q8. Select E.Fname, E.Lname, S.Fname, S.Lname  
From Employee E, Employee S  
Where E.SuperSSN = S.SSN

- Initial Query Tree



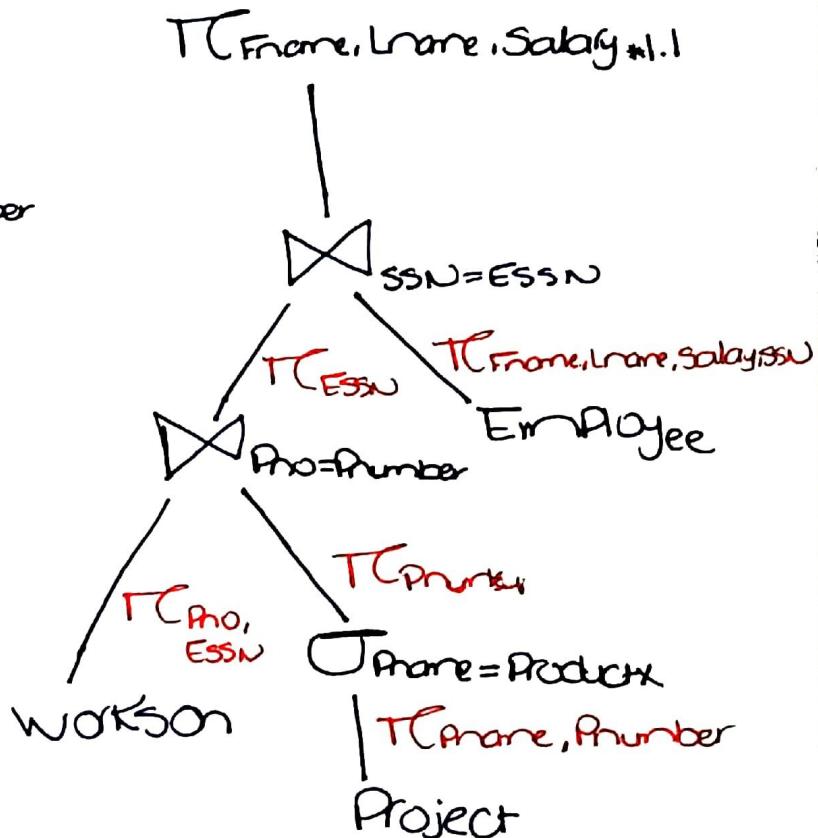
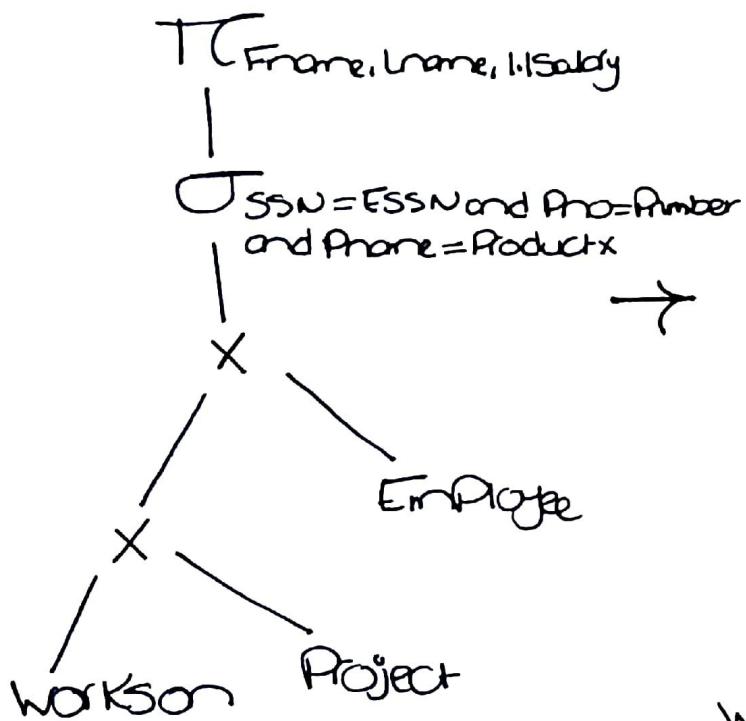
Optimize



Q27. Select Fname, Lname, I.I.\*Salary  
 From Employee, Works-on, Project  
 Where SSN=ESSN and Proj=Phumber and  
 Phone='Product X'

- works-on
- ESSN=SSN

Initial Query Tree



15.14.

- File of 4096 blocks
- needs to be sorted
- have 64 buffer blocks

- How many Passes for merge Phase & Sort-merge algo. (external Sort)?

Recall,

$$\# \text{IOs} = 2N \left( \overbrace{\log_{B-1} \lceil \frac{N}{B} \rceil}^{\text{No of Passes}} \right)$$

• For B=2 only.

$$N = 4096 \text{ blocks}$$

$$B = 64 \text{ blocks}$$

$$\text{No. of Passes} = 1 + \lceil \log_{63} \lceil \frac{4096}{64} \rceil \rceil = 3$$

15.15.

→ Develop approximate Cost functions for Project, Union, Intersection, Set-difference, Cartesian Product.

- let relations R and S be stored in M and N disk blocks respectively
- let the result take X blocks (or ignore final wr.)  
↳ unless known lecture

1. Project Operation on R

→ also eliminates duplicate rows  
hence,

- if one of the attributes is a Key or they form a Composite Key

$$\rightarrow \# IOs = M + M$$

- read all blocks
- write blocks again
- choose columns in memory

- else, we may have duplicates

$$\rightarrow \# IOs = \underbrace{2M(1 + \lceil \log_2 M \rceil)}_{\substack{\text{eliminate} \\ \text{duplicates} \\ \text{as you sort}}} \rightarrow \underbrace{\text{external sort}}_{\# IOs}$$

ITA

2. Set Operations (union\*, difference, intersection)

→ Implemented via Sort-merge join

$$\# IOs = O(M \log_2 M) + O(N \log_2 N) + M + N$$

+ X ← write join result

### 3. Cartesian Product

- Can be done via Sort-merge (like previous)
- Typically implemented with simple nested loop (refined) or block nested loop.

$$X = M \cdot N$$

- Simple nested loop (refined)
- Block nested loop

→ like simple nested loop but  $B > 3$

$$\# \text{IOs} = M + M \cdot N + X$$

↓                    ↓                    ↓  
 Scan R      scan S for      final  
 (outer loop)   every block   write

no. of records  
in R

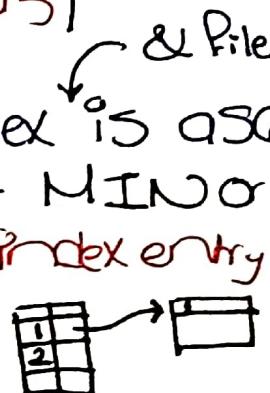
$$* \quad X = \frac{|R| \cdot |S|}{blk(R,S)} \quad \leftarrow \begin{array}{l} \text{no. of resulting records} \\ \text{\#REMEMBER ME (2 Pages)} \end{array}$$

when the record length is that of R + that of S how many records does a block hold

- 15.17. Can a sparse index be used to imp. an aggregate operator? (index-only plan)

- A sparse index doesn't contain all so AVG, Sum aren't possible (and generally can't unless its secondary non-key → block of pointers)

- Depending on whether the index is ascendingly or descendingly sorted can get MIN or MAX respectively by accessing first index entry.



15.21.

→ Extend Sort-Merge algo. to implement Left Outer Join.

•  $R \text{ } \bowtie_c S$

→ matches records in R and S according to C

→ If a record in R has no match, Put null for S columns and add it to result

• Sort-Merge algo. for Left Outer Join (<sup>general</sup><sub>idea</sub>)

1. Sort R, S via external sort

2. merge-Step

→ Advance S Pointer until its  $>$  that of R

→ If  $\text{IP} =$  then add the match to result  
if  $>$  (no match) → Put nulls for S then add.

→ Advance R Pointer by 1

\* Although this should be okay, TA attempted to join while 'Sorting'

15.22.

→ Compare the cost of two different query plans for:

$\sigma_{\text{Salary} > 40K} (E \bowtie_{\text{Dno}=\text{Dnumber}} D)$

• Given Statistics

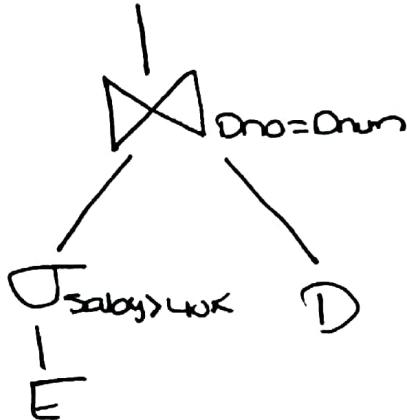
→ E has 2000 blocks with 10000 records

→ D has .5 blocks with 50 records

→ Salary ranges from 1 to 500K

## Tree 1 (Optimized)

$$\rightarrow \# \text{Keys}(\text{chun}) = \# \text{Keys}(\text{do}) \\ = 50$$



→ Can Conclude:

$$bFr(e) = \frac{10000}{2000} = 5 \text{ sec/block}$$

$$bFr(d) = \frac{50}{5} = 10 \text{ sec/block}$$

- Employee record must be twice as long
- One block should hold  $\boxed{7}$  records of concatenated records ( $bFr(e, d)$ )

For tree 1 :  $\underbrace{\text{read } E}_{\text{read}} \quad \underbrace{\text{write the result} = 1840}_{\text{write}}$

$$\# IOs_1 = 2000 + \frac{460K}{500K} \cdot 2000 \quad // \text{For } \sigma$$

$$\left. \begin{array}{l} \text{Sort merge} \\ \text{join before} \\ \text{Final write} \end{array} \right\} + 2 \times 5 \times (1 + \lceil \log_2 5 \rceil) + 2 \times 1840 \times (1 + \lceil \log_2 1840 \rceil) \quad \overbrace{\qquad \qquad \qquad}^{11}$$

$$\left. \begin{array}{l} \\ \\ \end{array} \right\} + 5 + 1840$$

$$\left. \begin{array}{l} \text{Final} \\ \text{write} \end{array} \right\} + \frac{10,000 \times 50 \times \frac{46}{50} \times \frac{1}{50}}{bFr(e, d)} \quad \leftarrow \text{no. of records from join}$$

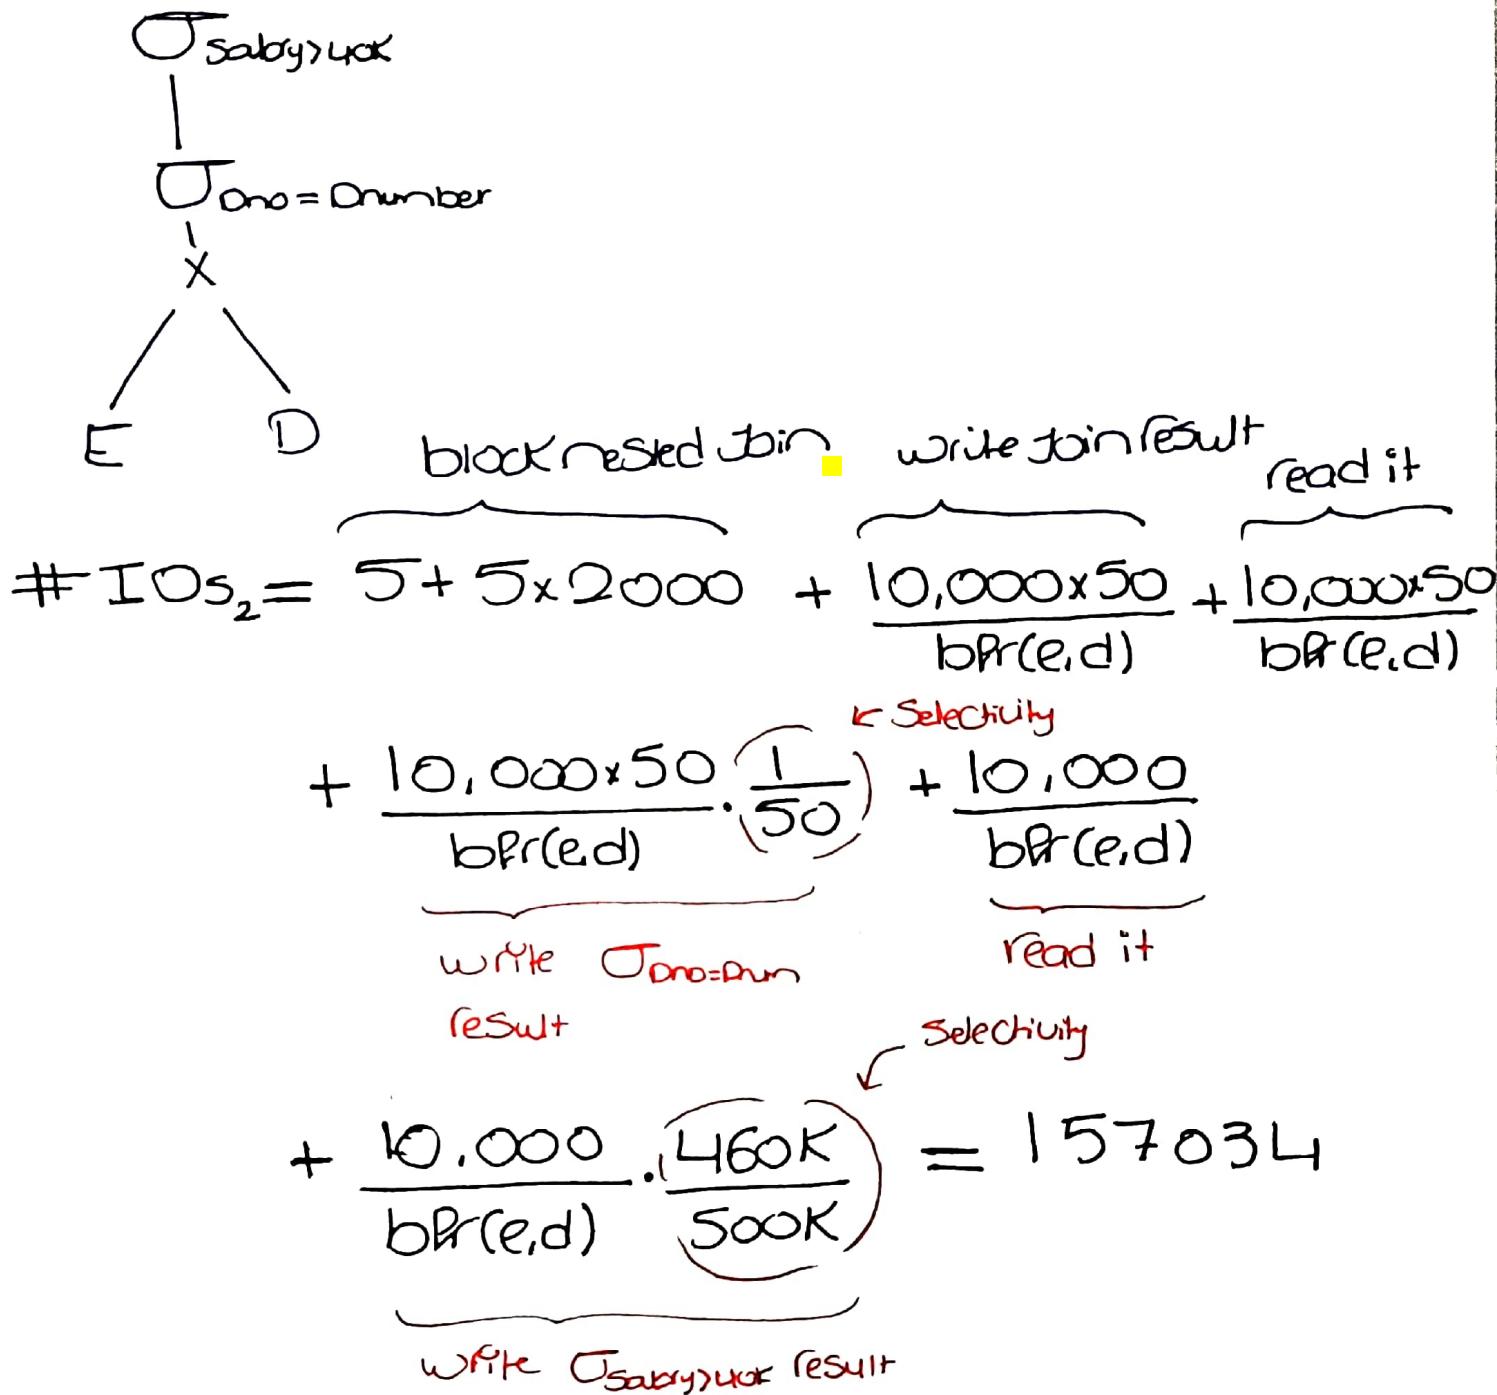
$$= 49355$$

→ Selectivities

$$S_{\text{Salary.}} = \frac{46}{50}$$

$$S_{\text{Dno=Dnum}} = \frac{1}{\max(50, 50)}$$

## Tree 2:



- Observe that

- We used block nested as this is how X is often implemented.
- For the specific M,N it does better than Sort merge