

Chapter 13: Disk Storage, Basic File Structures, and Hashing


Chapter Outline

- | Disk Storage Devices
- | Files of Records
- | Operations on Files
- | Unordered Files
- | Ordered Files
- | Hashed Files
 - Dynamic and Extendible Hashing Techniques

Disk Storage Devices

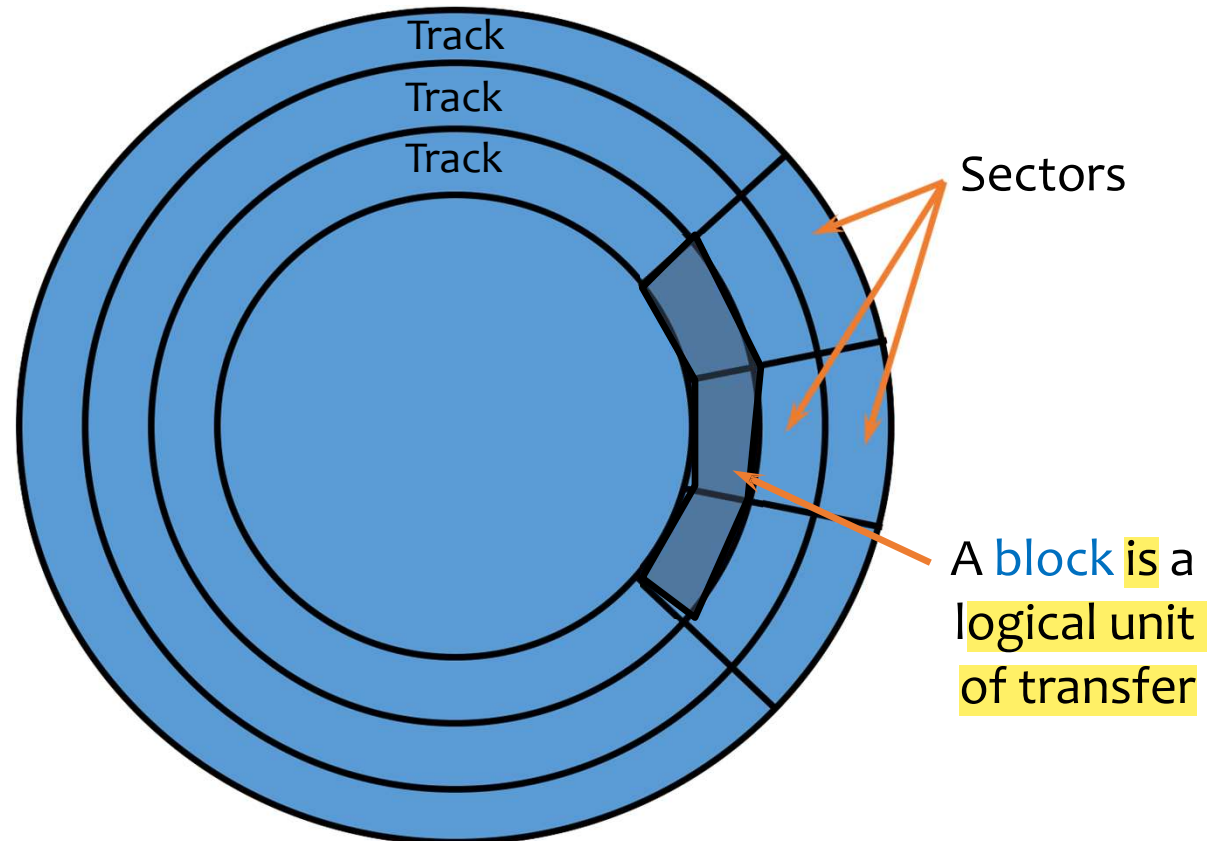
- | Data stored as magnetized areas on magnetic disk surfaces.
- | A disk pack contains several magnetic disks connected to a rotating spindle.
- | Disks are divided into concentric circular tracks on each disk surface.
 - | Track capacities vary typically from 4 to 50 Kbytes or more
- | A track is divided into smaller blocks or sectors
 - | because it usually contains a large amount of information

A typical hard drive



Top view

“Zoning”: more sectors/data on outer tracks



Disk access time

Sum of:

- **Seek time**: time for disk heads to move to the correct cylinder
- **Rotational delay**: time for the desired block to rotate under the disk head
- **Transfer time**: time to read/write data in the block
(= time for disk to rotate over the block)

Data on External Storage

- Data must persist on disk across program executions in a DBMS
 - Data is huge
 - Must persist across executions
 - But has to be fetched into main memory when DBMS processes the data
- The unit of information for reading data from disk, or writing data to disk, is a page
- ✓ • Disks: Can retrieve random page at fixed cost
 - But reading several consecutive pages is much cheaper than reading them in random order

Disk Space Management

- Lowest layer of DBMS software manages space on disk
- Higher levels call upon this layer to:
 - allocate/de-allocate a page
 - read/write a page
- Size of a page = size of a disk block
= data unit
- Request for a sequence of pages often satisfied by allocating contiguous blocks on disk
- Space on disk managed by Disk-space Manager
 - Higher levels don't need to know how this is done, or how free space is managed

Buffer Management

Suppose

- 1 million pages in db, but only space for 1000 in memory
- A query needs to scan the entire file
- DBMS has to
 - bring pages into main memory
 - decide which existing pages to replace to make room for a new page
 - called Replacement Policy
- Managed by the Buffer manager
 - Files and access methods ask the buffer manager to access a page mentioning the “record id” (soon)
 - Buffer manager loads the page if not already there

Files of Records

- Page or block is OK when doing I/O, but higher levels of DBMS operate on records, and files of records
- FILE: A collection of pages, each containing a collection of records
- Must support:
 - insert/delete/modify record
 - read a particular record (specified using record id)
 - scan all records (possibly with some conditions on the records to be retrieved)

File Organization

- **File organization:** Method of arranging a file of records on external storage
 - One file can have multiple pages
 - **Record id (rid)** is sufficient to physically locate the page containing the record on disk
 - **Indexes** are data structures that allow us to find the record ids of records with given values **in index search key** fields
- **NOTE: Several uses of “keys” in a database**
 - Primary/foreign/candidate/super keys
 - Index search keys

Alternative File Organizations

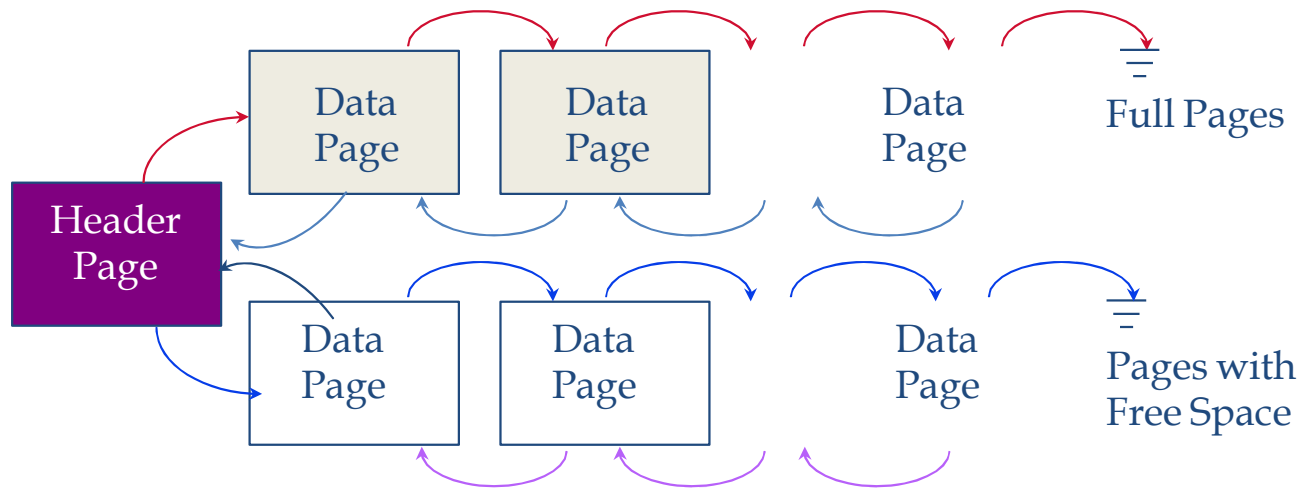
Many alternatives exist, each ideal for some situations, and not so good in others:

- **Heap (random order) files:** Suitable when typical access is a file scan retrieving all records
- **Sorted Files:** Best if records must be retrieved in some order, or only a “range” of records is needed.
- **Indexes:** Data structures to organize records via trees or hashing
 - Like sorted files, they speed up searches for a subset of records, based on values in certain (“search key”) fields
 - Updates are much faster than in sorted files

Unordered (Heap) Files

- Simplest file structure contains records in no particular order
- As file grows and shrinks, disk pages are allocated and de-allocated
- To support record level operations, we must:
 - keep track of the *pages* in a file
 - keep track of *free space* on pages
 - keep track of the *records* on a page
- There are many alternatives for keeping track of this

Heap File Implemented as a List



- The header page id and Heap file name must be stored someplace
- Each page contains 2 'pointers' plus data
- **Problem?**
 - to insert a new record, we may need to scan several pages on the free list to find one with sufficient space

How do we arrange a collection of records on a page?

- Each page contains several slots
 - one for each record
- Record is identified by <page-id, slot-number>
- Fixed-Length Records
- Variable-Length Records
- For both, there are options for
 - Record formats (how to organize the fields within a record)
 - Page formats (how to organize the records within a page)

Ordered Files

- . Also called a **sequential** file.
- . File records are kept sorted by the values of an ordering field.
- . Insertion is expensive: records must be inserted in the correct order.
 - It is common to keep a separate unordered *overflow* (or *transaction*) file for new records to improve insertion efficiency; this is periodically merged with the main ordered file.
- . A **binary search** can be used to search for a record on its ordering field value.
 - This requires reading and searching \log_2 of the file blocks on the average, an improvement over linear search.
- . Reading the records in order of the ordering field is quite efficient.

Ordered Files (contd.)

	NAME	SSN	BIRTHDATE	JOB	SALARY	SEX
block 1	Aaron, Ed					
	Abbott, Diane					
	Acosta, Marc					
block 2	Adams, John					
	Adams, Robin					
	Akers, Jan					
block 3	Alexander, Ed					
	Alfred, Bob					
	Allen, Sam					
block 4	Allen, Troy					
	Anders, Keith					
	Anderson, Rob					
block 5	Anderson, Zach					
	Angeli, Joe					
	Archer, Sue					
block 6	Arnold, Mack					
	Arnold, Steven					
	Atkins, Timothy					
block n-1	Wong, James					
	Wood, Donald					
	Woods, Manny					
block n	Wright, Pam					
	Wyatt, Charles					
	Zimmer, Byron					

Average Access Times

The following table shows the average access time to access a specific record for a given type of file

TABLE 13.2 AVERAGE ACCESS TIMES FOR BASIC FILE ORGANIZATIONS

TYPE OF ORGANIZATION	ACCESS/SEARCH METHOD	AVERAGE TIME TO ACCESS A SPECIFIC RECORD
Heap (Unordered)	Sequential scan (Linear Search)	$b/2$
Ordered	Sequential scan	$b/2$
Ordered	Binary Search	$\log_2 b$

Hashed Files

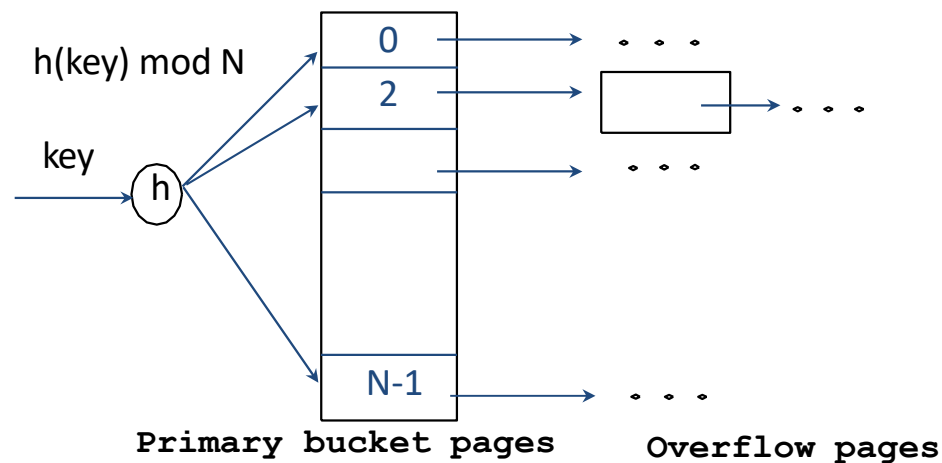
- Hashing for disk files is called **External Hashing**
- The file blocks are divided into **M** equal-sized **buckets**, numbered $\text{bucket}_0, \text{bucket}_1, \dots, \text{bucket}_{M-1}$.
 - Typically, a bucket corresponds to one (or a fixed number of) disk block.
- One of the file fields is designated to be the **hash key** of the file.
- The record with hash key value **K** is stored in bucket **i**, where $i=h(K)$, and **h** is the **hashing function**. E.g., $h(k)= K \bmod M$
- Search is very efficient on the hash key.
- Collisions occur when a new record hashes to a bucket that is already full.
 - An overflow file is kept for storing such records.
 - Overflow records that hash to each bucket can be linked together.

Hashed Files (contd.)

- There are numerous methods for collision resolution, including the following:
 - **Open addressing:** Proceeding from the occupied position specified by the hash address, the program checks the subsequent positions in order until an unused (empty) position is found.
 - **Chaining:** For this method, various overflow locations are kept, usually by extending the array with a number of overflow positions. In addition, a pointer field is added to each record location. A collision is resolved by placing the new record in an unused overflow location and setting the pointer of the occupied hash address location to the address of that overflow location.
 - **Multiple hashing:** The program applies a second hash function if the first results in a collision. If another collision results, the program uses open addressing or applies a third hash function and then uses open addressing if necessary.

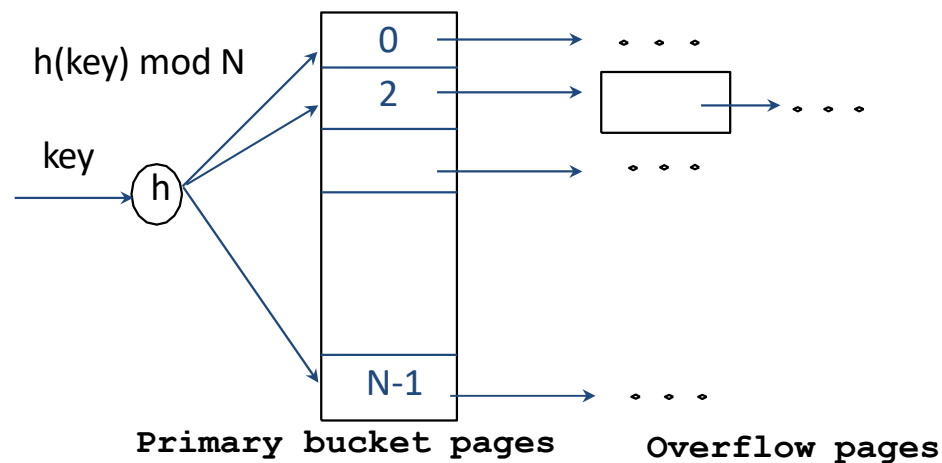
Static Hashing

- Pages containing data = a collection of **buckets**
 - each bucket has one primary page, also possibly overflow pages
 - buckets contain **data entries k^***



Static Hashing

- # primary pages fixed
 - allocated sequentially, never de-allocated, overflow pages if needed.
- $h(k) \bmod N$ = bucket to which data entry with key k belongs
 - N = # of buckets



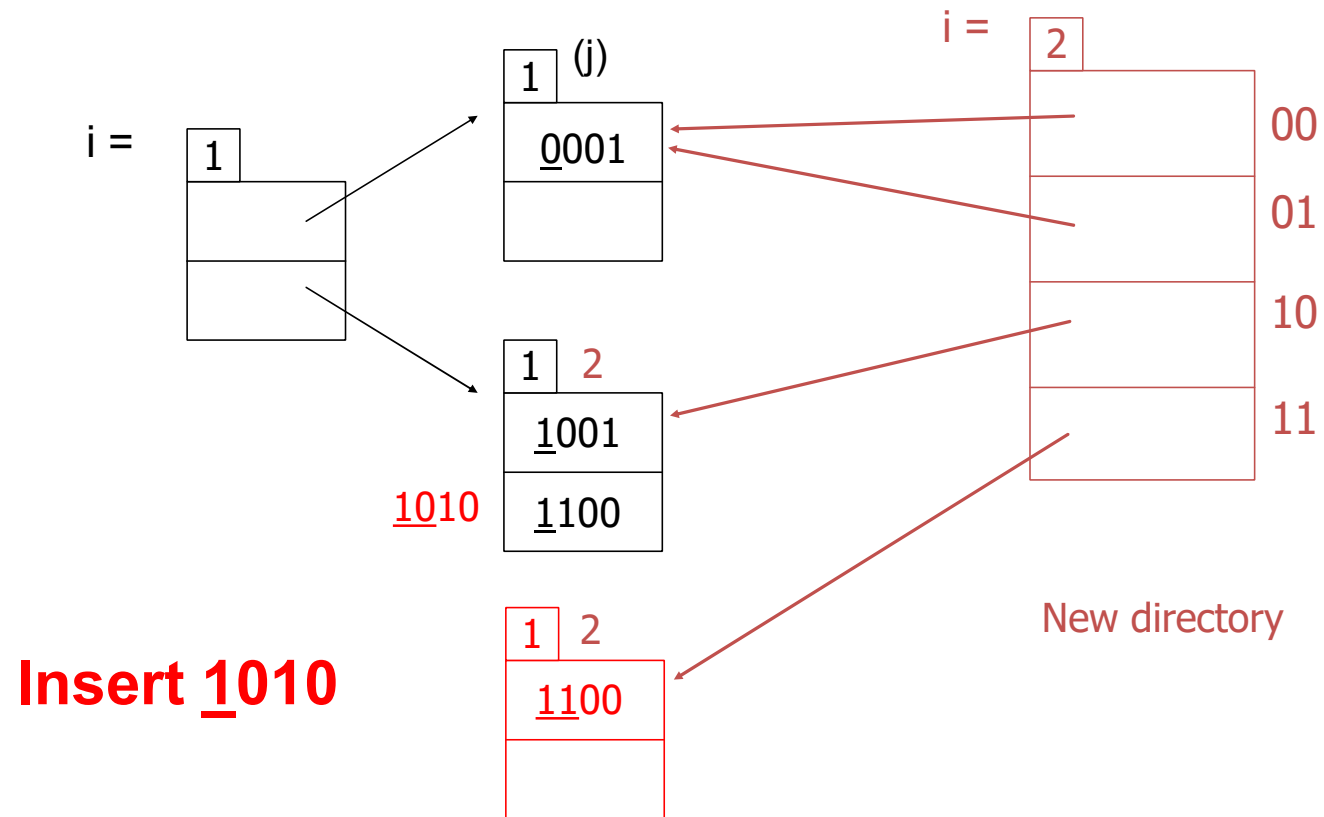
Static Hashing

- Hash function works on search key field of record r
 - Must distribute values over range $0 \dots N-1$
 - $h(\text{key}) = (a * \text{key} + b)$ usually works well
 - $\text{bucket} = h(\text{key}) \bmod N$
 - a and b are constants – chosen to tune h
- Advantage:
 - #buckets known – pages can be allocated sequentially
 - search needs 1 I/O (if no overflow page)
 - insert/delete needs 2 I/O (if no overflow page) (why 2?)
- Disadvantage:
 - Long overflow chains can develop if file grows and degrade performance (data skew)
 - Or waste of space if file shrinks
- Solutions:
 - keep some pages say 80% full initially
 - Periodically rehash if overflow pages (can be expensive)
 - or use Dynamic Hashing

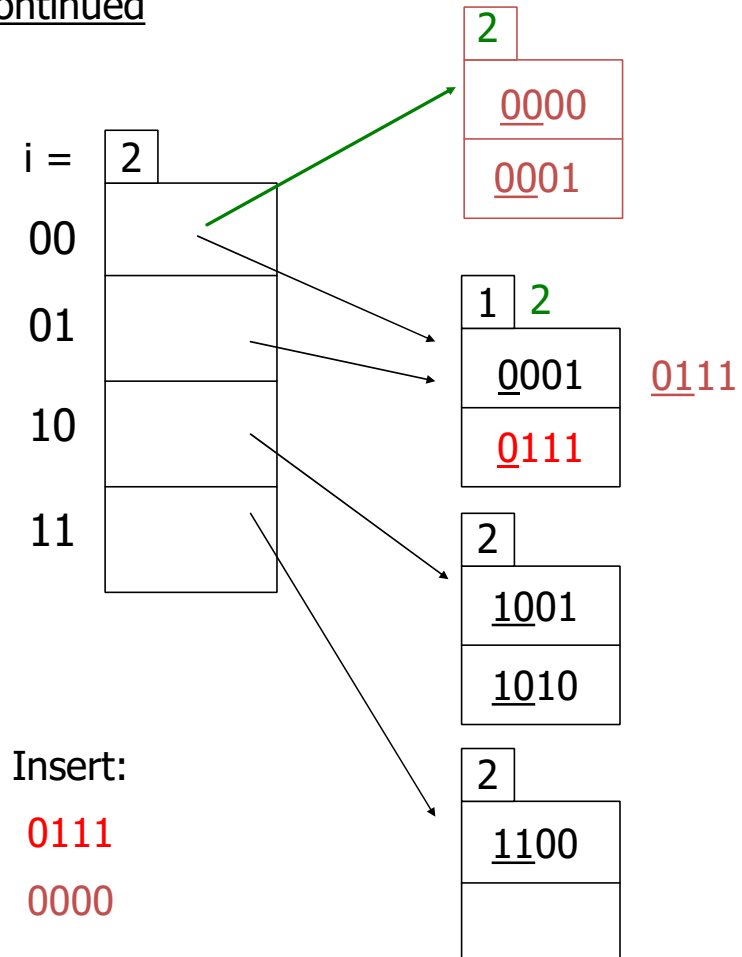
Extendible Hashing

- Consider static hashing
- Bucket (primary page) becomes full
- Why not re-organize file by doubling # of buckets?
 - Reading and writing (double #pages) all pages is expensive
- Idea: Use directory of pointers to buckets
 - double # of buckets by doubling the directory, splitting just the bucket that overflowed
 - Directory much smaller than file, so doubling it is much cheaper
 - Only one page of data entries is split
 - No overflow page (new bucket, no new overflow page)
 - Trick lies in how hash function is adjusted

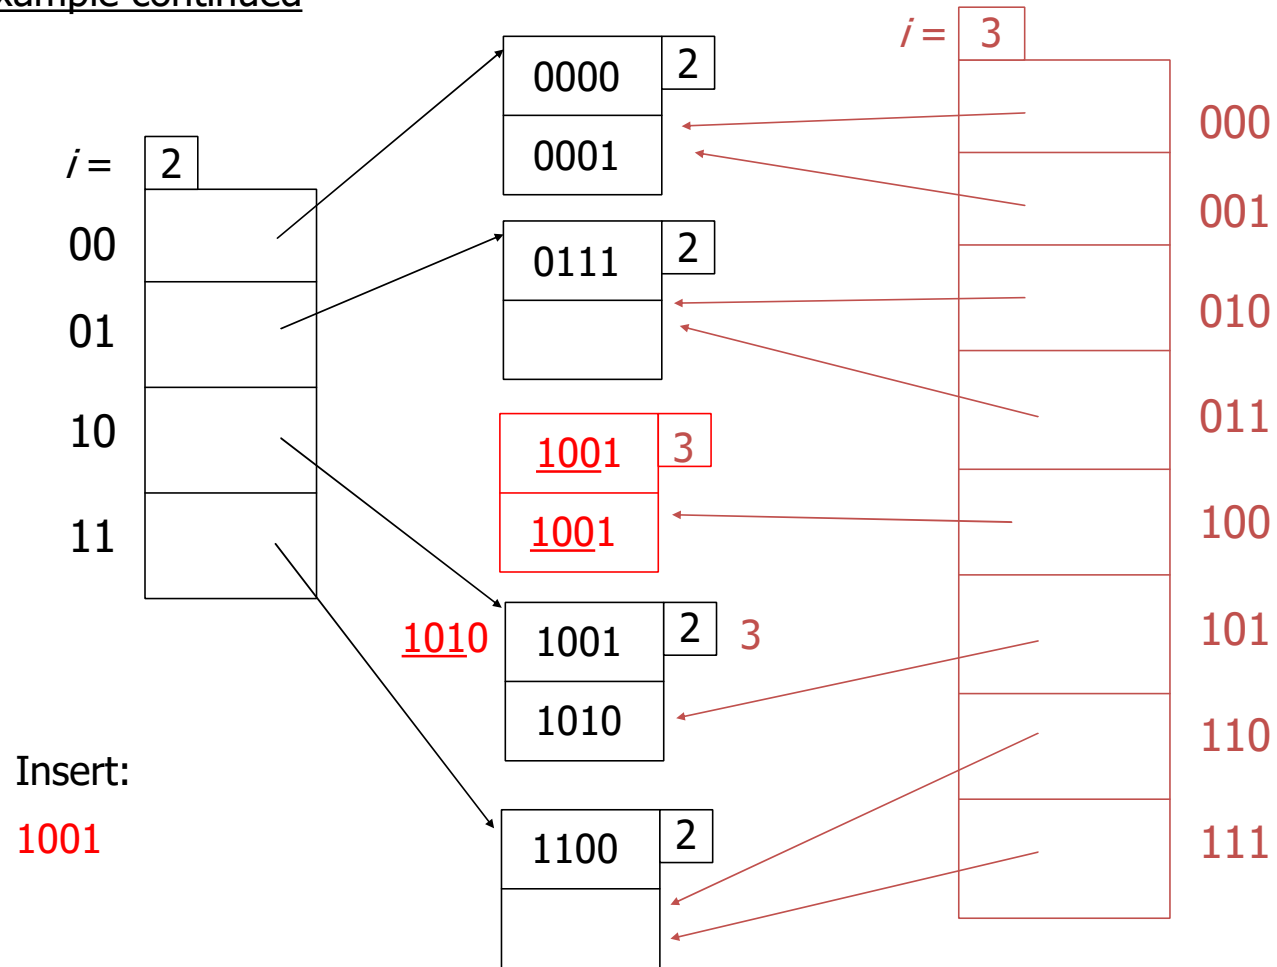
Example: $h(k)$ is 4 bits; 2 keys/block



Example continued



Example continued



When does bucket split cause directory doubling?

- Before insert, local depth of bucket = global depth
- Insert causes local depth to become > global depth
- directory is doubled by **copying it over** and **'fixing'** pointer to split image page

Comments on Extendible Hashing

- If directory fits in memory, equality search answered with one disk access (to access the bucket); else:
 - Directory grows in spurts, and, if the distribution of *hash values* is skewed, directory can grow large
 - Multiple entries with same hash value cause problems
- Delete:
 - If removal of data entry makes bucket empty, can be merged with 'split image'
 - If each directory element points to same bucket as its split image, can halve directory.

Extendible Hashing

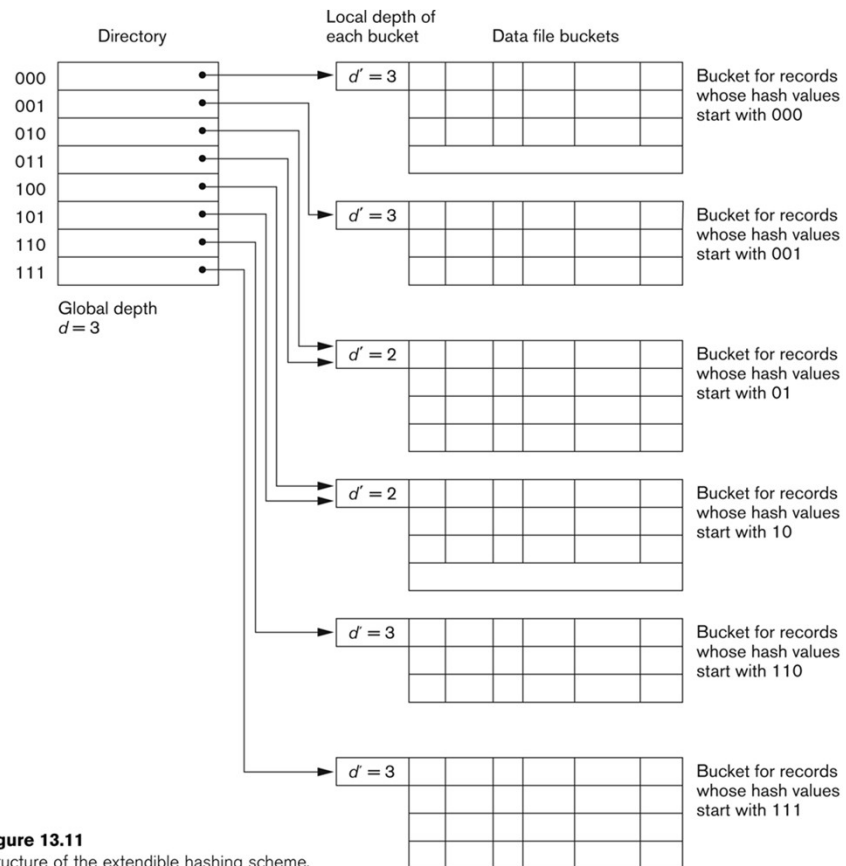


Figure 13.11
Structure of the extendible hashing scheme.

Dynamic Hashing

