

Algorithms for Query Processing and Optimization

Chapter Outline

- 0. Introduction to Query Processing
- 1. Translating SQL Queries into Relational Algebra
- 2. External Sorting
- 3. Algorithms for SELECT and JOIN Operations
- 4. Algorithms for PROJECT and SET Operations
- 5. Implementing Aggregate Operations and Outer Joins

Different File Organizations

We need to understand the importance
of appropriate file organization and index

Search key = <age, sal>

Consider following options:

- Heap files
 - random order; insert at end-of-file
- Sorted files
 - sorted on <age, sal>
- Clustered B+ tree file
 - search key <age, sal>
- Heap file with unclustered B⁺-tree index
 - on search key <age, sal>
- Heap file with unclustered hash index
 - on search key <age, sal>

Possible Operations

- Scan
 - Fetch all records from disk to buffer pool
- Equality search
 - Find all employees with age = 23 and sal = 50
 - Fetch page from disk, then locate qualifying record in page
- Range selection
 - Find all employees with age > 35
- Insert a record
 - identify the page, fetch that page from disk, inset record, write back to disk (possibly other pages as well)
- Delete a record
 - similar to insert

Understanding the Workload

- A workload is a mix of queries and updates
- For each query in the workload:
 - Which relations does it access?
 - Which attributes are retrieved?
 - Which attributes are involved in selection/join conditions? How selective are these conditions likely to be?
- For each update in the workload:
 - Which attributes are involved in selection/join conditions? How selective are these conditions likely to be?
 - The type of update (INSERT/DELETE/UPDATE), and the attributes that are affected

Choice of Indexes

- What indexes should we create?
 - Which relations should have indexes? What field(s) should be the search key? Should we build several indexes?
- For each index, what kind of an index should it be?
 - Clustered? Hash/tree?

More on Choice of Indexes

- One approach:
 - Consider the most important queries
 - Consider the best plan using the current indexes
 - See if a better plan is possible with an additional index.
 - If so, create it.
 - Obviously, this implies that we must understand how a DBMS evaluates queries and creates **query evaluation plans**
 - We will learn query execution and optimization later - For now, we discuss simple 1-table queries.
- Before creating an index, must also consider the impact on updates in the workload

Trade-offs for Indexes

- Indexes can make
 - queries go faster
 - updates slower
- Require disk space, too

Index Selection Guidelines

- Attributes in WHERE clause are candidates for index keys
 - Exact match condition suggests hash index
 - Range query suggests tree index
 - Clustering is especially useful for range queries
 - can also help on equality queries if there are many duplicates
- Try to choose indexes that benefit as many queries as possible
 - Since only one index can be clustered per relation, choose it based on important queries that would benefit the most from clustering
- Multi-attribute search keys should be considered when a WHERE clause contains several conditions
 - Order of attributes is important for range queries
- Note: clustered index should be used judiciously
 - expensive updates, although cheaper than sorted files

Examples of Clustered Indexes

- B+ tree index on E.age can be used to get qualifying tuples

What is a good indexing strategy?

- How selective is the condition?
 - everyone > 40, index not of much help, scan is as good
 - Suppose 10% > 40. Then?

```
SELECT E.dno  
FROM Emp E  
WHERE E.age>40
```

- Depends on if the index is clustered
 - otherwise can be more expensive than a linear scan
 - if clustered, 10% I/O (+ index pages)

Which attribute(s)?
Clustered/Unclustered?
B+ tree/Hash?

Examples of Clustered Indexes

Group-By query

What is a good indexing strategy?

- Use *E.age* as search key?
 - Bad If many tuples have *E.age* > 10 or if not clustered....
 - ...using *E.age* index and sorting the retrieved tuples by *E.dno* may be costly
- Clustered *E.dno* index may be better
 - First group by, then count tuples with age > 10
 - good when age > 10 is not too selective
- Note: the first option is good when the WHERE condition is highly selective (few tuples have age > 10), the second is good when not highly selective

```
SELECT E.dno, COUNT (*)  
FROM Emp E  
WHERE E.age>10  
GROUP BY E.dno
```

Which attribute(s)?
Clustered/Unclustered?
B+ tree/Hash?

Examples of Clustered Indexes

What is a good indexing strategy?

Equality queries and duplicates

- Clustering on *E.hobby* helps
 - hobby not a candidate key, several tuples possible
- Does clustering help now?
 - (eid = key)
 - Not much
 - at most one tuple satisfies the condition

```
SELECT E.dno  
FROM Emp E  
WHERE E.hobby='Stamps'
```

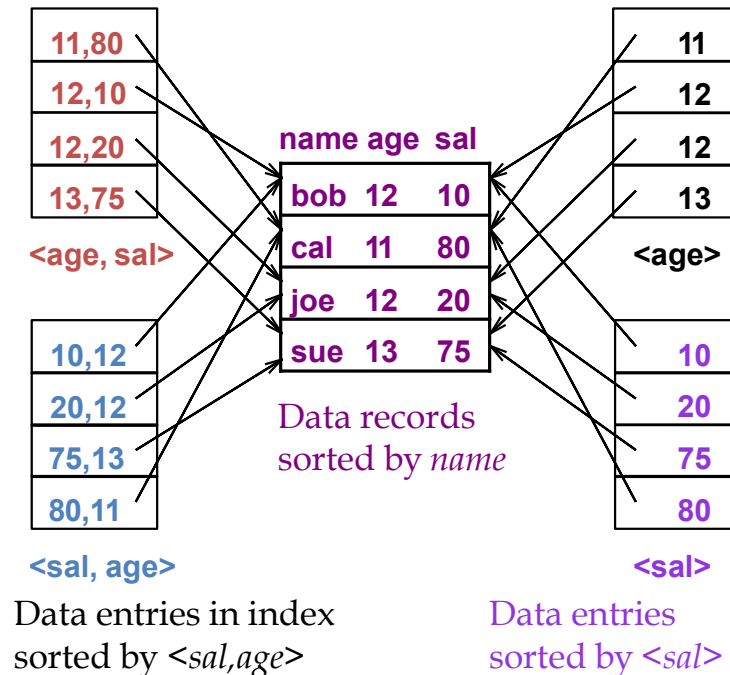
Which attribute(s)?
Clustered/Unclustered?
B+ tree/Hash?

```
SELECT E.dno  
FROM Emp E  
WHERE E.eid=50
```

Indexes with Composite Search Keys

- **Composite Search Keys:** Search on a combination of fields
- **Equality query:** Every field value is equal to a constant value. E.g. wrt $\langle \text{sal}, \text{age} \rangle$ index:
 - age=20 and sal =75
- **Range query:** Some field value is not a constant. E.g.:
 - sal > 10 – which combination(s) would help?
 - $\langle \text{age}, \text{sal} \rangle$ does not help
 - B+tree on $\langle \text{sal} \rangle$ or $\langle \text{sal}, \text{age} \rangle$ helps
 - has to be a prefix

Examples of composite key indexes using lexicographic order.



Composite Search Keys

- To retrieve Emp records with $age=30$ AND $sal=4000$, an index on $\langle age, sal \rangle$ would be better than an index on age or an index on sal
 - first find $age = 30$, among them search $sal = 4000$
- If condition is: $20 < age < 30$ AND $3000 < sal < 5000$:
 - Clustered tree index on $\langle age, sal \rangle$ or $\langle sal, age \rangle$ is best.
- If condition is: $age=30$ AND $3000 < sal < 5000$:
 - Clustered $\langle age, sal \rangle$ index much better than $\langle sal, age \rangle$ index
 - more index entries are retrieved for the latter
- Composite indexes are larger, updated more often

Index-Only Plans

- A number of queries can be answered without retrieving any tuples from one or more of the relations involved if a suitable index is available

```
SELECT E.dno, COUNT(*)  
FROM Emp E  
GROUP BY E.dno
```

<E.dno>

```
SELECT E.dno, MIN(E.sal)  
FROM Emp E  
GROUP BY E.dno
```

<E.dno,E.sal>

Tree index!

<E.age,E.sal>

Tree index!

- For index-only strategies, clustering is not important

```
SELECT AVG(E.sal)  
FROM Emp E  
WHERE E.age=25 AND  
E.sal BETWEEN 3000 AND 5000
```


External Sorting

Why Sort?

- A classic problem in computer science
- Data requested in sorted order
 - e.g., find students in increasing gpa order
- Sorting is first step in bulk loading B+ tree index
- Sorting useful for eliminating duplicate copies in a collection of records
- Sort-merge join algorithm involves sorting
- **Problem: sort 1Gb of data with 1Mb of RAM**
 - need to minimize the cost of disk access

2-Way Sort: Requires 3 Buffers

- Suppose $N = 2^k$ pages in the file
- Pass 0: Read a page, sort it, write it.
 - repeat for all 2^k pages
 - only one buffer page is used
- Pass 1:
 - Read two pages, sort (merge) them using one output page, write them to disk
 - repeat 2^{k-1} times
 - three buffer pages used
- Pass 2, 3, 4, continue

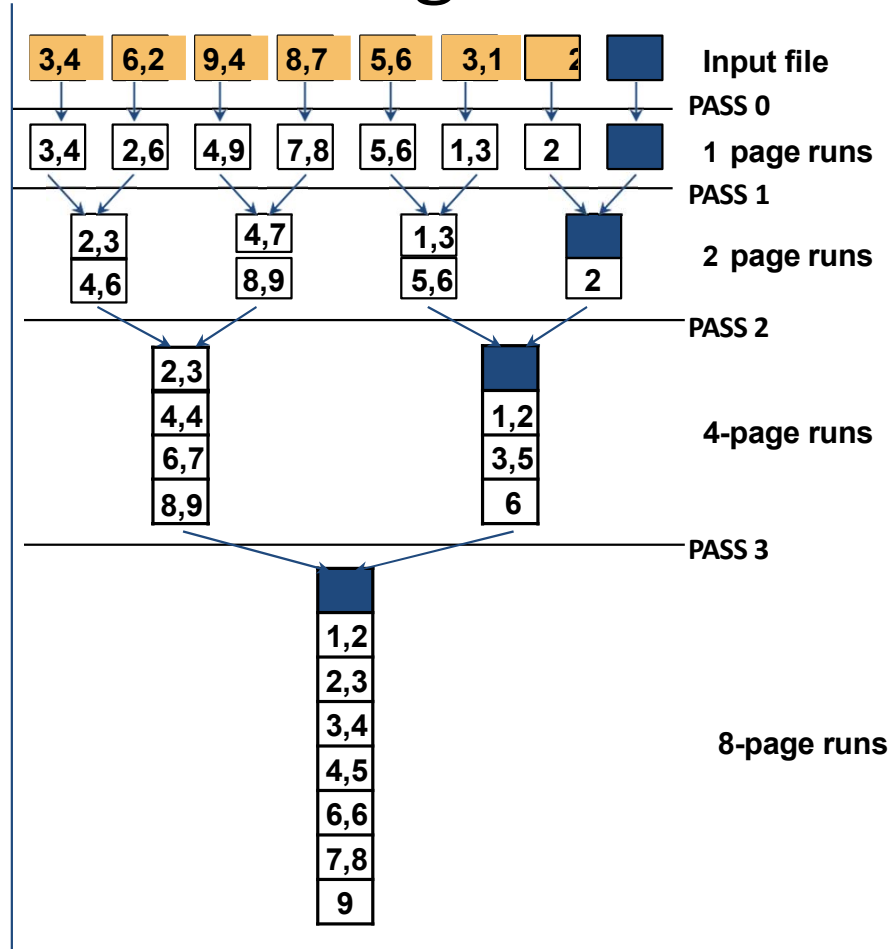


Two-Way External Merge Sort

- Each sorted sub-file is called a **run**
 - each run can contain multiple pages
- Each pass we read + write each page in file.
- N pages in the file,
- => the number of passes


$$= \lceil \log_2 N \rceil + 1$$
- So total cost is:

$$2N (\lceil \log_2 N \rceil + 1)$$
- Not too practical, but useful to learn basic concepts for external sorting



General External Merge Sort

- Suppose we have more than 3 buffer pages.
- How can we utilize them?
- To sort a file with N pages using B buffer pages:
 - Pass 0: use B buffer pages:
 - Produce $\lceil N/B \rceil$ sorted runs of B pages each.
 - Pass 1, 2, ..., etc.: merge $B-1$ runs to one output page
 - keep writing to disk once the output page is full



Cost of External Merge Sort

- Number of passes: $1 + \lceil \log_{B-1} [N/B] \rceil$
- Cost = $2N * (\text{\# of passes})$ – why 2 times?
- E.g., with 5 buffer pages, to sort 108 page file:
- Pass 0: sorting 5 pages at a time
 - $\lceil 108/5 \rceil = 22$ sorted runs of 5 pages each (last run is only 3 pages)
- Pass 1: 4-way merge
 - $\lceil 22/4 \rceil = 6$ sorted runs of 20 pages each (last run is only 8 pages)
- Pass 2: 4-way merge
 - (but 2-way for the last two runs)
 - $\lceil 6/4 \rceil = 2$ sorted runs, 80 pages and 28 pages
- Pass 3: 2-way merge (only 2 runs remaining)
 - Sorted file of 108 pages

Number of Passes of External Sort

High B is good, although CPU cost increases


N	B=3	B=5	B=9	B=17	B=129	B=257
100	7	4	3	2	1	1
1,000	10	5	4	3	2	2
10,000	13	7	5	4	2	2
100,000	17	9	6	5	3	3
1,000,000	20	10	7	5	3	3
10,000,000	23	12	8	6	4	3
100,000,000	26	14	9	7	4	4
1,000,000,000	30	15	10	8	5	4

I/O for External Merge Sort

- If 10 buffer pages
 - either merge 9 runs at a time with one output buffer
 - or 8 runs with two output buffers
- If #page I/O is the metric
 - goal is minimize the #passes
 - each page is read and written in each pass
- If we decide to read a block of b pages sequentially
 - Suggests we should make each buffer (input/output) be a block of pages
 - But this will reduce fan-out during merge passes
 - i.e. not as many runs can be merged again any more
 - In practice, most files still sorted in 2-3 passes

Double Buffering

- To reduce CPU wait time for I/O request to complete, can prefetch into 'shadow block'.



Overview of Query Evaluation

Overview of Query Evaluation

- How queries are evaluated in a DBMS
 - How DBMS describes data (tables and indexes)
- Relational Algebra Tree/Plan = Logical Query Plan
- Now Algorithms will be attached to each operator = Physical Query Plan
- Plan = Tree of RA ops, with choice of algorithm for each op.
 - Each operator typically implemented using a “pull” interface
 - when an operator is “pulled” for the next output tuples, it “pulls” on its inputs and computes them

Overview of Query Evaluation

- Two main issues in query optimization:
 1. For a given query, what plans are considered?
 - Algorithm to search plan space for cheapest (estimated) plan
 2. How is the cost of a plan estimated?
- Ideally: Want to find best plan
- Practically: Avoid worst plans!

Assumption: ignore final write

- i.e. assume that your final results can be left in memory
 - and does not be written back to disk
 - unless mentioned otherwise
- Why such an assumption?

Algorithms for Joins

Equality Joins With One Join Column

```
SELECT *  
FROM   Reserves R, Sailors S  
WHERE  R.sid=S.sid
```

- In algebra: $R \bowtie S$
 - Common! Must be carefully optimized
 - $R \times S$ is large; so, $R \times S$ followed by a selection is inefficient
- Cost metric: # of I/Os
 - Remember, we will ignore output costs (always)
 - = the cost to write the final result tuples back to the disk

Common Join Algorithms

1. Nested Loops Joins (NLJ)
 - Simple nested loop join
 - Block nested loop join
2. Sort Merge Join Very similar to external sort
3. Hash Join

Algorithms for Joins

1. NESTED LOOP JOINS

Simple Nested Loops Join

$R \bowtie S$

```
foreach tuple r in R do
    foreach tuple s in S where  $r_i == s_j$  do
        add  $\langle r, s \rangle$  to result
```

$M = 1000$ pages in R
 $p_R = 100$ tuples per page

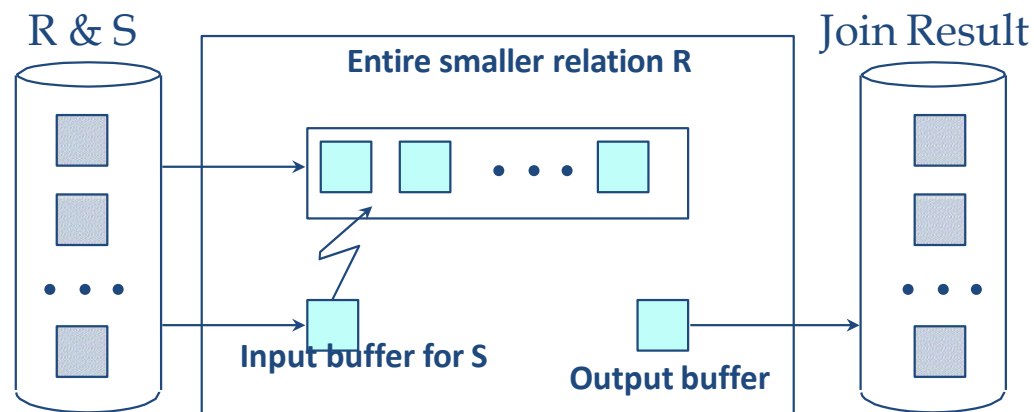
$N = 500$ pages in S
 $p_S = 80$ tuples per page

- For each tuple in the **outer** relation R , we scan the entire **inner** relation S .
 - Cost: $M + (p_R * M) * N = 1000 + 100 * 1000 * 500$ I/Os.
- **Page-oriented Nested Loops join:**
 - For each *page* of R , get each *page* of S
 - and write out matching pairs of tuples $\langle r, s \rangle$
 - where r is in R -page and S is in S -page.
 - Cost: $M + M * N = 1000 + 1000 * 500$
- If smaller relation (S) is outer
 - Cost: $N + M * N = 500 + 500 * 1000$

How many buffer pages
do you need?

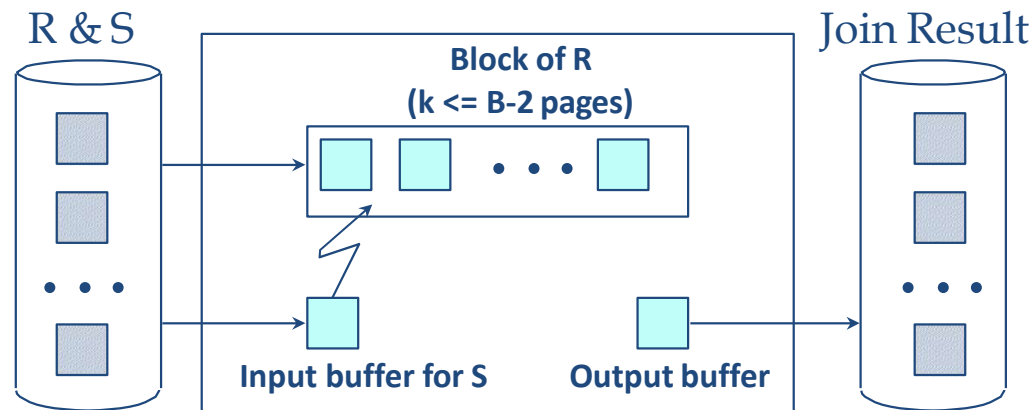
Block Nested Loops Join

- Simple-Nested does not properly utilize buffer pages (uses 3 pages)
- Suppose have enough memory to hold **the smaller relation R + at least two other pages**
 - e.g. in the example on previous slide (S is smaller), and we need $500 + 2 = 502$ pages in the buffer
- Then use one page as an input buffer for scanning the inner
 - one page as the output buffer
 - For each matching tuple r in R-block, s in S-page, add $\langle r, s \rangle$ to result
- Total I/O = $M+N$
- What if the entire smaller relation does not fit?



Block Nested Loops Join

- If R does not fit in memory,
 - Use one page as an input buffer for scanning the inner S
 - one page as the output buffer
 - and use all remaining pages to hold ``block'' of outer R.
 - For each matching tuple r in R-block, s in S-page, add $\langle r, s \rangle$ to result
 - Then read next R-block, scan S, etc.



Cost of Block Nested Loops

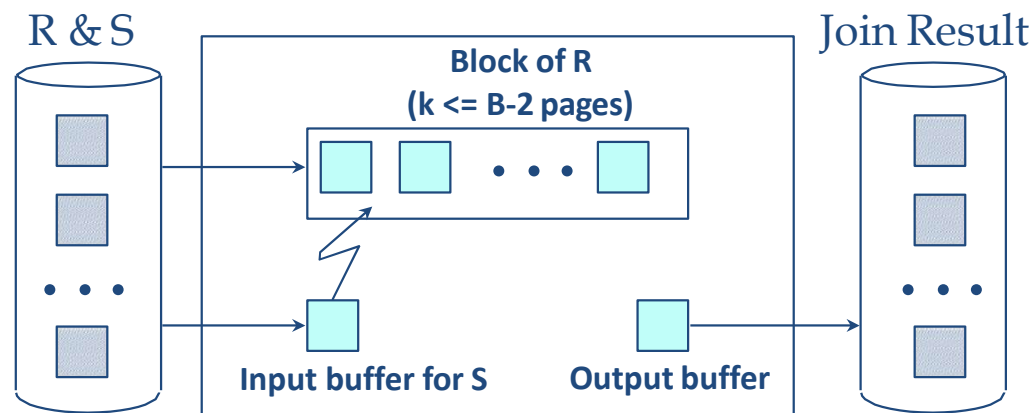
in class

- R is outer
- B-2 = 100-page blocks
- How many blocks of R?
- Cost to scan R?
- Cost to scan S?
- Total Cost?

M = 1000 pages in R
 $p_R = 100$ tuples per page

N = 500 pages in S
 $p_S = 80$ tuples per page

```
foreach block of B-2 pages of R do
  foreach page of S do {
    for all matching in-memory tuples r in R-
      block and s in S-page
      add <r, s> to result
```



Cost of Block Nested Loops

- R is outer
- B-2 = 100-page blocks
- How many blocks of R? 10
- Cost to scan R? 1000
- Cost to scan S? $10 * 500$
- Total Cost? $1000 + 5000 = 6000$
- (check yourself)
 - If space for just 90 pages of R, we would scan S 12 times, cost = 7000

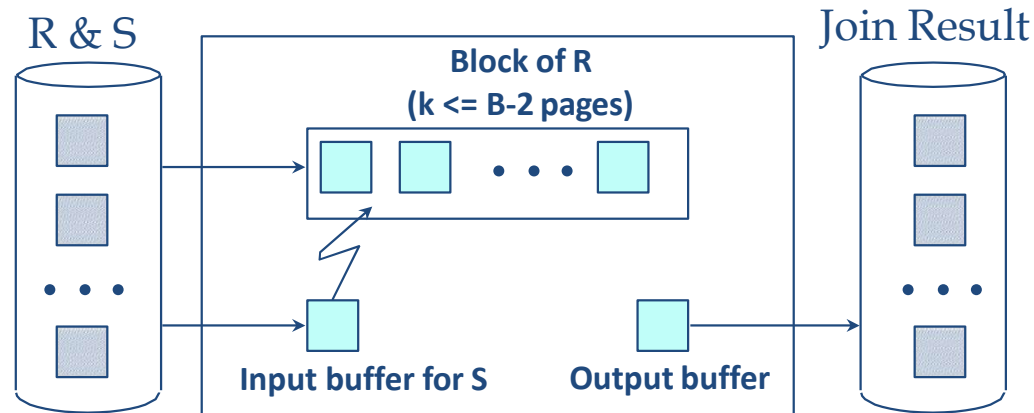
M = 1000 pages in R
 $p_R = 100$ tuples per page

N = 500 pages in S
 $p_S = 80$ tuples per page

```
foreach block of B-2 pages of R do
    foreach page of S do {
        for all matching in-memory tuples r in R-
        block and s in S-page
            add <r, s> to result
```

- Cost: Scan of outer + #outer blocks * scan of inner
 - #outer blocks = $\lceil \text{\#pages of outer relation} / \text{blocksize} \rceil$

for blocked access, it might be good to equally divide buffer pages among R and S ("seek time" less)



Algorithms for Joins

2. SORT-MERGE JOINS


Sort-Merge Join

- Sort R and S on the join column
- Then scan them to do a ``merge'' (on join col.)
- Output result tuples.

Sort-Merge Join: 1/3

- Advance scan of R until current R-tuple \geq current S tuple
 - then advance scan of S until current S-tuple \geq current R tuple
 - do this as long as current R tuple = current S tuple


Sailors



S

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

Reserves



R

<u>sid</u>	<u>bid</u>	<u>day</u>	rname
28	103	12/4/96	guppy
28	103	11/3/96	yuppy
31	101	10/10/96	dustin
31	102	10/12/96	lubber
31	101	10/11/96	lubber
58	103	11/12/96	dustin

Sort-Merge Join: 2/3

- At this point, all R tuples with same value in R_i (*current R group*) and all S tuples with same value in S_j (*current S group*)
 - match
 - find all the equal tuples
 - output $\langle r, s \rangle$ for all pairs of such tuples

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

S →

<u>sid</u>	<u>bid</u>	<u>day</u>	rname
28	103	12/4/96	guppy
28	103	11/3/96	yuppy
31	101	10/10/96	dustin
31	102	10/12/96	lubber
31	101	10/11/96	lubber
58	103	11/12/96	dustin

→ R

WRITE TWO OUTPUT TUPLES

Sort-Merge Join: 3/3

- Then resume scanning R and S

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

S →

WRITE THREE OUTPUT TUPLES

<u>sid</u>	<u>bid</u>	<u>day</u>	rname
28	103	12/4/96	guppy
28	103	11/3/96	yuppy
31	101	10/10/96	dustin
31	102	10/12/96	lubber
31	101	10/11/96	lubber
58	103	11/12/96	dustin

R →

Sort-Merge Join: 3/3

- ... and proceed till end

S	<u>sid</u>	sname	rating	age
	22	dustin	7	45.0
	28	yuppy	9	35.0
	31	lubber	8	55.5
→	44	guppy	5	35.0
	58	rusty	10	35.0

NO MATCH, CONTINUE SCANNING S

R	<u>sid</u>	<u>bid</u>	<u>day</u>	rname
	28	103	12/4/96	guppy
	28	103	11/3/96	yuppy
	31	101	10/10/96	dustin
	31	102	10/12/96	lubber
	31	101	10/11/96	lubber
→	58	103	11/12/96	dustin

Sort-Merge Join: 3/3

- ... and proceed till end

S	<u>sid</u>	sname	rating	age
	22	dustin	7	45.0
	28	yuppy	9	35.0
	31	lubber	8	55.5
	44	guppy	5	35.0
	58	rusty	10	35.0

→

WRITE ONE OUTPUT TUPLE

R	<u>sid</u>	<u>bid</u>	<u>day</u>	rname
	28	103	12/4/96	guppy
	28	103	11/3/96	yuppy
	31	101	10/10/96	dustin
	31	102	10/12/96	lubber
	31	101	10/11/96	lubber
	58	103	11/12/96	dustin

→

Example of Sort-Merge Join

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

<u>sid</u>	<u>bid</u>	<u>day</u>	rname
28	103	12/4/96	guppy
28	103	11/3/96	yuppy
31	101	10/10/96	dustin
31	102	10/12/96	lubber
31	101	10/11/96	lubber
58	103	11/12/96	dustin

- Typical Cost: $O(M \log M) + O(N \log N) + (M+N)$
 - ignoring B (as the base of log)
 - cost of sorting R + sorting S + merging R, S
 - The cost of scanning in merge-sort, $M+N$, could be $M*N$!
 - assume the same single value of join attribute in both R and S
 - but it is extremely unlikely

Cost of Sort-Merge Join

$M = 1000$ pages in R
 $p_R = 100$ tuples per page

$N = 500$ pages in S
 $p_S = 80$ tuples per page

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

<u>sid</u>	<u>bid</u>	<u>day</u>	rname
28	103	12/4/96	guppy
28	103	11/3/96	yuppy
31	101	10/10/96	dustin
31	102	10/12/96	lubber
31	101	10/11/96	lubber
58	103	11/12/96	dustin

- 100 buffer pages
- Sort R:
 - (pass 0) $1000/100 = 10$ sorted runs
 - (pass 1) merge 10 runs
 - read + write, 2 passes
 - $4 * 1000 = 4000$ I/O
- Similarly, Sort S: $4 * 500 = 2000$ I/O
- Second merge phase of sort-merge join
 - another $1000 + 500 = 1500$ I/O
 - assume uniform ~ 2.5 matches per sid, so $M+N$ is sufficient
- Total 7500 I/O

- Check yourself:
 - Consider #buffer pages 35, 100, 300
 - Cost of sort-merge = 7500 in all three
 - Cost of block nested 16500, 6000, 2500

Algorithms for Joins

3. HASH JOINS

Two Phases

1. Partition Phase


- partition R and S using the same hash function **h**

2. Probing Phase


- join tuples from the same partition (same $h(..)$ value) of R and S
- tuples in different partition of h will never join
- use a “different” hash function **h2** for joining these tuples
 - (why different – see next slide first)

Hash-Join

- Partition both relations using hash function h
- R tuples in partition i will only match S tuples in partition i



- ❖ Read in a partition of R, hash it using $h_2 (\neq h)$.
- ❖ Scan matching partition of S, search for matches.



Cost of Hash-Join

- In partitioning phase
 - read+write both relns; $2(M+N)$
 - In matching phase, read both relns; $M+N$ I/Os
 - remember – we are not counting final write
- In our running example, this is a total of 4500 I/Os
 - $3 * (1000 + 500)$
 - Compare with the previous joins

Sort-Merge Join vs. Hash Join

- Both can have a cost of $3(M+N)$ I/Os
 - if sort-merge gets enough buffer
- Hash join holds smaller relation in buffer-better if limited buffer
- Hash Join shown to be highly parallelizable
- Sort-Merge less sensitive to data skew
 - also result is sorted

Other operator algorithms

Algorithms for Selection

```
SELECT *  
FROM   Reserves R  
WHERE  R.rname = 'Joe'
```

- No index, unsorted data
 - Scan entire relation
 - May be expensive if not many 'Joe's
- No index, sorted data (on 'rname')
 - locate the first tuple, scan all matching tuples
 - first binary search, then scan depends on matches
- B+-tree index, Hash index
 - Discussed earlier
 - Cost of accessing data entries + matching data records
 - Depends on clustered/unclustered
- More complex condition like `day<8/9/94 AND bid=5 AND sid=3`
 - Either use one index, then filter
 - Or use two indexes, then take intersection, then apply third condition
 - etc.

Algorithms for Projection

```
SELECT DISTINCT  
      R.sid, R.bid  
FROM   Reserves R
```

- Two parts
 - Remove fields: **easy**
 - Remove duplicates (if distinct is specified): **expensive**
- Sorting-based
 - Sort, then scan adjacent tuples to remove duplicates
 - Can eliminate unwanted attributes in the first pass of merge sort
- Hash-based
 - Exactly like hash join
 - Partition only one relation in the first pass
 - Remove duplicates in the second pass
- Sort vs Hash
 - Sorting handles skew better, returns results sorted
 - Hash table may not fit in memory – sorting is more standard
- Index-only scan may work too
 - If all required attributes are part of index

Algorithms for Set Operations

- Intersection, cross product are special cases of joins
- Union, Except
 - Sort-based
 - Hash-based
 - Very similar to joins and projection

Algorithms for Aggregate Operations

- SUM, AVG, MIN etc.
 - again similar to previous approaches
- Without grouping:
 - In general, requires scanning the relation.
 - Given index whose search key includes all attributes in the SELECT or WHERE clauses, can do index-only scan
- With grouping:
 - Sort on group-by attributes
 - or, hash on group-by attributes
 - can combine sort/hash and aggregate
 - can do index-only scan here as well