

# Algorithms for Query Processing and Optimization

# Chapter Outline

- 0. Introduction to Query Processing
- 1. Translating SQL Queries into Relational Algebra
- 2. External Sorting
- 3. Algorithms for SELECT and JOIN Operations
- 4. Algorithms for PROJECT and SET Operations
- 5. Implementing Aggregate Operations and Outer Joins

# Different File Organizations

We need to understand the importance of appropriate file organization and index

Search key = <age, sal>

Consider following options:

- Heap files
  - random order; insert at end-of-file
- Sorted files
  - sorted on <age, sal>
- Clustered B+ tree file
  - search key <age, sal>
- Heap file with unclustered B+-tree index
  - on search key <age, sal>
- Heap file with unclustered hash index
  - on search key <age, sal>

# Possible Operations

- Scan
  - Fetch all records from disk to buffer pool
- Equality search
  - Find all employees with age = 23 and sal = 50
  - Fetch page from disk, then locate qualifying record in page
- Range selection
  - Find all employees with age > 35
- Insert a record
  - identify the page, fetch that page from disk, inset record, write back to disk (possibly other pages as well)
- Delete a record
  - similar to insert

# Understanding the Workload

- A workload is a mix of queries and updates
- For each query in the workload:
  - Which relations does it access?
  - Which attributes are retrieved?
  - Which attributes are involved in selection/join conditions? How selective are these conditions likely to be?
- For each update in the workload:
  - Which attributes are involved in selection/join conditions? How selective are these conditions likely to be?
  - The type of update (INSERT/DELETE/UPDATE), and the attributes that are affected

# Choice of Indexes

- What indexes should we create?
  - Which relations should have indexes? What field(s) should be the search key? Should we build several indexes?
- For each index, what kind of an index should it be?
  - Clustered? Hash/tree?

# More on Choice of Indexes

- One approach:
  - Consider the most important queries
  - Consider the best plan using the current indexes
  - See if a better plan is possible with an additional index.
  - If so, create it.
  - Obviously, this implies that we must understand how a DBMS evaluates queries and creates query evaluation plans
  - We will learn query execution and optimization later - For now, we discuss simple 1-table queries.
- Before creating an index, must also consider the impact on updates in the workload

# Trade-offs for Indexes

- Indexes can make
  - queries go faster
  - updates slower
- Require disk space, too



# Index Selection Guidelines

- Attributes in WHERE clause are candidates for index keys
  - Exact match condition suggests hash index
  - Range query suggests tree index
  - Clustering is especially useful for range queries
    - can also help on equality queries if there are many duplicates
- Try to choose indexes that benefit as many queries as possible
  - Since only one index can be clustered per relation, choose it based on important queries that would benefit the most from clustering
- Multi-attribute search keys should be considered when a WHERE clause contains several conditions
  - Order of attributes is important for range queries
- Note: clustered index should be used judiciously
  - expensive updates, although cheaper than sorted files

# Examples of Clustered Indexes

- B+ tree index on E.age can be used to get qualifying tuples

What is a good indexing strategy?

- How selective is the condition?
  - everyone > 40, index not of much help, scan is as good
  - Suppose 10% > 40. Then?

```
SELECT E.dno  
FROM Emp E  
WHERE E.age>40
```

- Depends on if the index is clustered
  - otherwise can be more expensive than a linear scan
  - if clustered, 10% I/O (+ index pages)

Which attribute(s)?  
Clustered/Unclustered?  
B+ tree/Hash?

# Examples of Clustered Indexes

## Group-By query

What is a good indexing strategy?

- Use *E.age* as search key?
  - Bad If many tuples have *E.age* > 10 or if not clustered....
  - ...using *E.age* index and sorting the retrieved tuples by *E.dno* may be costly
- Clustered *E.dno* index may be better
  - First group by, then count tuples with age > 10
  - good when age > 10 is not too selective
- Note: the first option is good when the WHERE condition is highly selective (few tuples have age > 10), the second is good when not highly selective

```
SELECT E.dno, COUNT (*)  
FROM Emp E  
WHERE E.age>10  
GROUP BY E.dno
```

Which attribute(s)?  
Clustered/Unclustered?  
B+ tree/Hash?

# Examples of Clustered Indexes

What is a good indexing strategy?

## Equality queries and duplicates

- Clustering on *E.hobby* helps
  - hobby not a candidate key, several tuples possible
- Does clustering help now?
  - (eid = key)
  - Not much
  - at most one tuple satisfies the condition

```
SELECT E.dno  
FROM Emp E  
WHERE E.hobby='Stamps'
```

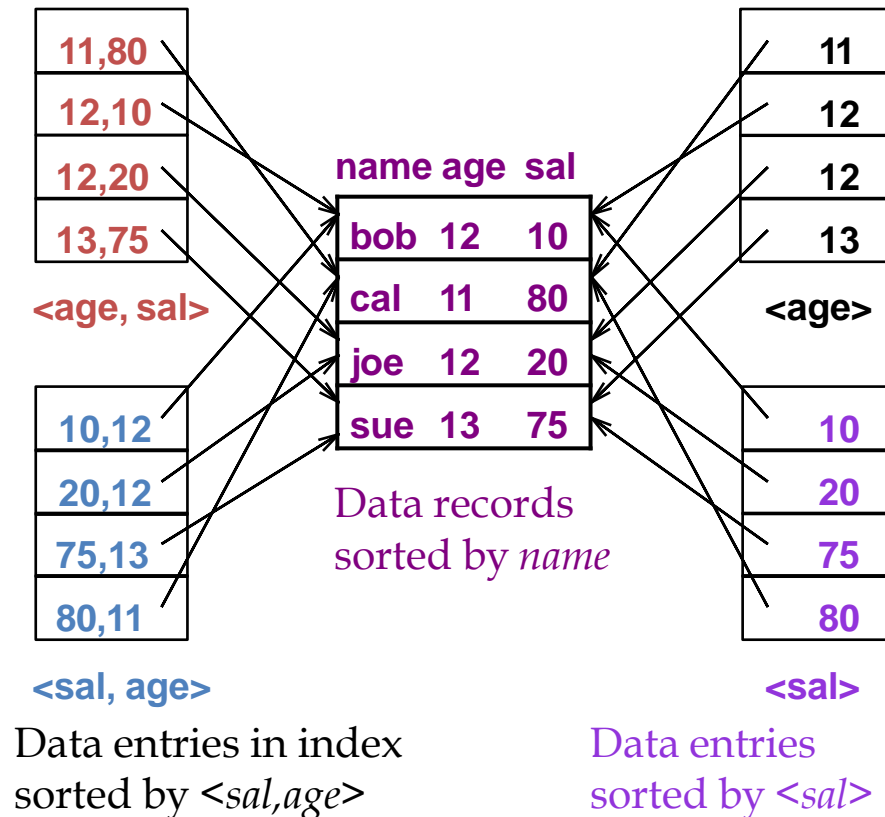
Which attribute(s)?  
Clustered/Unclustered?  
B+ tree/Hash?

```
SELECT E.dno  
FROM Emp E  
WHERE E.eid=50
```

# Indexes with Composite Search Keys

- **Composite Search Keys:** Search on a combination of fields
- **Equality query:** Every field value is equal to a constant value. E.g. wrt  $\langle \text{sal}, \text{age} \rangle$  index:
  - age=20 and sal =75
- **Range query:** Some field value is not a constant. E.g.:
  - sal > 10 – which combination(s) would help?
  - $\langle \text{age}, \text{sal} \rangle$  does not help
  - B+tree on  $\langle \text{sal} \rangle$  or  $\langle \text{sal}, \text{age} \rangle$  helps
  - has to be a prefix

Examples of composite key indexes using lexicographic order.



# Composite Search Keys

- To retrieve Emp records with  $age=30$  AND  $sal=4000$ , an index on  $\langle age, sal \rangle$  would be better than an index on  $age$  or an index on  $sal$ 
  - first find  $age = 30$ , among them search  $sal = 4000$
- If condition is:  $20 < age < 30$  AND  $3000 < sal < 5000$ :
  - Clustered tree index on  $\langle age, sal \rangle$  or  $\langle sal, age \rangle$  is best.
- If condition is:  $age=30$  AND  $3000 < sal < 5000$ :
  - Clustered  $\langle age, sal \rangle$  index much better than  $\langle sal, age \rangle$  index
  - more index entries are retrieved for the latter
- Composite indexes are larger, updated more often

# Index-Only Plans

- A number of queries can be answered without retrieving any tuples from one or more of the relations involved if a suitable index is available

```
SELECT E.dno, COUNT(*)  
FROM Emp E  
GROUP BY E.dno
```

*<E.dno>*

```
SELECT E.dno, MIN(E.sal)  
FROM Emp E  
GROUP BY E.dno
```

*<E.dno, E.sal>*

*Tree index!*

*<E.age, E.sal>*

*Tree index!*

- For index-only strategies, clustering is not important

```
SELECT AVG(E.sal)  
FROM Emp E  
WHERE E.age=25 AND  
E.sal BETWEEN 3000 AND 5000
```

# External Sorting




# Why Sort?

- A classic problem in computer science
- Data requested in sorted order
  - e.g., find students in increasing gpa order
- Sorting is first step in bulk loading B+ tree index
- Sorting useful for eliminating duplicate copies in a collection of records
- Sort-merge join algorithm involves sorting
- **Problem: sort 1Gb of data with 1Mb of RAM**
  - need to minimize the cost of disk access

# 2-Way Sort: Requires 3 Buffers

- Suppose  $N = 2^k$  pages in the file
- Pass 0: Read a page, sort it, write it.
  - repeat for all  $2^k$  pages
  - only one buffer page is used
- Pass 1:
  - Read two pages, sort (merge) them using one output page, write them to disk
  - repeat  $2^{k-1}$  times
  - three buffer pages used
- Pass 2, 3, 4, ..... continue

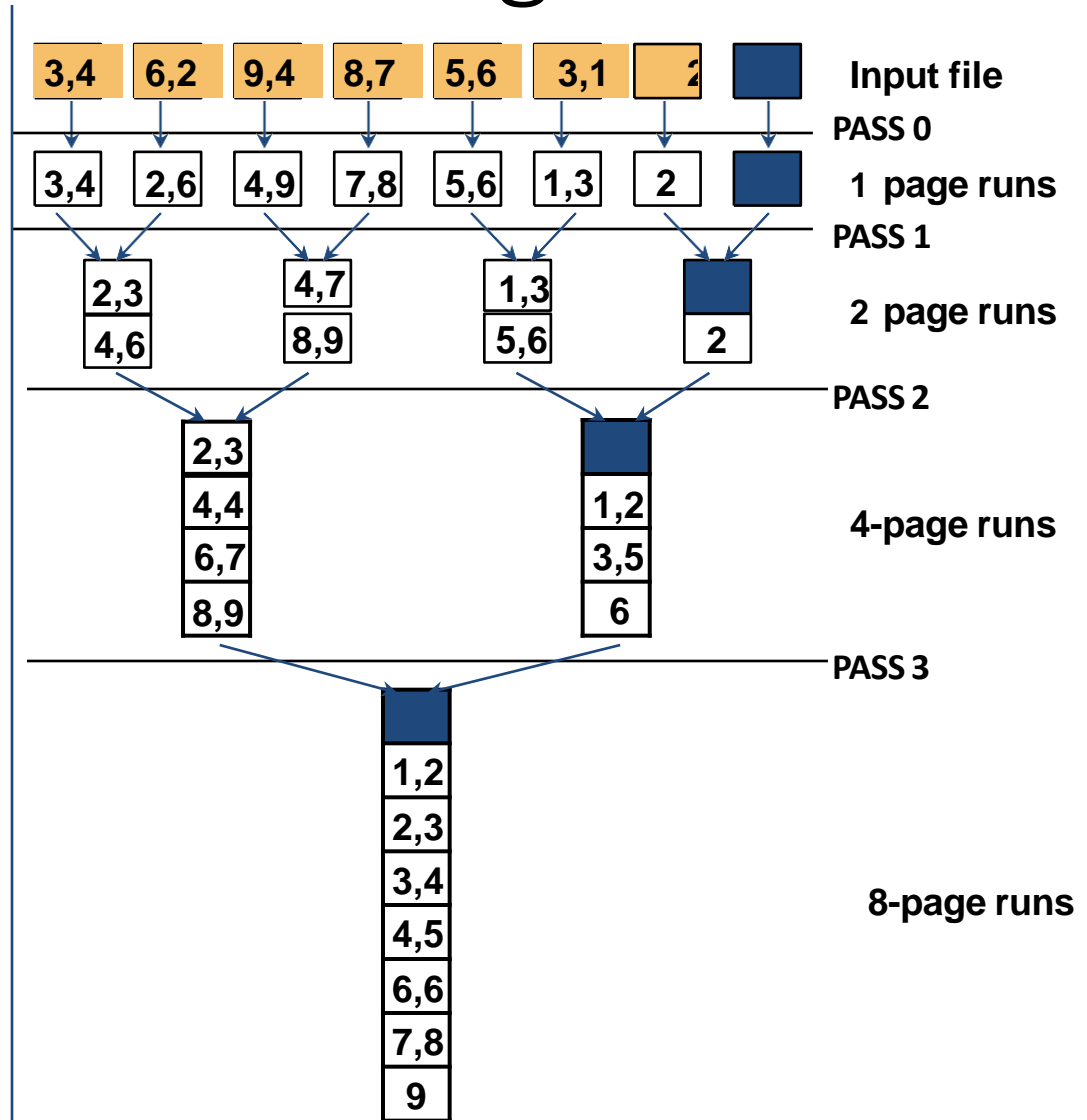


# Two-Way External Merge Sort

- Each sorted sub-file is called a **run**
  - each run can contain multiple pages
- Each pass we read + write each page in file.
- N pages in the file,
- => the number of passes  

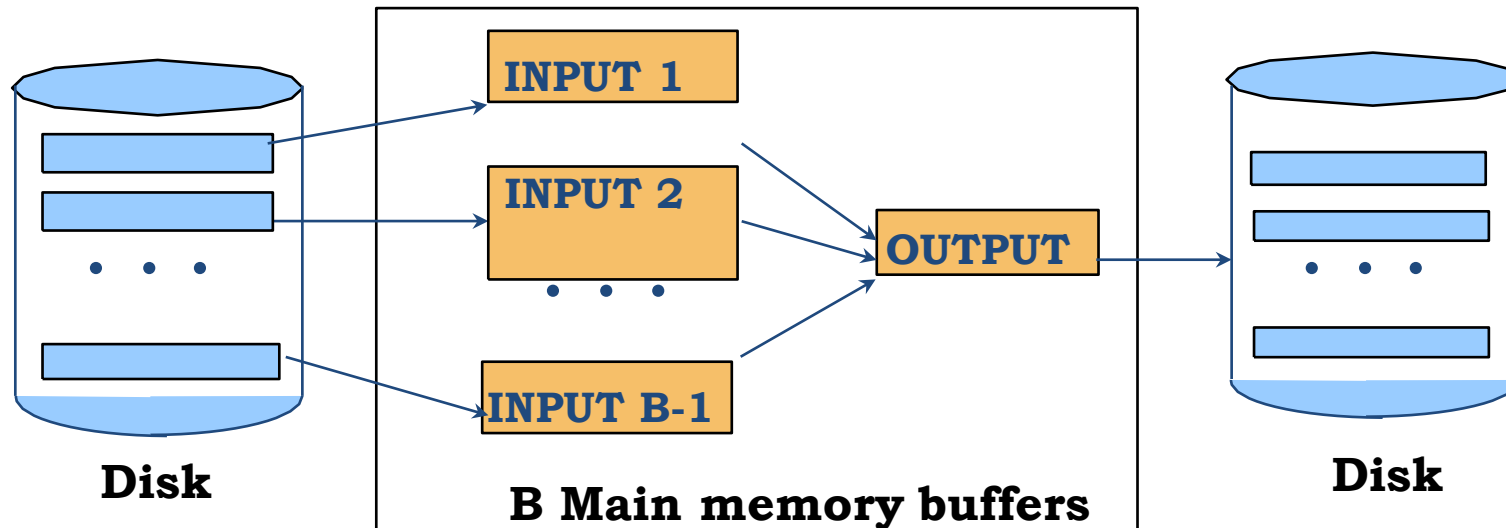
$$= \lceil \log_2 N \rceil + 1$$
- So total cost is:  

$$2N (\lceil \log_2 N \rceil + 1)$$
- Not too practical, but useful to learn basic concepts for external sorting



# General External Merge Sort

- Suppose we have more than 3 buffer pages.
- How can we utilize them?
- To sort a file with  $N$  pages using  $B$  buffer pages:
  - Pass 0: use  $B$  buffer pages:
    - Produce  $\lceil N/B \rceil$  sorted runs of  $B$  pages each.
  - Pass 1, 2, ..., etc.: merge  $B-1$  runs to one output page
    - keep writing to disk once the output page is full



# Cost of External Merge Sort

- Number of passes:  $1 + \lceil \log_{B-1} [N/B] \rceil$
- Cost =  $2N * (\text{\# of passes})$  – why 2 times?
- E.g., with 5 buffer pages, to sort 108 page file:
- Pass 0: sorting 5 pages at a time
  - $\lceil 108/5 \rceil = 22$  sorted runs of 5 pages each (last run is only 3 pages)
- Pass 1: 4-way merge
  - $\lceil 22/4 \rceil = 6$  sorted runs of 20 pages each (last run is only 8 pages)
- Pass 2: 4-way merge
  - (but 2-way for the last two runs)
  - $\lceil 6/4 \rceil = 2$  sorted runs, 80 pages and 28 pages
- Pass 3: 2-way merge (only 2 runs remaining)
  - Sorted file of 108 pages

# Number of Passes of External Sort

High B is good, although CPU cost increases

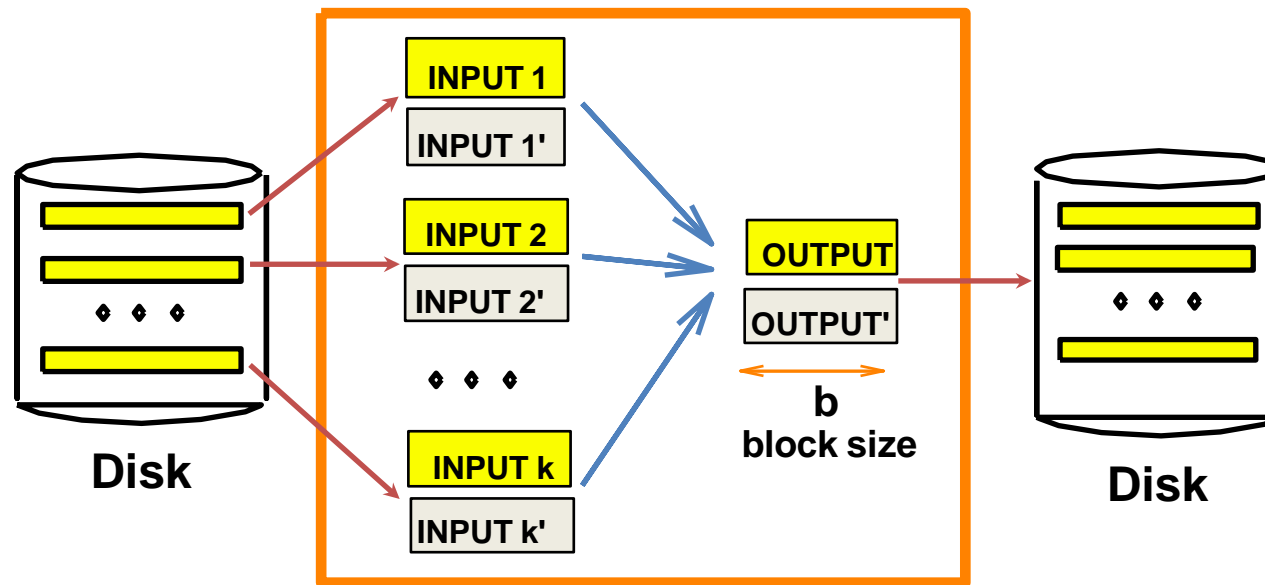
N	B=3	B=5	B=9	B=17	B=129	B=257
100	7	4	3	2	1	1
1,000	10	5	4	3	2	2
10,000	13	7	5	4	2	2
100,000	17	9	6	5	3	3
1,000,000	20	10	7	5	3	3
10,000,000	23	12	8	6	4	3
100,000,000	26	14	9	7	4	4
1,000,000,000	30	15	10	8	5	4

# I/O for External Merge Sort

- If 10 buffer pages
  - either merge 9 runs at a time with one output buffer
  - or 8 runs with two output buffers
- If #page I/O is the metric
  - goal is minimize the #passes
  - each page is read and written in each pass
- If we decide to read a block of  $b$  pages sequentially
  - Suggests we should make each buffer (input/output) be a block of pages
  - But this will reduce fan-out during merge passes
    - i.e. not as many runs can be merged again any more
  - In practice, most files still sorted in 2-3 passes

# Double Buffering

- To reduce CPU wait time for I/O request to complete, can prefetch into 'shadow block'.



B main memory buffers, k-way merge



# Overview of Query Evaluation

# Overview of Query Evaluation

- How queries are evaluated in a DBMS
  - How DBMS describes data (tables and indexes)
- Relational Algebra Tree/Plan = Logical Query Plan
- Now Algorithms will be attached to each operator = Physical Query Plan
- Plan = Tree of RA ops, with choice of algorithm for each op.
  - Each operator typically implemented using a “pull” interface
  - when an operator is “pulled” for the next output tuples, it “pulls” on its inputs and computes them

# Overview of Query Evaluation

- Two main issues in query optimization:
  1. For a given query, what plans are considered?
    - Algorithm to search plan space for cheapest (estimated) plan
  2. How is the cost of a plan estimated?
- Ideally: Want to find best plan
- Practically: Avoid worst plans!

# Assumption: ignore final write

- i.e. assume that your final results can be left in memory
  - and does not be written back to disk
  - unless mentioned otherwise
- Why such an assumption?

# Algorithms for Joins

# Equality Joins With One Join Column

```
SELECT *  
FROM   Reserves R, Sailors S  
WHERE  R.sid=S.sid
```

- In algebra:  $R \bowtie S$ 
  - Common! Must be carefully optimized
  - $R \times S$  is large; so,  $R \times S$  followed by a selection is inefficient
- Cost metric: # of I/Os
  - Remember, we will ignore output costs (always)  
= the cost to write the final result tuples back to the disk

# Common Join Algorithms

## 1. Nested Loops Joins (NLJ)

- Simple nested loop join
- Block nested loop join

## 2. Sort Merge Join Very similar to external sort

## 3. Hash Join

# Algorithms for Joins

## 1. NESTED LOOP JOINS



# Simple Nested Loops Join

$R \bowtie S$

```
foreach tuple r in R do
    foreach tuple s in S where  $r_i == s_j$  do
        add  $\langle r, s \rangle$  to result
```

$M = 1000$  pages in R  
 $p_R = 100$  tuples per page

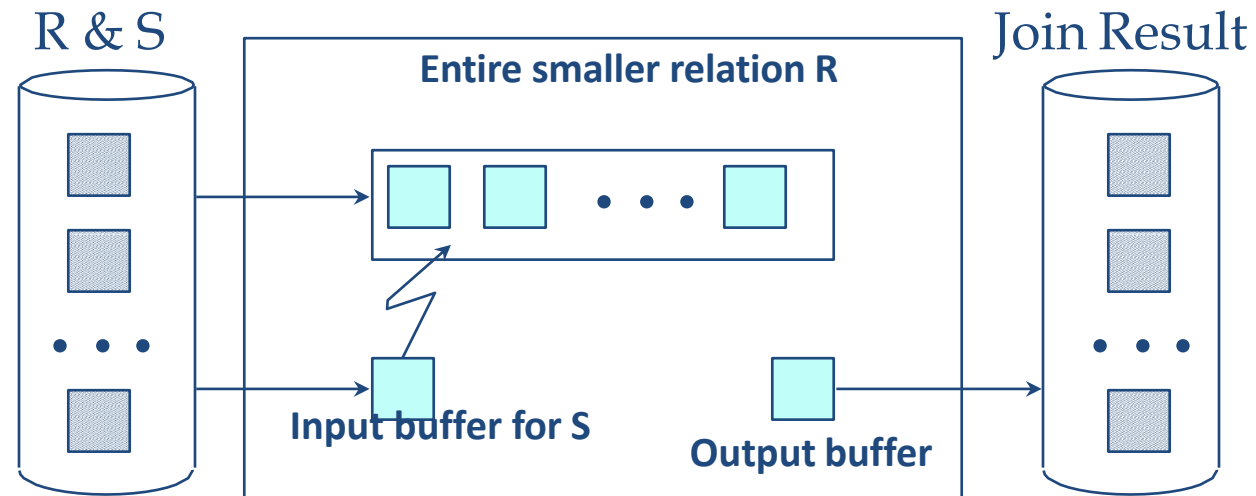
$N = 500$  pages in S  
 $p_S = 80$  tuples per page

- For each tuple in the **outer** relation R, we scan the entire **inner** relation S.
  - Cost:  $M + (p_R * M) * N = 1000 + 100 * 1000 * 500$  I/Os.
- **Page-oriented Nested Loops join:**
  - For each *page* of R, get each *page* of S
  - and write out matching pairs of tuples  $\langle r, s \rangle$
  - where r is in R-page and S is in S-page.
  - Cost:  $M + M * N = 1000 + 1000 * 500$
- If smaller relation (S) is outer
  - Cost:  $N + M * N = 500 + 500 * 1000$

How many buffer pages  
do you need?

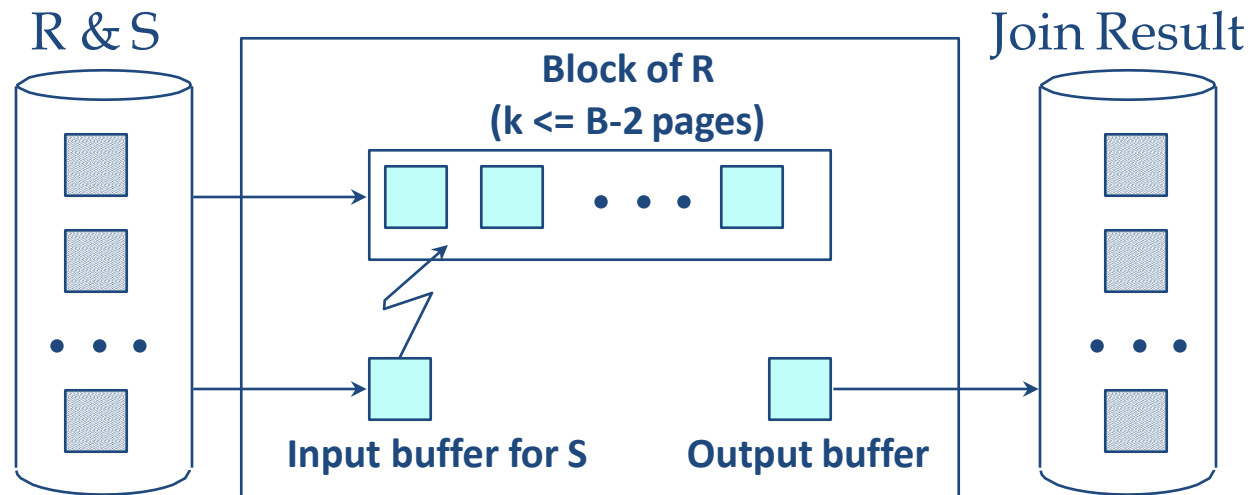
# Block Nested Loops Join

- Simple-Nested does not properly utilize buffer pages (uses 3 pages)
- Suppose have enough memory to hold **the smaller relation R + at least two other pages**
  - e.g. in the example on previous slide (S is smaller), and we need  $500 + 2 = 502$  pages in the buffer
- Then use one page as an input buffer for scanning the inner
  - one page as the output buffer
  - For each matching tuple  $r$  in R-block,  $s$  in S-page, add  $\langle r, s \rangle$  to result
- Total I/O =  $M+N$
- What if the entire smaller relation does not fit?



# Block Nested Loops Join

- If R does not fit in memory,
  - Use one page as an input buffer for scanning the inner S
  - one page as the output buffer
  - and use all remaining pages to hold ``block'' of outer R.
  - For each matching tuple r in R-block, s in S-page, add  $\langle r, s \rangle$  to result
  - Then read next R-block, scan S, etc.



# Cost of Block Nested Loops

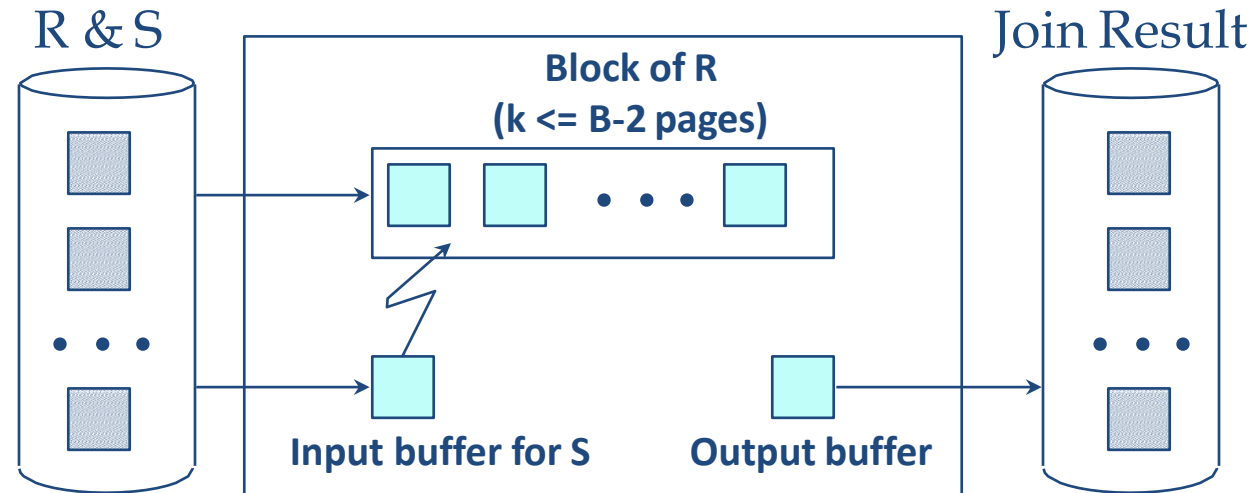
$M = 1000$  pages in  $R$   
 $p_R = 100$  tuples per page

$N = 500$  pages in  $S$   
 $p_S = 80$  tuples per page

in class

- $R$  is outer
- $B-2 = 100$ -page blocks
- How many blocks of  $R$ ?
- Cost to scan  $R$ ?
- Cost to scan  $S$ ?
- Total Cost?

```
foreach block of  $B-2$  pages of  $R$  do
  foreach page of  $S$  do {
    for all matching in-memory tuples  $r$  in  $R$ -
      block and  $s$  in  $S$ -page
      add  $\langle r, s \rangle$  to result
```



# Cost of Block Nested Loops

$M = 1000$  pages in R  
 $p_R = 100$  tuples per page

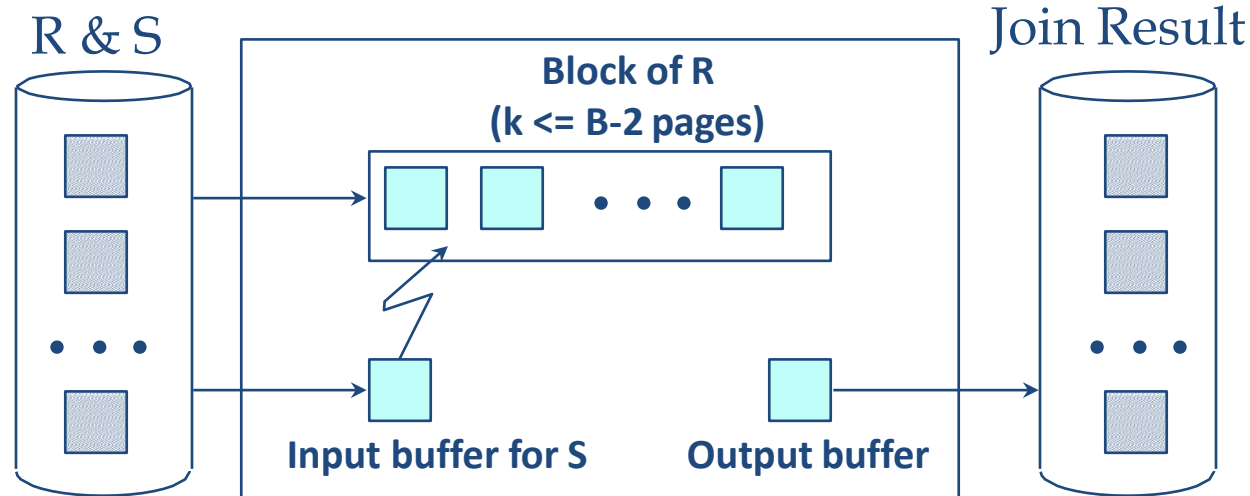
$N = 500$  pages in S  
 $p_S = 80$  tuples per page

- R is outer
- B-2 = 100-page blocks
- How many blocks of R? 10
- Cost to scan R? 1000
- Cost to scan S?  $10 * 500$
- Total Cost?  $1000 + 5000 = 6000$
- (check yourself)
  - If space for just 90 pages of R, we would scan S 12 times, cost = 7000

```
foreach block of B-2 pages of R do
  foreach page of S do {
    for all matching in-memory tuples r in R-
      block and s in S-page
      add <r, s> to result
```

- Cost: Scan of outer + #outer blocks \* scan of inner
  - #outer blocks =  $\lceil \text{\#pages of outer relation} / \text{blocksize} \rceil$

for blocked access, it might be good to equally divide buffer pages among R and S ("seek time" less)



# Algorithms for Joins

## 2. SORT-MERGE JOINS


# Sort-Merge Join

- Sort R and S on the join column
- Then scan them to do a ``merge'' (on join col.)
- Output result tuples.

# Sort-Merge Join: 1/3

- Advance scan of R until current R-tuple  $\geq$  current S tuple
  - then advance scan of S until current S-tuple  $\geq$  current R tuple
  - do this as long as current R tuple = current S tuple

**Sailors**



**S**

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

**Reserves**



**R**

<u>sid</u>	<u>bid</u>	<u>day</u>	rname
28	103	12/4/96	guppy
28	103	11/3/96	yuppy
31	101	10/10/96	dustin
31	102	10/12/96	lubber
31	101	10/11/96	lubber
58	103	11/12/96	dustin



## Sort-Merge Join: 2/3

- At this point, all R tuples with same value in  $R_i$  (*current R group*) and all S tuples with same value in  $S_j$  (*current S group*)
  - match
  - find all the equal tuples
  - output  $\langle r, s \rangle$  for all pairs of such tuples

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

<u>sid</u>	<u>bid</u>	<u>day</u>	rname
28	103	12/4/96	guppy
28	103	11/3/96	yuppy
31	101	10/10/96	dustin
31	102	10/12/96	lubber
31	101	10/11/96	lubber
58	103	11/12/96	dustin

WRITE TWO OUTPUT TUPLES

# Sort-Merge Join: 3/3

- Then resume scanning R and S

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

**S** →

WRITE THREE OUTPUT TUPLES

<u>sid</u>	<u>bid</u>	<u>day</u>	rname
28	103	12/4/96	guppy
28	103	11/3/96	yuppy
31	101	10/10/96	dustin
31	102	10/12/96	lubber
31	101	10/11/96	lubber
58	103	11/12/96	dustin

→ **R**

# Sort-Merge Join: 3/3

- ... and proceed till end

S	<u>sid</u>	sname	rating	age
	22	dustin	7	45.0
	28	yuppy	9	35.0
	31	lubber	8	55.5
→	44	guppy	5	35.0
	58	rusty	10	35.0

NO MATCH, CONTINUE SCANNING S


R	<u>sid</u>	<u>bid</u>	<u>day</u>	rname
	28	103	12/4/96	guppy
	28	103	11/3/96	yuppy
	31	101	10/10/96	dustin
	31	102	10/12/96	lubber
	31	101	10/11/96	lubber
→	58	103	11/12/96	dustin

# Sort-Merge Join: 3/3

- ... and proceed till end

**S**

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0



WRITE ONE OUTPUT TUPLE

**R**

<u>sid</u>	<u>bid</u>	<u>day</u>	rname
28	103	12/4/96	guppy
28	103	11/3/96	yuppy
31	101	10/10/96	dustin
31	102	10/12/96	lubber
31	101	10/11/96	lubber
58	103	11/12/96	dustin



# Example of Sort-Merge Join

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

<u>sid</u>	<u>bid</u>	<u>day</u>	rname
28	103	12/4/96	guppy
28	103	11/3/96	yuppy
31	101	10/10/96	dustin
31	102	10/12/96	lubber
31	101	10/11/96	lubber
58	103	11/12/96	dustin

- Typical Cost:  $O(M \log M) + O(N \log N) + (M+N)$ 
  - ignoring B (as the base of log)
  - cost of sorting R + sorting S + merging R, S
  - The cost of scanning in merge-sort,  $M+N$ , could be  $M*N$ !
    - assume the same single value of join attribute in both R and S
    - but it is extremely unlikely

# Cost of Sort-Merge Join

$M = 1000$  pages in R  
 $p_R = 100$  tuples per page

$N = 500$  pages in S  
 $p_S = 80$  tuples per page

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

<u>sid</u>	<u>bid</u>	<u>day</u>	rname
28	103	12/4/96	guppy
28	103	11/3/96	yuppy
31	101	10/10/96	dustin
31	102	10/12/96	lubber
31	101	10/11/96	lubber
58	103	11/12/96	dustin

- 100 buffer pages
- Sort R:
  - (pass 0)  $1000/100 = 10$  sorted runs
  - (pass 1) merge 10 runs
  - read + write, 2 passes
  - $4 * 1000 = 4000$  I/O
- Similarly, Sort S:  $4 * 500 = 2000$  I/O
- Second merge phase of sort-merge join
  - another  $1000 + 500 = 1500$  I/O
  - assume uniform  $\sim 2.5$  matches per sid, so  $M+N$  is sufficient
- Total 7500 I/O

- Check yourself:
  - Consider #buffer pages 35, 100, 300
  - Cost of sort-merge = 7500 in all three
  - Cost of block nested 16500, 6500, 2500

# Algorithms for Joins

## 3. HASH JOINS

# Two Phases

## 1. Partition Phase

- partition R and S using the same hash function  $h$

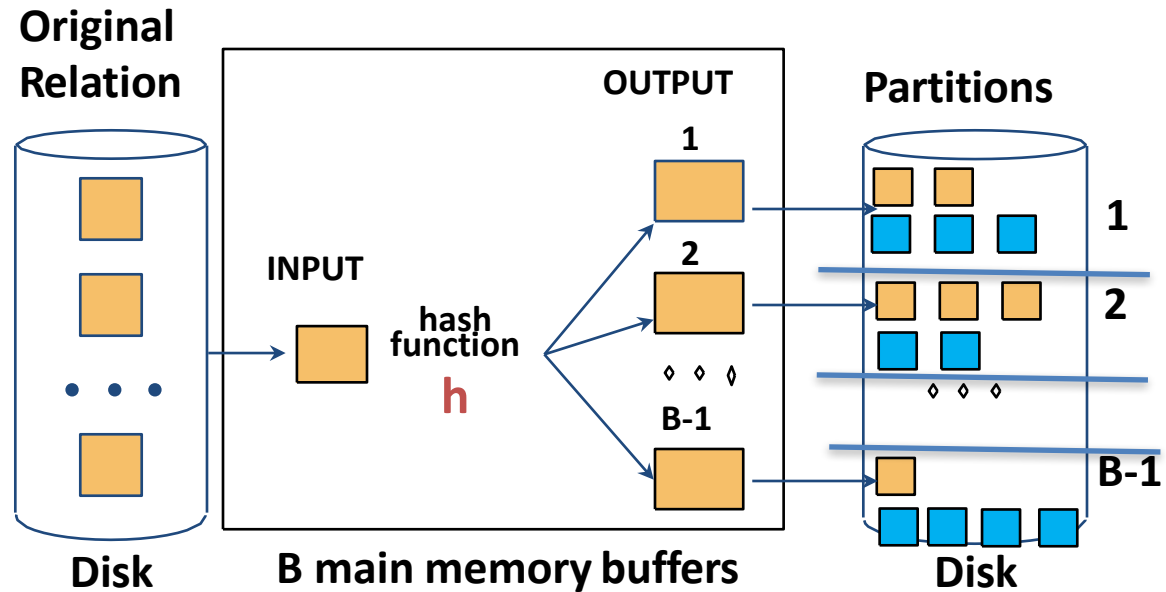
## 2. Probing Phase

- join tuples from the same partition (same  $h(..)$  value) of R and S
- tuples in different partition of  $h$  will never join
- use a “different” hash function  $h_2$  for joining these tuples
  - (why different – see next slide first)

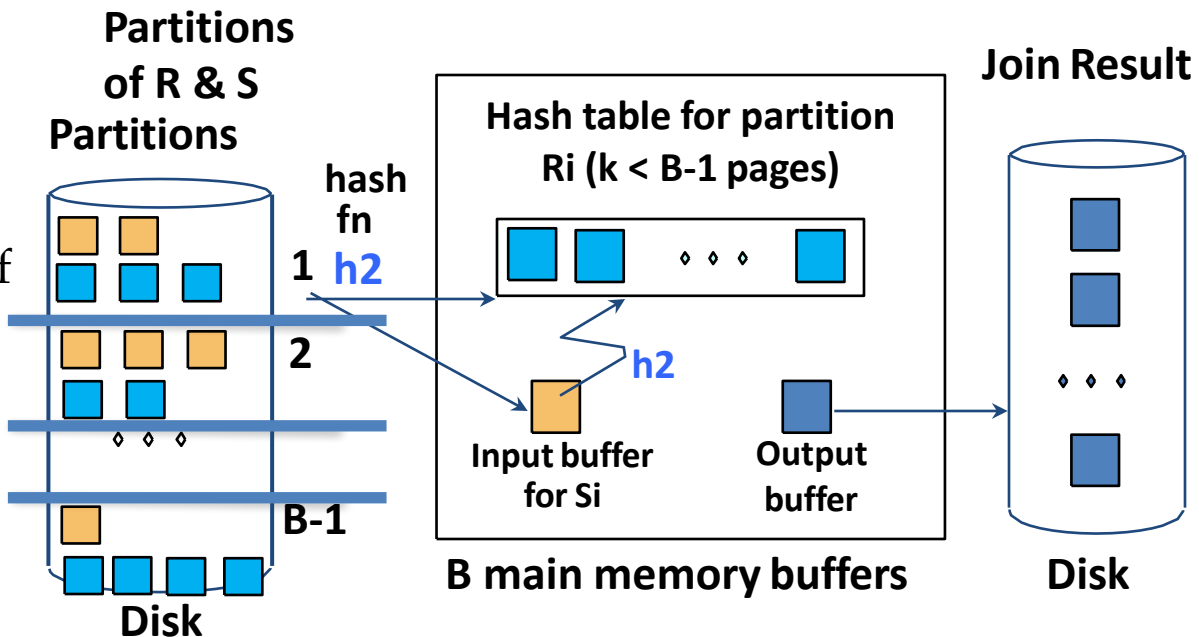


# Hash-Join

- Partition both relations using hash function  $h$
- $R$  tuples in partition  $i$  will only match  $S$  tuples in partition  $i$



- ❖ Read in a partition of  $R$ , hash it using  $h_2 (\neq h)$ .
- ❖ Scan matching partition of  $S$ , search for matches.



# Cost of Hash-Join

- In partitioning phase
  - read+write both relns;  $2(M+N)$
  - In matching phase, read both relns;  $M+N$  I/Os
  - remember – we are not counting final write
- In our running example, this is a total of 4500 I/Os
  - $3 * (1000 + 500)$
  - Compare with the previous joins

# Sort-Merge Join vs. Hash Join

- Both can have a cost of  $3(M+N)$  I/Os
  - if sort-merge gets enough buffer
- Hash join holds smaller relation in buffer-better if limited buffer
- Hash Join shown to be highly parallelizable
- Sort-Merge less sensitive to data skew
  - also result is sorted

Other operator algorithms

# Algorithms for Selection

```
SELECT *  
FROM Reserves R  
WHERE R.rname = 'Joe'
```

- No index, unsorted data
  - Scan entire relation
  - May be expensive if not many 'Joe's
- No index, sorted data (on 'rname')
  - locate the first tuple, scan all matching tuples
  - first binary search, then scan depends on matches
- B+-tree index, Hash index
  - Discussed earlier
  - Cost of accessing data entries + matching data records
  - Depends on clustered/unclustered
- More complex condition like `day<8/9/94 AND bid=5 AND sid=3`
  - Either use one index, then filter
  - Or use two indexes, then take intersection, then apply third condition
  - etc.

# Algorithms for Projection

```
SELECT DISTINCT  
       R.sid, R.bid  
FROM   Reserves R
```

- Two parts
  - Remove fields: **easy**
  - Remove duplicates (if distinct is specified): **expensive**
- Sorting-based
  - Sort, then scan adjacent tuples to remove duplicates
  - Can eliminate unwanted attributes in the first pass of merge sort
- Hash-based
  - Exactly like hash join
  - Partition only one relation in the first pass
  - Remove duplicates in the second pass
- Sort vs Hash
  - Sorting handles skew better, returns results sorted
  - Hash table may not fit in memory – sorting is more standard
- Index-only scan may work too
  - If all required attributes are part of index

# Algorithms for Set Operations

- Intersection, cross product are special cases of joins
- Union, Except
  - Sort-based
  - Hash-based
  - Very similar to joins and projection

# Algorithms for Aggregate Operations

- SUM, AVG, MIN etc.
  - again similar to previous approaches
- Without grouping:
  - In general, requires scanning the relation.
  - Given index whose search key includes all attributes in the `SELECT` or `WHERE` clauses, can do index-only scan
- With grouping:
  - Sort on group-by attributes
  - or, hash on group-by attributes
  - can combine sort/hash and aggregate
  - can do index-only scan here as well