

Cognitive Robotics

06. Non-Parametric Filters: Discrete Filter, Particle Filter, Monte Carlo Localization





AbdElMoniem Bayoumi, PhD

Spring 2022


Acknowledgment

- These slides have been created by Wolfram Burgard, Dieter Fox, Cyrill Stachniss and Maren Bennewitz

$$Bel(x \mid z, u) = \alpha p(z \mid x) \int_{x'} p(x \mid u, x') Bel(x') dx'$$



Discrete Filter: Piecewise Constant




Bayes Filter Algorithm

1. Algorithm **Bayes_filter**($Bel(x), d$):
 2. $\eta = 0$
 3. If d is a **perceptual** data item z then
 4. For all x do
 5. $Bel'(x) = P(z | x) Bel(x)$
 6. $\eta = \eta + Bel'(x)$
 7. For all x do
 8. $Bel'(x) = \eta^{-1} Bel'(x)$
 9. Else if d is an **action** data item u then
 10. For all x do
 11. $Bel'(x) = \sum_{x'} P(x | u, x') Bel(x')$
 12. Return $Bel'(x)$
- sum over all discrete states
- motion model, Ch. 4

Piecewise Constant Representation





$$Bel(x_t = \langle x, y, \theta \rangle)$$








Implementation

- To update the belief, one has to iterate over all cells of the grid
- When the belief is peaked, one wants to avoid updating irrelevant aspects of the state space
- Monitor whether the robot is de-localized or not
- Consider the likelihood of the observation in the relevant components of the state space
- Assume a bounded Gaussian for the motion uncertainty

Grid-Based Localization



Sonars and Occupancy Grid Map




Summary: Discrete Filters

- Discrete filters are an alternative way for implementing the Bayes Filter
- Histograms for representing the density
- Can represent multi-modal beliefs and recover from localization errors
- Huge memory and processing requirements
- Accuracy depends on the resolution of the grid
- In practice: approximations needed

Motivation: Particle Filter


- Discrete filter
 - High memory complexity
 - In general: fixed resolution
- Particle filters are a way to **efficiently** represent **non-Gaussian distributions**
- Basic principle
 - Set of state hypotheses (“particles”)
 - Survival-of-the-fittest

Example: Sample-Based Localization (Sonar)




Motivation

Goal: approach for dealing with **arbitrary distributions**



Key Idea: Samples

Use **a set of weighted samples** to represent arbitrary distributions



Particle Set

- Set of weighted samples

$$\mathcal{X} = \left\{ \left\langle x^{[j]}, w^{[j]} \right\rangle \right\}_{j=1, \dots, J}$$

**state
hypothesis**



**importance
weight**

- The samples represent the posterior

$$p(x) = \sum_{j=1}^J w^{[j]} \delta_{x^{[j]}}(x)$$

Particles for Approximation

- Particles for function approximation




- The more particles fall into a region, the higher the probability of the region

How to obtain such samples?

Closed Form Sampling is Only Possible for Few Distributions

- Example: Gaussian




$$x \leftarrow \frac{1}{2} \sum_{i=1}^{12} \text{rand}(-\sigma, \sigma)$$

- How to sample from **other** distributions?

Importance Sampling Principle

- We can use a different distribution π to generate samples from f
- Account for the “differences between π and f ” using a weight $\omega = f(x)/\pi(x)$
- target f
- proposal π
- Pre-condition:

$$f(x) > 0 \rightarrow \pi(x) > 0$$



Particle Filter

- Recursive Bayes filter
- Non-parametric approach
- Models the distribution by weighted samples
- **Prediction:** draw from the proposal
- **Correction:** weigh particles by the ratio of target and proposal

**The more samples we use,
the better is the estimate!**

Particle Filter Algorithm

1. Sample the particles using the proposal distribution

$$x_t^{[j]} \sim \text{proposal}(x_t \mid \dots)$$

2. Compute the importance weights

$$w_t^{[j]} = \frac{\text{target}(x_t^{[j]})}{\text{proposal}(x_t^{[j]})}$$

3. Resampling: Draw sample i with probability $w_t^{[i]}$ and repeat J times

Monte Carlo Localization

- **Each particle is a pose hypothesis**
- **Prediction:** For each particle, sample a new pose from the the motion model

$$x_t^{[j]} \sim p(x_t \mid x_{t-1}^{[j]}, u_t)$$


- **Correction:** Weigh samples according to the observation model

$$w_t^{[j]} \propto p(z_t \mid x_t^{[j]})$$


- **Resampling:** Draw sample i with probability $w_t^{[i]}$ and repeat J times

Reminder: Odometry Motion Model

start pose




end pose



According to the estimated motion


Reminder: Odometry Motion Model



Decompose the motion into

- Traveled distance
- Start rotation
- End rotation

Reminder: Odometry Motion Model




- Uncertainty in the translation of the robot:
Gaussian over the traveled distance
- Uncertainty in the rotation of the robot:
Gaussians over start and end rotation
- **For each particle, draw a new pose by sampling from three normal distributions**


Reminder: Odometry Motion Model

- Noise in odometry $u = (\delta_{rot1}, \delta_{trans}, \delta_{rot2})$
- Example: Gaussian noise

$$u \sim \mathcal{N}(0, \Sigma)$$



Reminder: Proximity Sensor Model

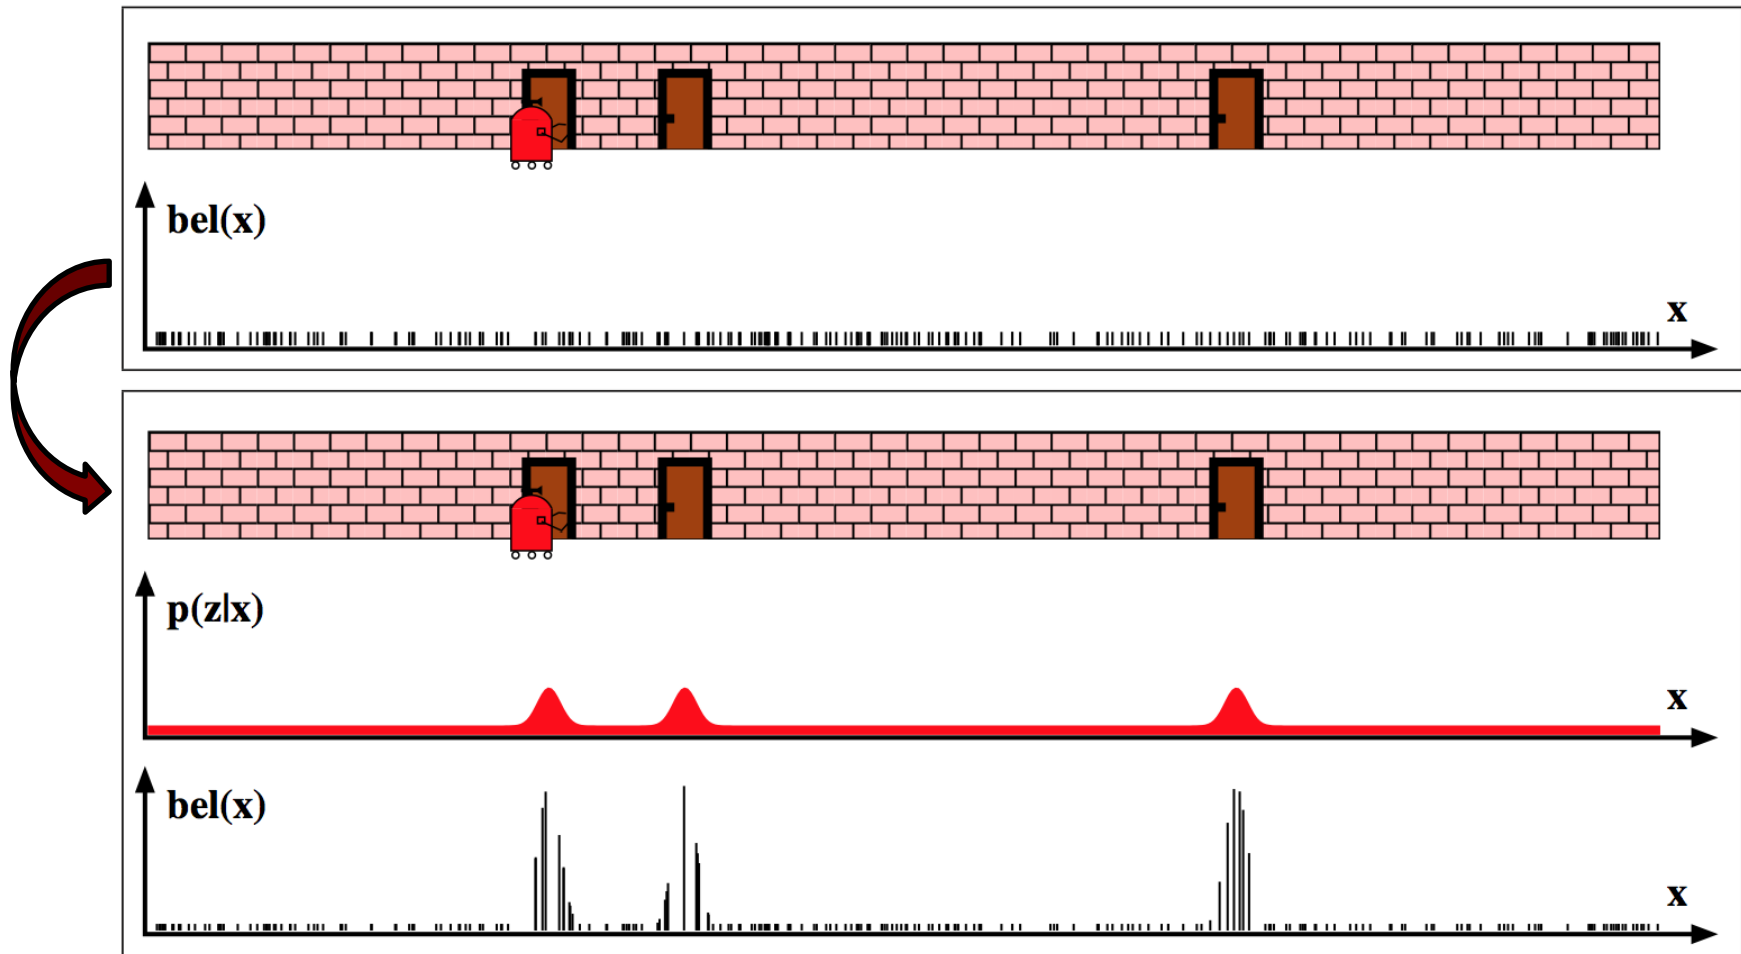


Particle Filter for Localization

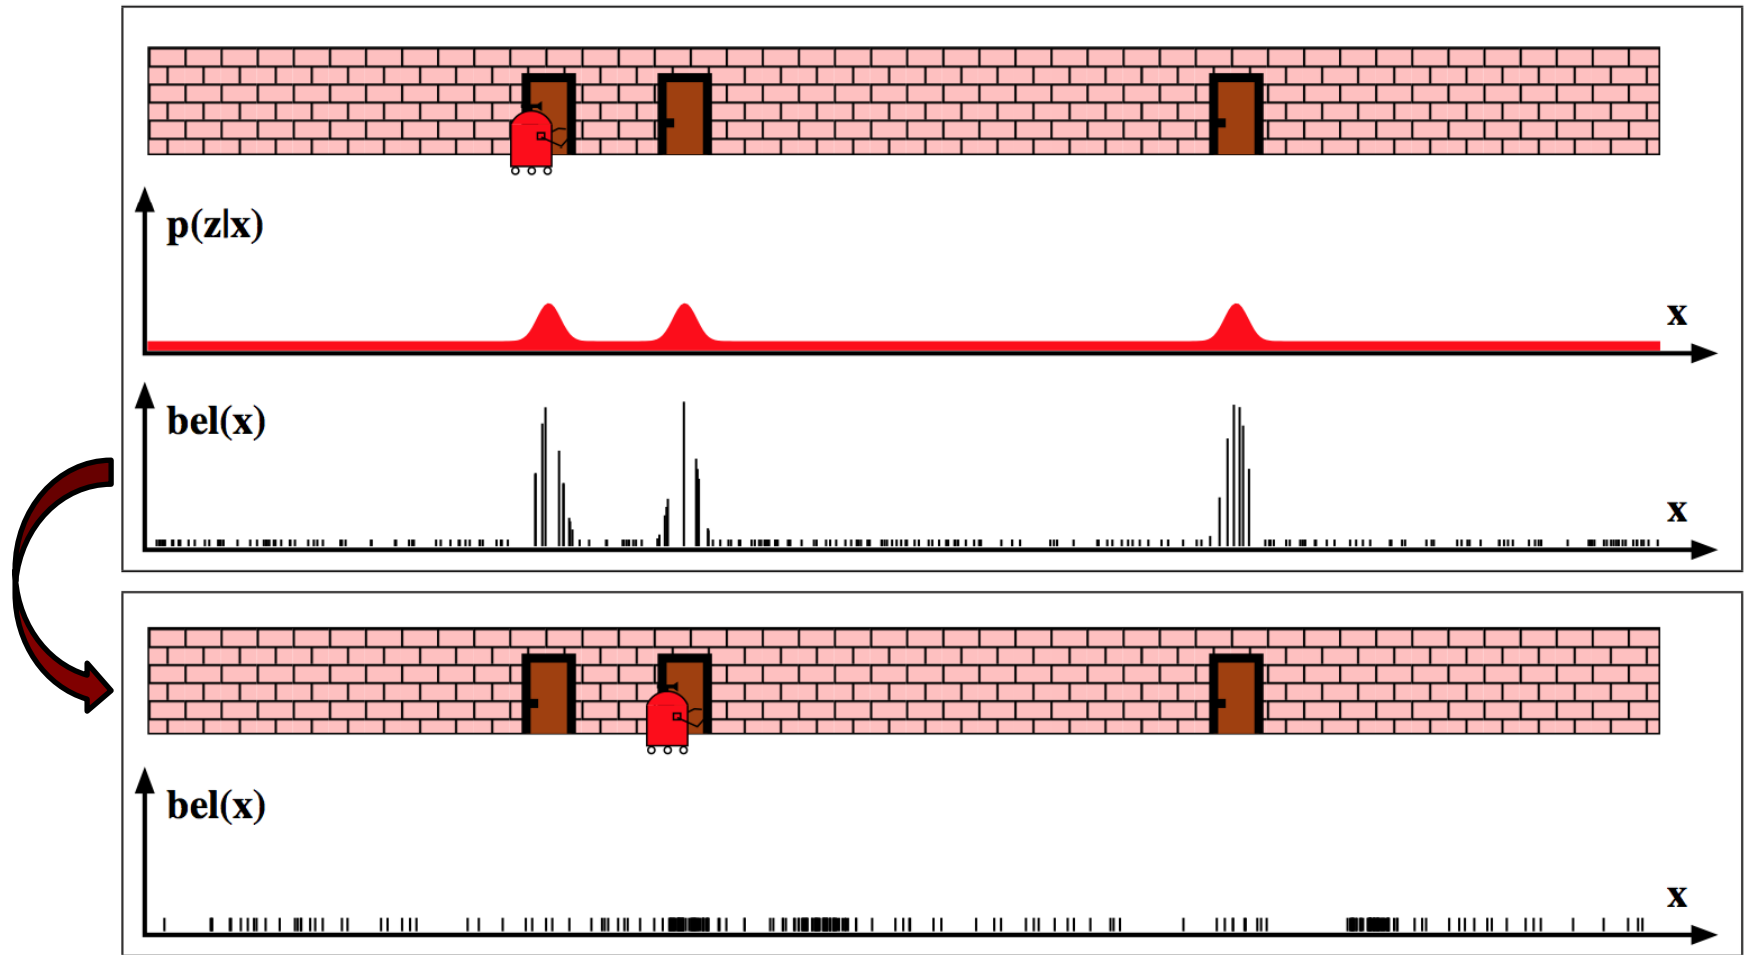
Particle_filter($\mathcal{X}_{t-1}, u_t, z_t$):

```
1:    $\bar{\mathcal{X}}_t = \mathcal{X}_t = \emptyset$ 
2:   for  $j = 1$  to  $J$  do
3:       sample  $x_t^{[j]} \sim p(x_t \mid u_t, x_{t-1}^{[j]})$ 
4:        $w_t^{[j]} = p(z_t \mid x_t^{[j]})$ 
5:        $\bar{\mathcal{X}}_t = \bar{\mathcal{X}}_t + \langle x_t^{[j]}, w_t^{[j]} \rangle$ 
6:   endfor
7:   for  $j = 1$  to  $J$  do
8:       draw  $i \in 1, \dots, J$  with probability  $\propto w_t^{[i]}$ 
9:       add  $x_t^{[i]}$  to  $\mathcal{X}_t$ 
10:  endfor
11:  return  $\mathcal{X}_t$ 
```

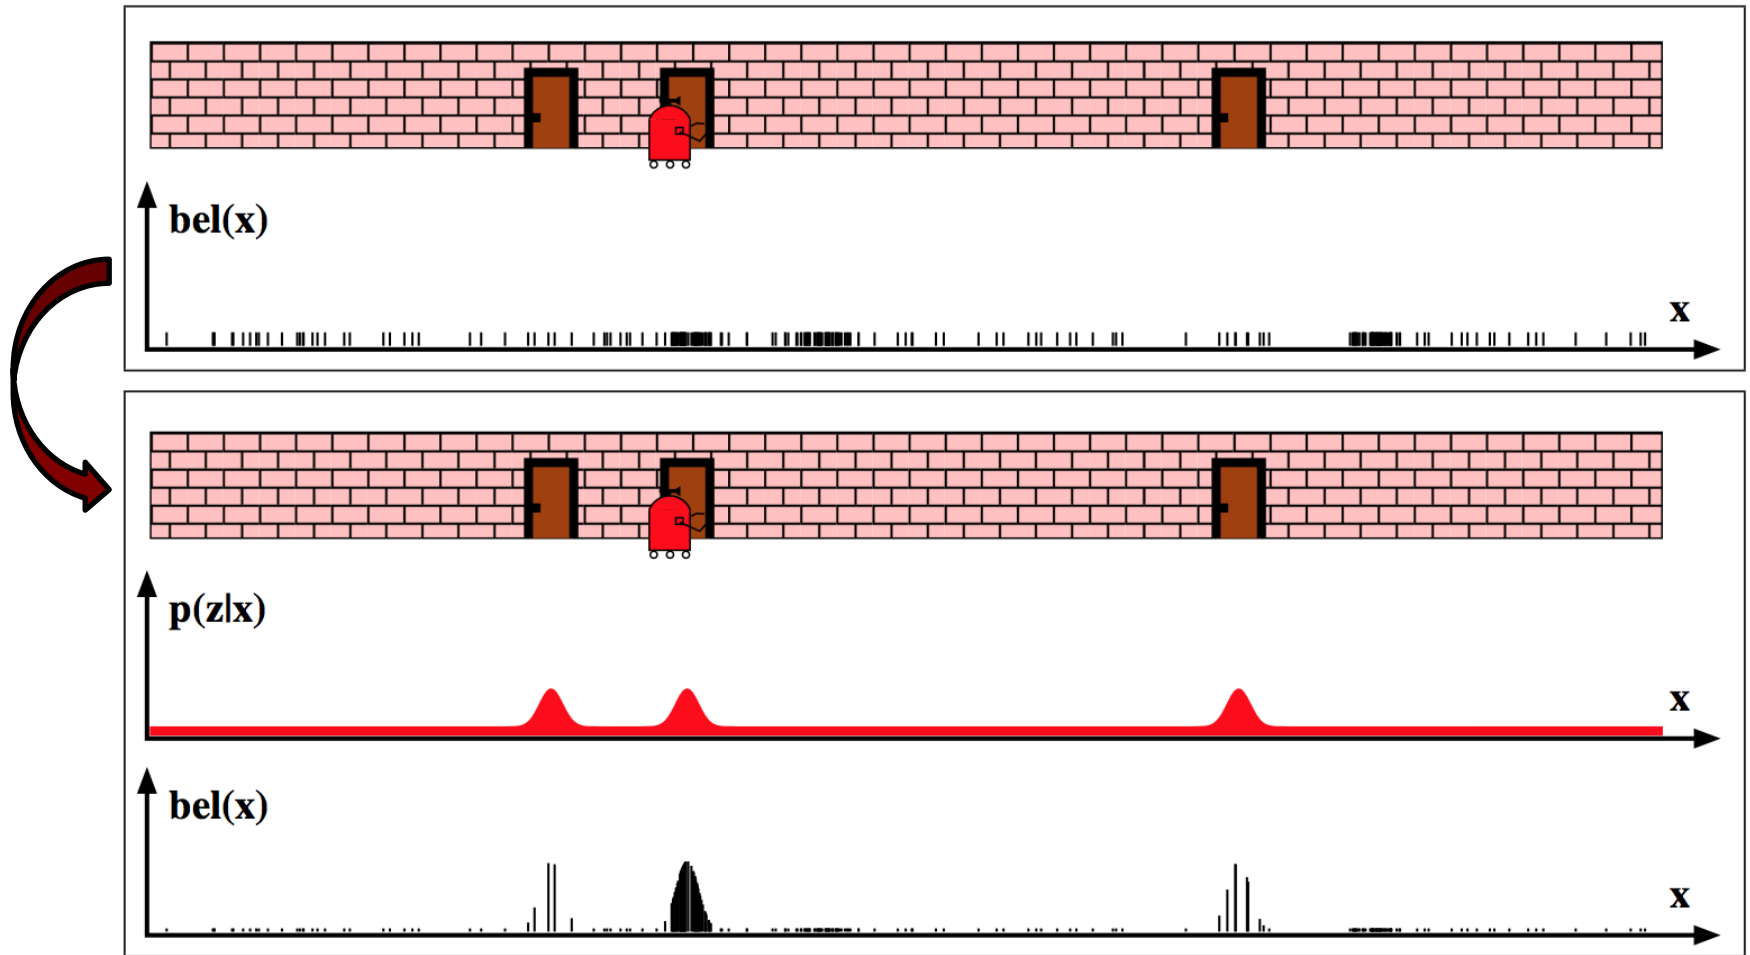
MCL – Correction Step




MCL – Prediction Step




MCL – Correction Step



MCL – Prediction Step



Example: Sample-Based Localization (Sonar)




Resampling

- Repeat J times: Draw sample i with probability $w_t^{[i]}$
- Informally: “Replace unlikely samples by more likely ones”
- Survival-of-the-fittest principle
- “Trick” to avoid that many samples cover unlikely states
- Needed as we have a limited number of samples

Assumption: normalized weights

Resampling


“roulette wheel”



- Draw randomly between 0 and 1
- Binary search
- Repeat J times
- $O(J \log J)$

initial value between 0 and $1/J$

step size = $1/J$





- Systematic resampling
- Low variance resampl.
- $O(J)$
- Also called “stochastic universal resampling”

Low Variance Resampling

Low_variance_resampling($\mathcal{X}_t, \mathcal{W}_t$):


```
1:    $\bar{\mathcal{X}}_t = \emptyset$ 
2:    $r = \text{rand}(0; J^{-1})$  initialization
3:    $c = w_t^{[1]}$  cumulative sum of weights
4:    $i = 1$ 
5:   for  $j = 1$  to  $J$  do J = #particles
6:        $U = r + (j - 1)J^{-1}$  step size = 1/J
7:       while  $U > c$  decide whether or not
8:            $i = i + 1$  to take particle i
9:            $c = c + w_t^{[i]}$ 
10:      endwhile
11:      add  $x_t^{[i]}$  to  $\bar{\mathcal{X}}_t$ 
12:   endfor
13:   return  $\bar{\mathcal{X}}_t$ 
```







Courtesy: Thrun, Burgard, Fox 37



























Courtesy: Thrun, Burgard, Fox 47

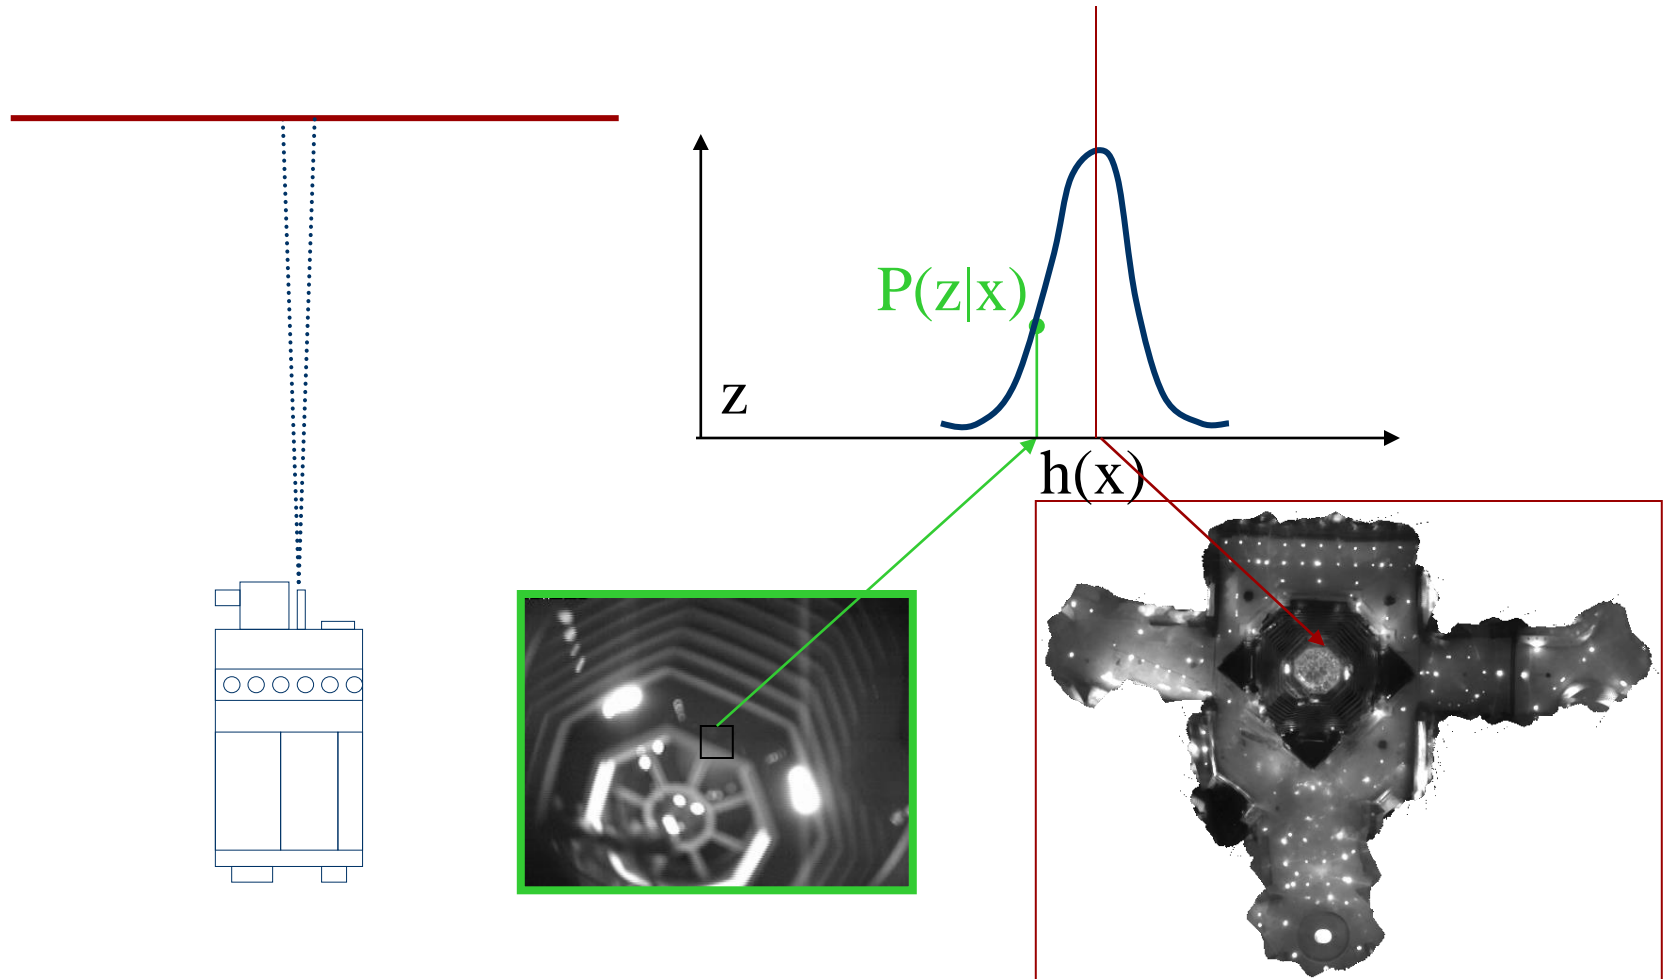


Using Ceiling Maps for Localization




[Dellaert et al. 99]

Vision-Based Localization




Under a Light

Measurement z :




$P(z/x)$:




Next to a Light

Measurement z :



$P(z/x)$:




Elsewhere


Measurement z :



$P(z/x)$:



Global Localization Using Vision



How to deal with localization errors?

- The approach described so far is able
 - to track the pose of a mobile robot and
 - to globally localize the robot
- How can we deal with localization errors (i.e., the kidnapped robot problem)?

Approaches

- At each time step, randomly insert a fixed number of samples
- Alternatively, insert random samples proportional to the average likelihood of the particles

Summary – Particle Filters

- Particle filters are non-parametric Bayes filters
- Belief represented by a set of weighted samples
- Proposal distribution to draw the samples for the next time step
- Particle weight to account for the differences between the proposal and the target
- Re-sampling: Draw new particles with a probability proportional to the weight

Summary – PF Localization

- Particles are propagated according to the motion model
- Particles are weighted according to the likelihood of the observation
- Called: Monte-Carlo localization (MCL)
- Used in many practical localization systems
- **The art is to design appropriate motion and sensor models**

Acknowledgment

- These slides have been created by Wolfram Burgard, Dieter Fox, Cyrill Stachniss and Maren Bennewitz