

Chapter 3- Part 5

The Data Link Layer

Sliding Windows Protocols (2)

Many protocols/algorithms discussed in this chapter apply to other layers



Selective Repeat _ARQ Algorithms



Selective Repeat ARQ

*Go-Back-N ARQ is inefficient in the case of a **noisy** link.*

- *In a noisy link frames have higher probability of damage , which means the resending of multiple frames.*
- *this resending consumes the bandwidth and slow down the transmission .*


Solution:

- *Selective Repeat ARQ protocol : resent only the damage frame*
- *It defines a **negative Acknolgment (NAK)** that report the sequence number of a damaged frame before the timer expires*
- *It is more efficient for noisy link, but the processing at the receiver is more complex*




Selective Repeat ARQ

- The window size is reduced to one half of 2^m
- Sender window size = receiver window size = $2^m / 2$
- Window size = (Max sequence number +1) / 2
- If $m = 2$, Window size = $4/2=2$
- Sequence number = 0, 1, 2, 3



a. Sender window



b. Receiver window

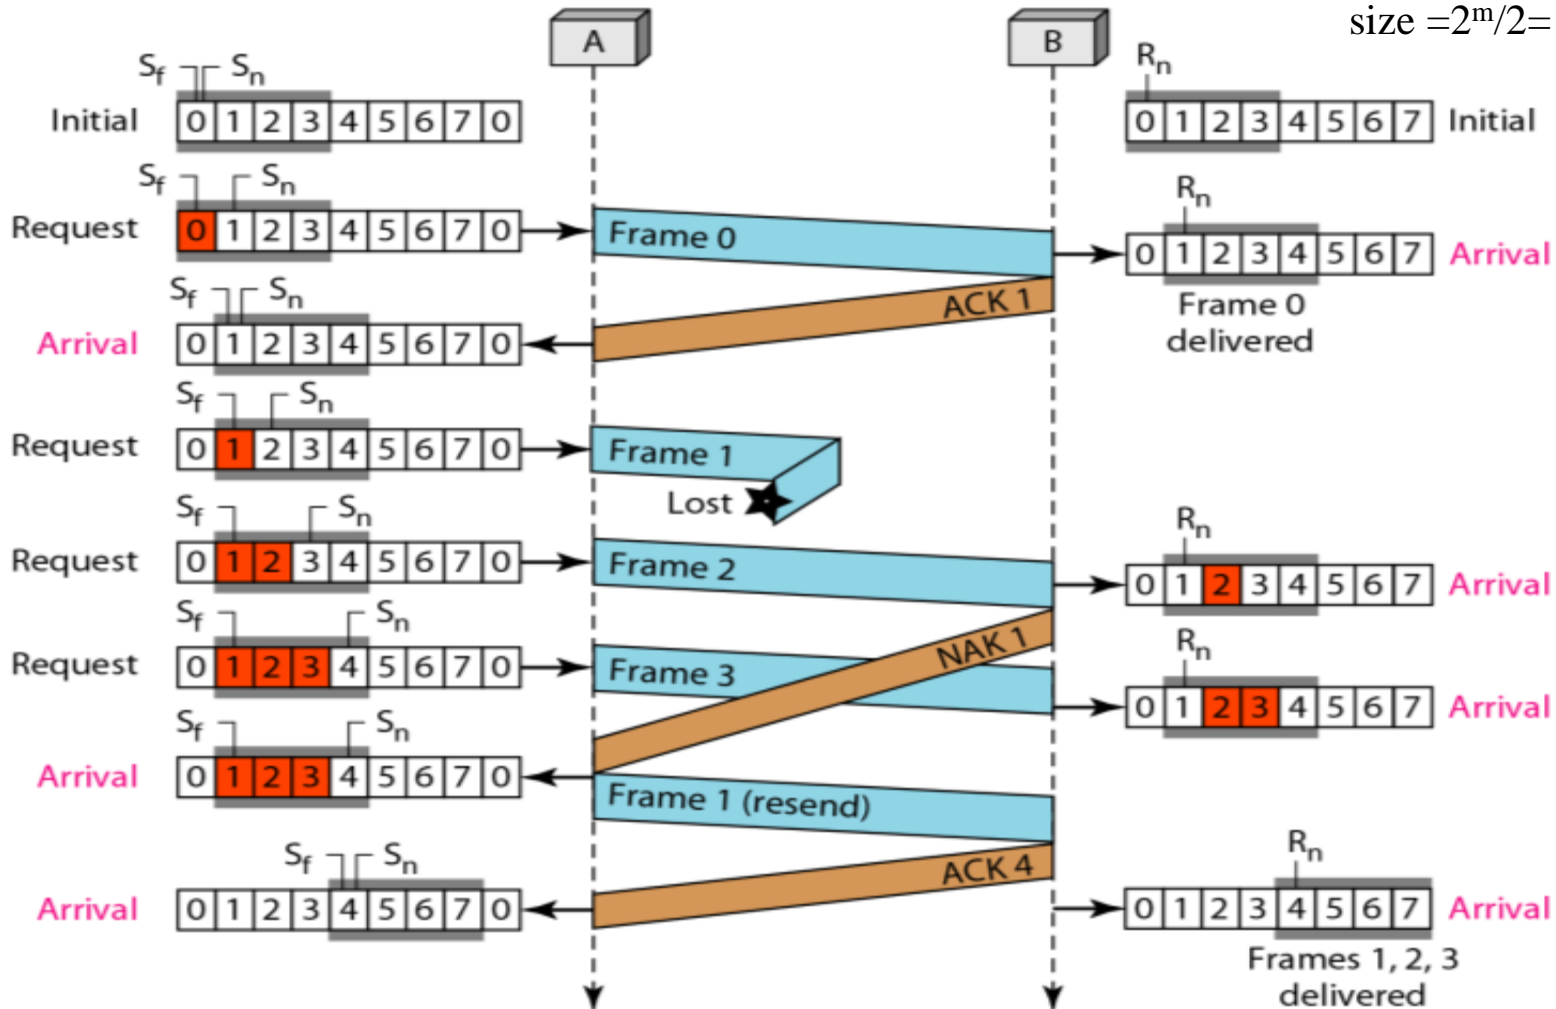


Selective Repeat ARQ

Lost Frame

$m=3$

Sequences no = $2^m = 8$:
0,1,2,3,4,5,6,7
Window
size = $2^{m-1} = 8/2 = 4$





Selective Repeat ARQ

- *At the receiver site we need to distinguish between the acceptance of a frame and its delivery to the network layer.*
- *At the second arrival , frame 2 arrives and is stored and marked , but it can not be delivered because frame 1 is missing .*
- *At the next arrival , frame 3 arrives and is marked and stored , but still none of the frames can be delivered .*
- *Only at the last arrival , when finally a copy of frame 1 arrives , can frames 1 , 2 , and 3 be delivered to the network layer.*
- ***There are two conditions for the delivery of frames to the network layer: First , a set of consecutive frames must have arrived. Second, the set starts from the beginning of the window .***



Selective Repeat ARQ

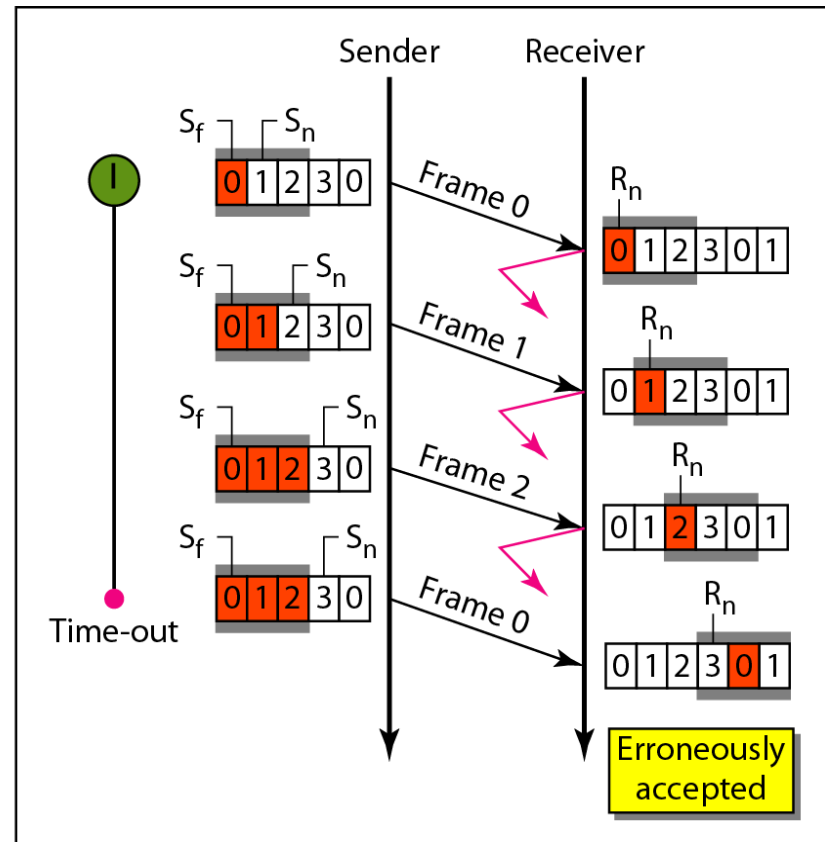

The next point is about the ACKs:

- *Notice that only two ACKs are sent here. The first one acknowledges only the first frame; the second one acknowledges three frames.*
- *In Selective Repeat, ACKs are sent when data are delivered to the network layer. If the data belonging to n frames are delivered in one shot, only one ACK is sent for all of them.*



Selective Repeat ARQ

m=2, Window size =2 , seq numbers is 4 (0,1,2,3)





Note

In Selective Repeat ARQ, the size of In the sender and receiver window must be at most one-half of 2^m .



SR_ARQ Algorithms



Selective Repeat Algorithm

- Accepts *out of order* frames
- **One timer per frame**
- On timer expire, only the frame associated with timer is re-sent \Rightarrow *selective repeat*
- Receiver window size is **fixed** at $(\text{MAX_SEQ}+1)/2$
- Sender window size starts at 0 and *grows* to $(\text{MAX_SEQ}+1)/2$
- We have timer for ACK \Rightarrow No need for reverse traffic to piggyback ACK
- For each sequence number within window, receiver has
 - a buffer
 - “arrived” bit. If set, means buffer is full
- When a packet arrives at receiver
 - Check to see if it falls within window of sequence numbers
 - Check to see if arrive bit is clear (i.e. buffer is empty)
- If both conditions are satisfied, store the arrived packet
- Receiver delivers packets to network layer in order and advances window
- Assume that the function *start_ack_timer()* **resets** the expiration time of the ACK timer to the current time.



A Sliding Window Protocol Using Selective Repeat

Sender & Receiver window size is **half** of MAX_SEQ

/* Protocol 6 (nonsequential receive) accepts frames out of order, but passes packets to the network layer in order. Associated with each outstanding frame is a timer. When the timer expires, only that frame is retransmitted, not all the outstanding frames, as in protocol 5. */

```
#define MAX_SEQ 7 /* should be 2^n - 1 */
#define NR_BUFS ((MAX_SEQ + 1)/2)
typedef enum {frame_arrival, cksum_err, timeout, network_layer_ready, ack_timeout} event_type;
#include "protocol.h"
boolean no_nak = true; /* no nak has been sent yet */
seq_nr oldest_frame = MAX_SEQ + 1; /* initial value is only for the simulator */

static boolean between(seq_nr a, seq_nr b, seq_nr c)
{
    /* Same as between in protocol5, but shorter and more obscure. */
    return ((a <= b) && (b < c)) || ((c < a) && (a <= b)) || ((b < c) && (c < a));
}
```

If frame type is **not** data, **s.seq** is **ignored**
⇒ We can put anything

```
static void send_frame(frame_kind fk, seq_nr frame_nr, seq_nr frame_expected, packet buffer[])
{
    /* Construct and send a data, ack, or nak frame. */
    frame s; /* scratch variable */

    s.kind = fk; /* kind == data, ack, or nak */
    if (fk == data) s.info = buffer[frame_nr % NR_BUFS];
    s.seq = frame_nr; /* only meaningful for data frames */
    s.ack = (frame_expected + MAX_SEQ) % (MAX_SEQ + 1);
    if (fk == nak) no_nak = false; /* one nak per frame, please */
    to_physical_layer(&s); /* transmit the frame */
    if (fk == data) start_timer(frame_nr % NR_BUFS);
    stop_ack_timer(); /* no need for separate ack frame */
}
```

- **s.ack** contains the sequence number of the **last received** frame in the **contiguous** sequence of frames
- you can think of **s.ack** as **circular(frame_expected-1, MAX_SEQ)**

If the packet that we are sending now is a NAK, then clear **no_nak**. This way if the next packet is **NOT frame_expected**, we will start the **ACK timer** instead of sending a NAK one more time

Continued → 12



A Sliding Window Protocol Using Selective Repeat (2)

```
void protocol6(void)
{
    seq_nr ack_expected;           /* lower edge of sender's window */
    seq_nr next_frame_to_send;     /* upper edge of sender's window + 1 */
    seq_nr frame_expected;         /* lower edge of receiver's window */
    seq_nr too_far;                /* upper edge of receiver's window + 1 */
    int i;                         /* index into buffer pool */
    frame r;                       /* scratch variable */
    packet out_buf[NR_BUFS];       /* buffers for the outbound stream */
    packet in_buf[NR_BUFS];        /* buffers for the inbound stream */
    boolean arrived[NR_BUFS];      /* inbound bit map */
    seq_nr nbuffered;              /* how many output buffers currently used */
    event_type event;

    enable_network_layer();        /* initialize */
    ack_expected = 0;              /* next ack expected on the inbound stream */
    next_frame_to_send = 0;        /* number of next outgoing frame */
    frame_expected = 0;
    too_far = NR_BUFS;
    nbuffered = 0;                /* initially no packets are buffered */
    for (i = 0; i < NR_BUFS; i++) arrived[i] = false;
}
```

Need buffer at **both** sender and receiver

Sender window is between **ack_expected** and **next_frame_to_send**

receiver window is **fixed** size NR_BUFS. It is *initialized* between **0** and **NR_BUFS**

Continued →



A Sliding Window Protocol Using Selective Repeat (3)

out_buf[] is **circular**

```

while (true) {
    wait_for_event(&event);
    switch(event) {
        case network_layer_ready:
            nbuffered = nbuffered + 1;
            from_network_layer(&out_buf[next_frame_to_send % NR_BUFS]); /* fetch new packet */
            send_frame(data, next_frame_to_send, frame_expected, out_buf); /* transmit the frame */
            inc(next_frame_to_send); /* advance upper window edge */
            break;

        case frame_arrival:
            from_physical_layer(&r); /* a data or control frame has arrived */
            if (r.kind == data) { /* fetch incoming frame from physical layer */
                /* An undamaged frame has arrived. */
                if ((r.seq != frame_expected) && no_nak)
                    send_frame(nak, 0, frame_expected, out_buf); else start_ack_timer();
                if (between(frame_expected, r.seq, too_far) && (arrived[r.seq % NR_BUFS] == false)) {
                    /* Frames may be accepted in any order. */
                    arrived[r.seq % NR_BUFS] = true; /* mark buffer as full */
                    in_buf[r.seq % NR_BUFS] = r.info; /* insert data into buffer */
                    while (arrived[frame_expected % NR_BUFS]) {
                        /* Pass frames and advance window. */
                        to_network_layer(&in_buf[frame_expected % NR_BUFS]);
                        no_nak = true;
                        arrived[frame_expected % NR_BUFS] = false;
                        inc(frame_expected); /* advance lower edge of receiver's window */
                        inc(too_far); /* advance upper edge of receiver's window */
                        start_ack_timer(); /* to see if a separate ack is needed */
                    }
                }
            }
    }
}

```

- We did NOT receive the lower end of receive window
⇒ Send NAK immediately
- Because **send_frame()** sets **s.ack** to **circular(frame_expected-1)** then the NAK frame contains the sequence number of the last received frame
- Thus a when a sender receives a NAK, it should re-send the frame with sequence number **circular(r.ack+1)**

Because we are sending a NAK, this value is meaningless
⇒ Put any thing

Restart ACK timer every time we receive out of order packet ([See WHY later?](#))

- Increment **both** lower and upper receiver window boundaries
⇒ Receiver window size always fixed at **NR_BUFS**

Deliver **contiguous** packet sequence to network layer starting from **bottom** of receiver window.
frame_expected is **one more than** the sequence number of the **last received frame** in **contiguous** sequence

Allow sending NAK because we send NAK for **frame_expected** and **frame_expected** will be incremented soon

Reset ACK timer to the current time



A Sliding Window Protocol Using Selective Repeat (4)

- When the other side sends a NAK, the other side sets the “ack” field to **circular (frame_expected - 1)**.
- Hence **r.ack** contains sequence number of the last received frame by the other side.
- Thus a sender should re-send the frame whose sequence number is **circular(r.ack+1)**
- Thus we have to test **between** for **circular(r.ack+1)**

- ACK means all frames in the **contiguous** sequence of frames **before and including** ACKed frame have been received. Do the following **between ack_expected** and **r.ack**

⇒ Free buffers
⇒ Stop all retransmit timers
⇒ Advance sender **lower** window

We assume that the timeout event causes the variable **oldest_frame** to be set according to the timeout that just expired

- On ACK timeout, we send ACK for the frame **before** bottom of the receiver window because **frame_expected** is one more than the sequence number of the last frame received in **contiguous** sequence
- Because **send_frame()** sets **s.ack** to **circular(frame_expected-1)**, we will end up **re-ACKing** the last frame received in the **contiguous** sequence

Because we have ACK timeout
⇒ No need for reverse traffic
⇒ Solves the blocking problem of Go back N

```
if((r.kind==nak) && between(ack_expected,(r.ack+1)%(MAX_SEQ+1),next frame to send))  
    send_frame(data, (r.ack+1) % (MAX_SEQ + 1), frame_expected, out_buf);
```

```
while (between(ack_expected, r.ack, next_frame_to_send)) {  
    nbuffered = nbuffered - 1;          /* handle piggybacked ack */  
    stop_timer(ack_expected % NR_BUFS); /* frame arrived intact */  
    inc(ack_expected);                  /* advance lower edge of sender's window */  
}
```

```
break;
```

```
case cksum_err:
```

```
if (no_nak) send_frame(nak, 0, frame_expected, out_buf); /* damaged frame */  
break;
```

```
case timeout:
```

```
send_frame(data, oldest_frame, frame_expected, out_buf); /* we timed out */  
break;
```

```
case ack_timeout:
```

```
send_frame(ack, 0, frame_expected, out_buf); /* ack timer expired; send ack */
```

```
if (nbuffered < NR_BUFS) enable_network_layer(); else disable_network_layer();
```

Re-send NAK for last frame received if we get a corrupted frame

Window size is NR_BUF
⇒ half of MAX_SEQ



A Sliding Window Protocol Using Selective Repeat

Discussion Points

- Why is NR_BUFS defined $(MAX_SEQ+1)/2$?
- How many buffers must receiver have?
- How many timers are needed at the receiver?
- How to get away without reverse traffic?
- How is NAK making protocol more efficient
- How to adjust timers
- *Why do we have to restart ACK timer every time we receive out of order packet?*



A Sliding Window Protocol Using Selective Repeat

Why $NR_BUFS (MAX_SEQ+1)/2$

- Suppose we allow NR_BUFS to be MAX_SEQ-1 (just like Go Back N)
- Consider the following sequence of events
- Sender window is 0,1,2,3,4,5,6
 - $next_frame_to_send = 7, ack_expected = 0$
- Receiver window is 0,1,2,3,4,5,6
- Sender sends frames 0,1,2,3,4,5,6
- Receiver receives all seven frames and sends ACK for all of them
 - All frames are valid frames because their sequence numbers lie within receiver window
 - Receiver delivers all 7 frames to network layer and sends ACK for 0,1,2,3,4,5,6
 - Receiver advances its receive window to 7,0,1,2,3,4,5
 - $\Rightarrow Frame_expected = 7, too_far = 6$
- Suppose all ACKs are lost
- Sender times out. Thus it **resends** frame 0 because frame 0 timer is the first timer started and hence it will be the first timer to expire
- Receiver receives **resent** data frame with $s.seq = 0$
 - $S.ack = 0$ is within the new receiver window
 - Receiver accepts frame 0 \Rightarrow **DUPLICATE** packet because this packet is already delivered to network layer
 - $arrived[0] = TRUE$
 - Does NOT advance $frame_expected$ because $frame_expected$ not received yet



A Sliding Window Protocol Using Selective Repeat

Why $NR_BUFS (MAX_SEQ+1)/2$

- Now receiver ACK timeout expires
 - Resends ACK with $s.ack = circular(frame_expected - 1) = 6$
- Sender receives ACK with $s.ack = 6$
 - Sender assumes that all frames 0,1,2,3,4,5,6 have been received
 - Sender advances its window to so that the sender window = 7,0,1,2,3,4,5
 - Sender sends frames 7,0,1,2,3,4,5
 - ⇒ $ack_expected = 7, next_frame_to_send = 6$
- Receiver receives frame 7,0,1,2,3,4,5
 - Receiver **accepts** frame 7 and puts it in buffer
 - Receiver **rejects** frame 0 because $arrived[0]=TRUE \Rightarrow$ The new frame 0 is **LOST**
 - Receiver **Accepts** 1,2,3,4,5 and puts them in buffer
 - Now receiver has a contiguous sequence of frames in the receive window so it delivers 7,0,1,2,3,4,5 to network layer
 - The delivered packets have the following problems
 - Frame 0 is **duplicate** and **out of sequence** because it belongs to the first batch
 - Frame 0 from the second batch is **lost**
- NOTE:
 - When Receiver receives frame 0 for the second time, it will send a NAK for $frame_expected$. Thus $s.ack=6$ in the NAK
 - The sender receives a NAK with $r.ack = 6$
 - The sender checks if **$circular(r.ack+1)$** lies between **$ack_expected$** and **$next_frame_to_send$** before resending the NAKed frame
 - $circular(r.ack+1) = 7, ack_expected = 0,$ and $next_frame_to_send = 7$
 - Thus the check **Fails**
 - Hence the NAK will have **no effect**



A Sliding Window Protocol Using Selective Repeat

How many buffers and timers at receiver

- Receiver only accepts frames within the window size
- Window size is **half** the sequence number*
- Receiver needs buffers equal to window size **not** MAX_SEQ
- We need one timer **ONLY** at receiver: This is the ACK timer
- We need timers equal to window size, not MAX_SEQ at the sender



A Sliding Window Protocol Using Selective Repeat Auxiliary Timer and NAK

- Auxiliary ACK timer
- No need for reverse traffic
- Auxiliary ACK timer should be *shorter* than retransmit timer
- NAK to expedite retransmit (not needed for correctness)
- Send NAK on
 - Out of sequence: $r.seq \neq frame_expected$
 - Checksum error: possible damage
 - In both cases, we send NAK for **circular(frame_expected - 1)**
 - Remember that $circular(x-1) = (x > 0) ? (x--) : MAX_SEQ$
- How to avoid multiple NAKs per packet
 - Variable `no_nak` is set only if receiver has **not** sent NAK for `frame_expected`



A Sliding Window Protocol Using Selective Repeat

Adjusting timers

- Sender timeout at sender should be *slightly larger* than roundtrip time
- Problem if round trip time is variable
- If reverse traffic is sporadic, ACK will be irregular, making it hard for receiver to estimate roundtrip time
- Receiver auxiliary ACK timer should be appreciably **shorter** than sender retransmit timeout



Why Restart ACK timer For Out of Order packet?

- Consider the following sequence of events
 - Sender sends packets 0,1,2,3
 - Receiver receives packets 0,1,2,3
 - Receiver sends ACK for all of them and advances window to 4,5,6,7
 - All ACKs are lost
 - *How will the protocol recover?*



Go Back N vs. Selective Repeat

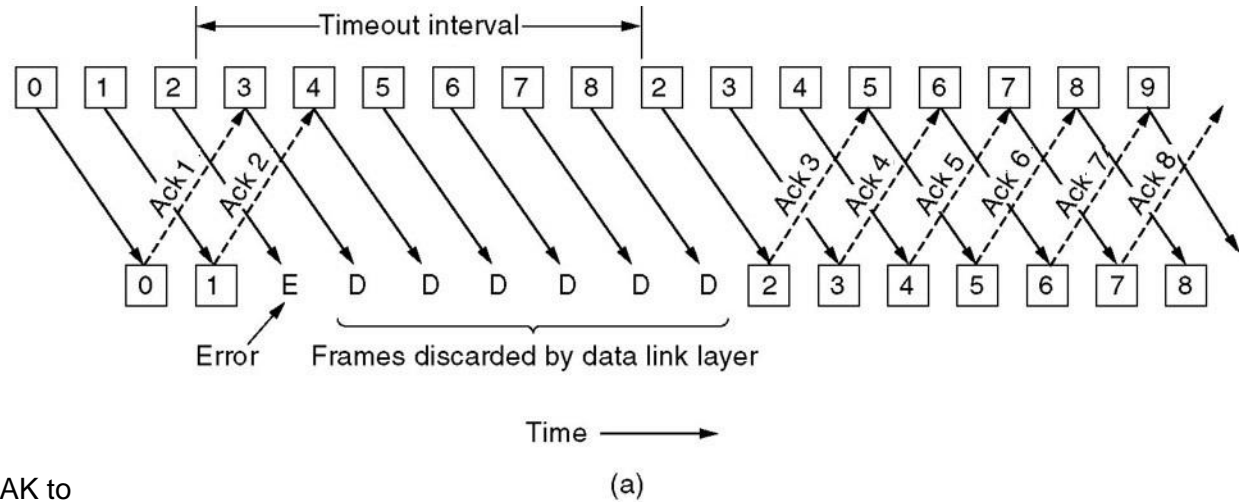
- Go Back N
 - Receiver discards all frames after lost or damaged frame
 - Receiver acknowledges received frame
 - Receiver does NOT send any ACK for frames received after lost or damaged frame
 - Relies on sender timeout
 - Equivalent to receiver window of size 1
 - Wastes a lot of bandwidth in case of error
 - Receiver needs to *buffer 1 frame* only
- Selective Repeat
 - Receiver buffers all frames received within window
 - Receiver acknowledges *last received in sequence frame* for every out of sequence frame
 - ⇒ **Cumulative** ACK
 - On timeout, sender only re-transmits *oldest unacknowledged* frame
 - Receiver can deliver all buffered frames, *in sequence*, to the network layer
 - Receiver often use NAK* to stimulate retransmission before timeout
 - Receiver needs to *buffer multiple frames* up to window size
- *Tradeoff between buffer space and link bandwidth utilization*



Go Back N vs. Selective Repeat


Expected frame not
last received frame

Go Back N



Selective Repeat
with NAK

Use NAK to
stimulate retransmit






Example Data Link Protocols

- HDLC – High-Level Data Link Control
- PPP - The Data Link Layer in the Internet



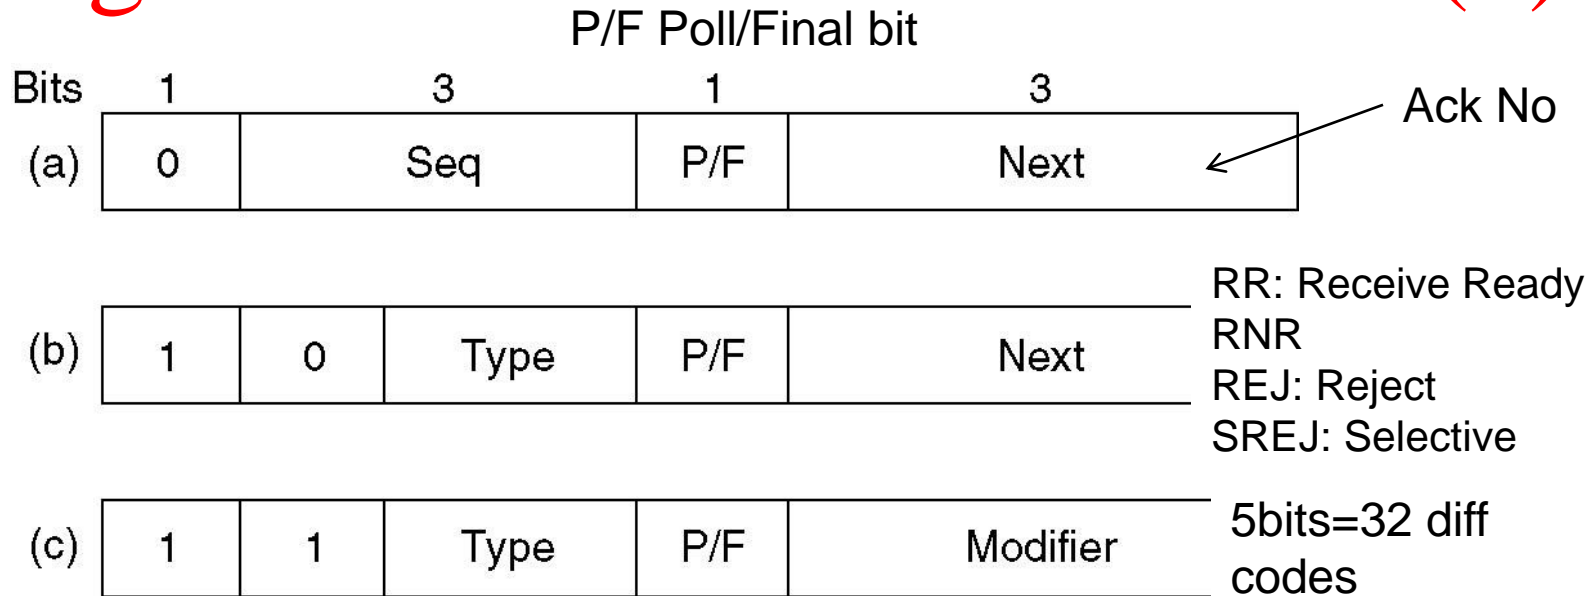
High-Level Data Link Control

- Frame format for bit-oriented protocols.
- HDLC uses bit stuffing and allows non-integer frames of, say, 30.25 bytes
- HDLC provides reliable transmission with a sliding window





High-Level Data Link Control (2)

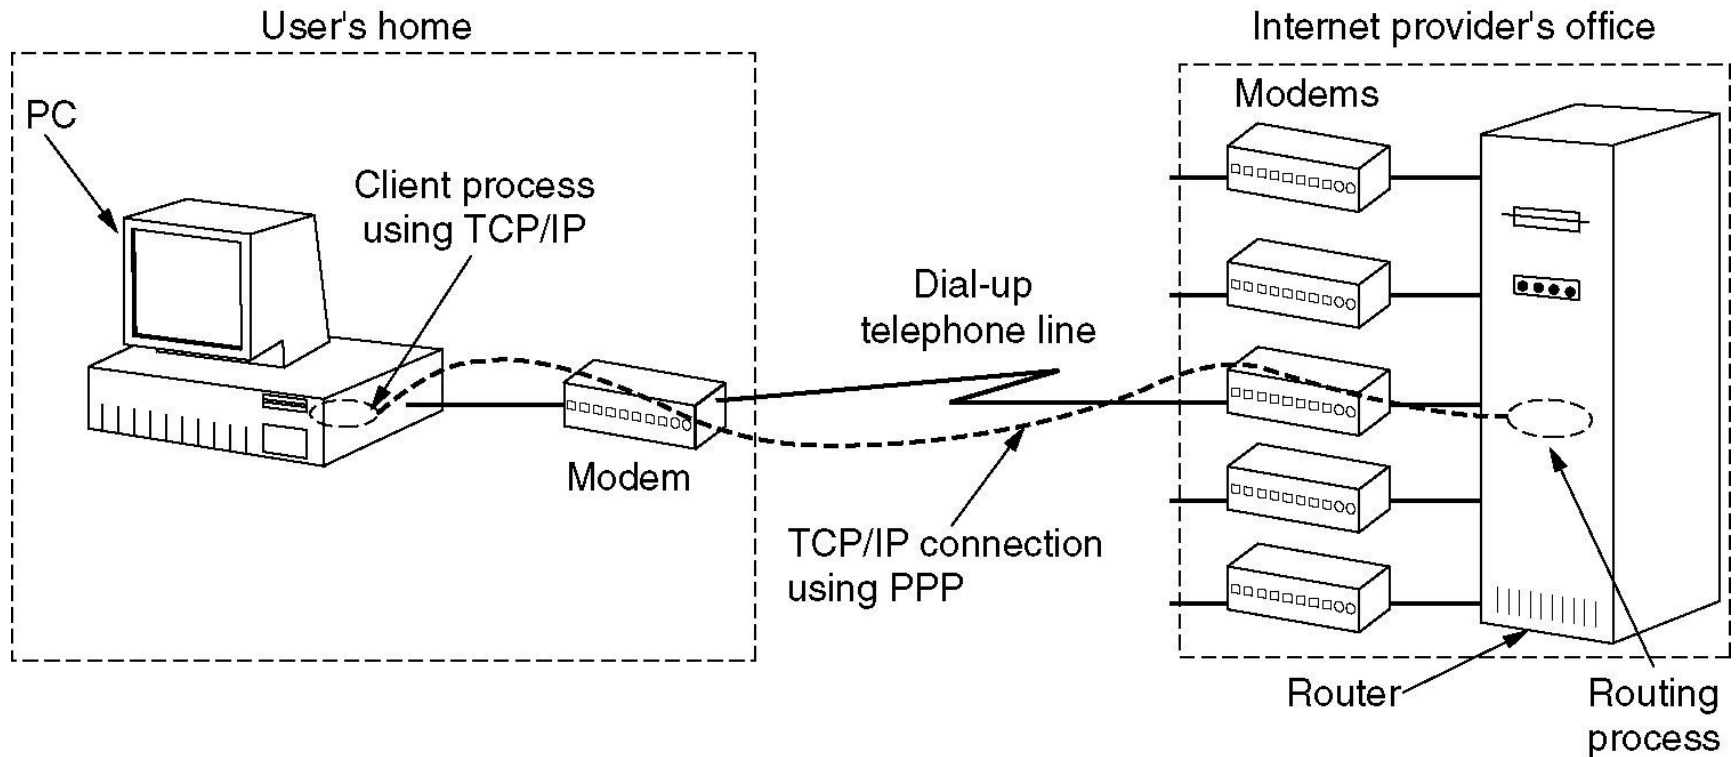


- (a) An information frame starts with 0.
- (b) A supervisory frame 10: for flow and error control when piggybacking is not required. Control functions such as acknowledgment of frames, requests for re-transmission, and requests for temporary suspension of frames being transmitted
- (c) An unnumbered frame 11: also used for control purposes. It is used to perform link initialization, link disconnection and other link control functions.



The Data Link Layer in the Internet


A home personal computer acting as an internet host.





PPP – Point to Point Protocol

- The PPP full frame format for unnumbered mode operation.
- PPP uses byte stuffing and all frames are an integral number of bytes



Address field: This field is always set to the binary value 11111111 to indicate that all stations are to accept the frame


Control field: the default value of which is 00000011. This value indicates an unnumbered frame

Protocol field: To tell what kind of packet is in the Payload field

- Codes starting with 0 for information
- Codes starting with 1 are used for configurations (Link Control Protocol LCP, Network control Protocol NCP)



PPP – Point to Point Protocol (2)



A simplified phase diagram for bring a line up and down.



PPP – Point to Point Protocol (3)

The LCP frame types.

Name	Direction	Description
Configure-request	I → R	List of proposed options and values
Configure-ack	I ← R	All options are accepted
Configure-nak	I ← R	Some options are not accepted
Configure-reject	I ← R	Some options are not negotiable
Terminate-request	I → R	Request to shut the line down
Terminate-ack	I ← R	OK, line shut down
Code-reject	I ← R	Unknown request received
Protocol-reject	I ← R	Unknown protocol requested
Echo-request	I → R	Please send this frame back
Echo-reply	I ← R	Here is the frame back
Discard-request	I → R	Just discard this frame (for testing)

