



Cairo University

Cairo University

Cairo University
Faculty of Engineering
Computer Engineering Department

Natural Language Processing

0100 0100
0110 0110
1001 1001
0100110
0100001
1101
1101

Dr. Sandra Wahid

learning processing
natural language NLP
text linguistics interaction
automatic programming technologies
understanding data evolution
linguistics science cloud automation
interaction machine learning
interaction communication
output systems
output communication
systems
output communication
systems

Text Categorization/Classification

Text Categorization: is the task of assigning a label or category to an entire text or document.

- **Sentiment Analysis**: is one common text categorization task, concerned with the extraction of sentiment, the positive or negative orientation that a writer expresses toward some object.
 - Examples: a review of a movie, book, or product on the web expresses the author's sentiment toward the product, while an editorial or political text expresses sentiment toward a candidate or political action.
 - Extracting consumer or public sentiment is thus relevant for fields from marketing to politics.
- **Spam detection**: is the binary classification task of assigning an email to one of the two classes *spam* or *not-spam*.
- **Language id**: is the task concerned with knowing the language a certain text is written in.
- **Authorship attribution**: is a task related to text classification tasks concerned with determining a text's author.
- **Subject category classification**: is the task concerned with assigning a library subject category or topic label to a text.

Text Categorization/Classification

- Classification is essential for tasks below the level of the document as well.
 - Period **disambiguation**: deciding if a period is the end of a sentence or part of a word.
 - Word **tokenization**: deciding if a character should be a word boundary.
 - **Language modeling**: can be viewed as classification, each word can be thought of as a class, and so predicting the next word is classifying the context-so-far into a class for each next word.
 - A **part-of-speech tagger**: classifies each occurrence of a word in a sentence as, e.g., a noun or a verb.
- The goal of classification is to take a **single observation**, extract some useful **features**, and thereby **classify** the observation into one of a set of discrete **classes**.

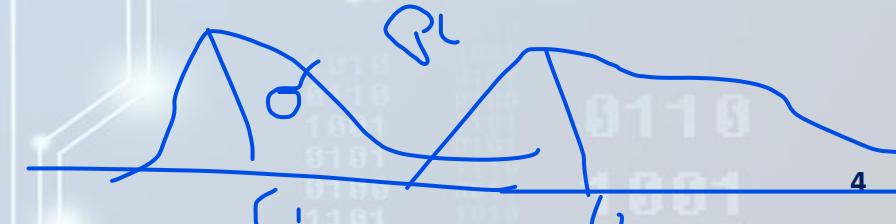


ba5ud word, atl3 menha m3lomat, w b3den a7aded hya
harbotha b class mn el classes elly 3andy

Classifiers

- **Supervised classification:** we have a training set of N documents that have each been hand-labeled with a class: $(d_1, c_1), \dots, (d_N, c_N)$. Our goal is to learn a classifier that is capable of mapping from a new document d to its correct class $c \in C$.
- **Probabilistic classifier:** additionally will tell us the probability of the observation being in the class. This full distribution over the classes can be useful information for downstream decisions; avoiding making discrete decisions early which can be useful when combining systems.

Iama bnst5dm el probability, bn5aly el model ya5ud el model bta3o mt25r, fa ye2dr ye3ml processing I more details, 3ks el discrete, momken el model yegy 3nd layer fl nos masln, w tla2y eno 5las hwa 3aml el prediction, fa ay propagation odam fl layers, msh htfr2 m3ak f 7aga, hwa keda keda khad kraro 5las.



Classifiers

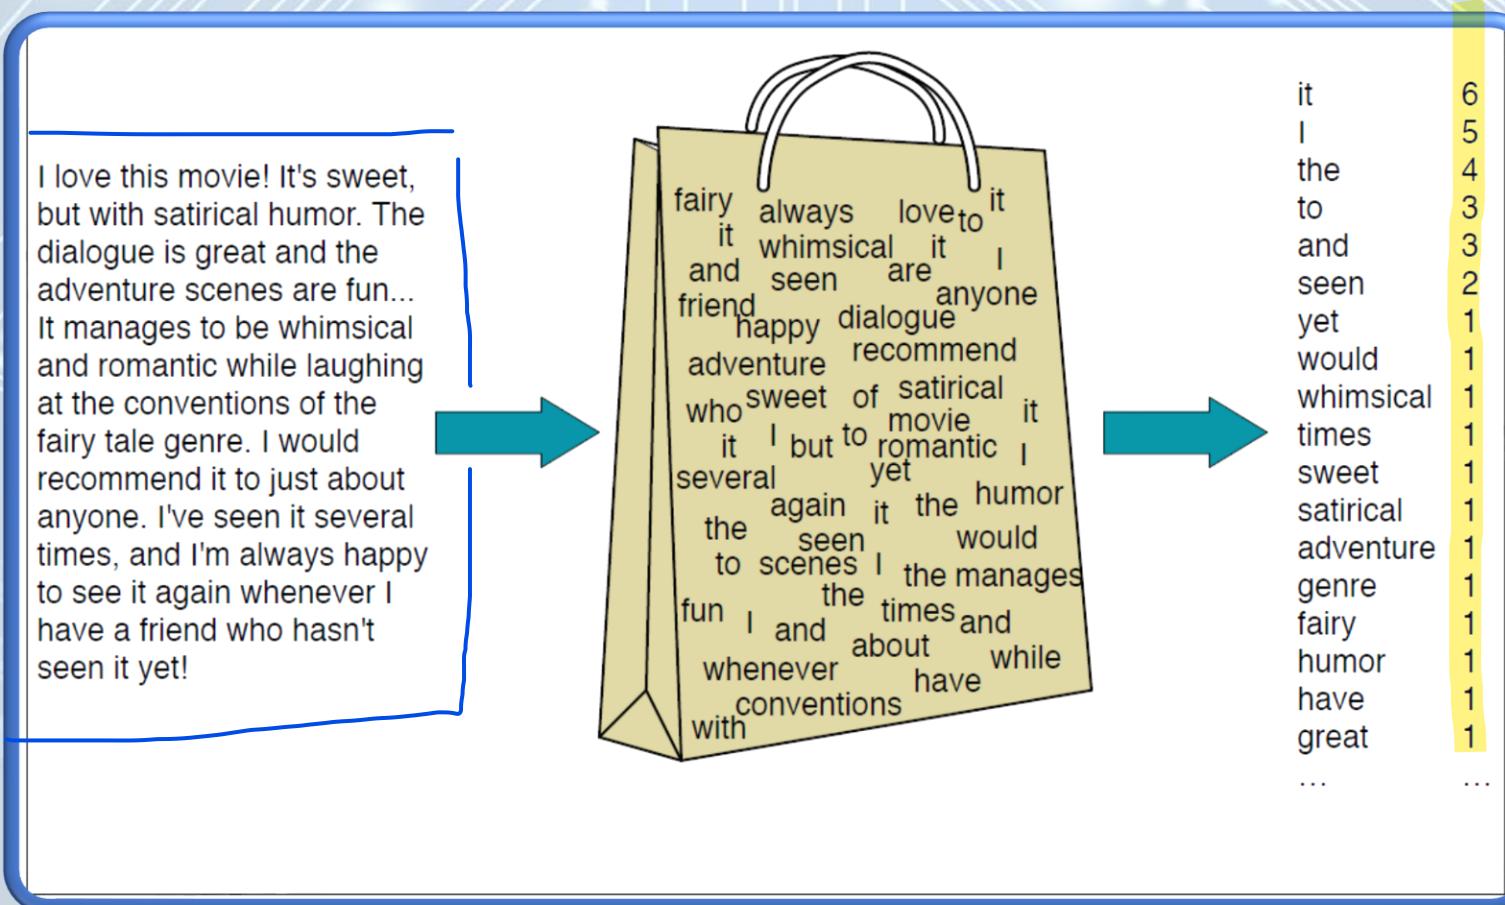
- Many kinds of machine learning algorithms are used to build classifiers:
 - **Generative classifiers**: build a model of how a class could generate some input data. Given an observation, they return the class most likely to have generated the observation, e.g. Naïve Bayes.
 - **Discriminative classifiers**: learn what features from the input are most useful to discriminate between the different possible classes, e.g. Logistic Regression.

Imagine we're trying to distinguish dog images from cat images:

- A generative model would have the goal of understanding what dogs look like and what cats look like.
- A discriminative model, by contrast, is only trying to learn to distinguish the classes (perhaps without learning much about them). So maybe all the dogs in the training data are wearing collars and the cats aren't.
- A generative model like naive Bayes makes use of the likelihood term $P(d|c)$, which expresses how to generate the features of a document if we knew it was of class c.
- By contrast a discriminative model in this text categorization scenario attempts to directly compute $P(c|d)$. It will learn to assign a high weight to document features that directly improve its ability to discriminate between possible classes, even if it couldn't generate an example of one of the classes.
- While discriminative systems are often more accurate and hence more commonly used, generative classifiers still have a role.

Naive Bayes Classifiers

- We will introduce the **multinomial naive Bayes** classifier.
- We represent a **text document** as if it were a **bag-of-words**, that is, an **unordered** set of words with their **position ignored**, keeping only their **frequency** in the document.



sebha delw2ty

e7na 3auzen ne3ml a
hena aslun?

a **probability** bta3t en el
document de tkon related
le class mo3ayn given el
document.

example: sentiment
analysis, m3ak post, w
3auz te3rf hwa eh el
probability en el post da
ykon el classification
bta3to eno happy post
masln.

we have $P(d|c)$ but we
need to get $P(c|d)$.

Naive Bayes Classifiers

- Naive Bayes is a **probabilistic classifier**, meaning that for a document d , out of all classes $c \in C$, the classifier returns the class \hat{c} which has the **maximum posterior probability** given the document.

$\hat{c} = \operatorname{argmax}_{c \in C} P(c|d)$ max index

$P(x|y) = \frac{P(y|x)P(x)}{P(y)}$

$\hat{c} = \operatorname{argmax}_{c \in C} P(c|d) = \operatorname{argmax}_{c \in C} \frac{P(d|c)P(c)}{P(d)}$

$\hat{c} = \operatorname{argmax}_{c \in C} P(d|c) P(c)$

VEC

3ndk maslan, happy, sad, Neutral.

fanta m3ak el doc, bytl3lk hwa baa probability en el doc de tkon tb3 class mo3yn.

enta m3ak $P(d|c)$ w menha btb2a 3auz te7sb $P(c|d)$ ne3ml a teftkr? ;)

This is possible because we will be computing $\frac{P(d|c)P(c)}{P(d)}$ for each possible class. But $P(d)$ doesn't change for each class → we are always asking about the most likely class for the same document d , which must have the same probability $P(d)$.

Naive Bayes Classifiers

- We can represent a document d as a set of features f_1, f_2, \dots, f_n :

$$\hat{c} = \underset{c \in C}{\operatorname{argmax}} \underbrace{P(f_1, f_2, \dots, f_n | c)}_{\text{likelihood}} \underbrace{P(c)}_{\text{prior}}$$

- This is still **too hard to compute directly**: without some simplifying assumptions, estimating the probability of every possible combination of features (for example, every possible set of words and positions) would require huge numbers of parameters and impossibly large training sets.

- Naive Bayes classifiers therefore make two simplifying assumptions:

- bag of words**: we assume position doesn't matter, and that the word "happy" has the same effect on classification whether it occurs as the 1st, 30th, or last word in the document. Thus, we assume that the features f_1, f_2, \dots, f_n only encode word identity and not position.
 - naive Bayes assumption**: this is the conditional independence assumption that the probabilities $P(f_i | c)$ are independent given the class c and hence can be 'naively' multiplied as follows:
$$P(f_1, f_2, \dots, f_n | c) = P(f_1 | c) \cdot P(f_2 | c) \cdot \dots \cdot P(f_n | c)$$

Naive Bayes Classifiers

- The final equation for the class chosen by a naive Bayes classifier is thus:

el class elly ha5taro hwa akbur result mn
el equation bta3t el mutliplication.

el equation de httnfz 3la kol el classes
elly 3ndy, w b3den n3ml MLE.

$$c_{NB} = \operatorname{argmax}_{c \in C} P(c) \prod_{f \in F} P(f|c)$$

es2l el doctor 3la 7war el word positions
da.

- To apply the naive Bayes classifier to text, we need to consider **word positions**,
by simply walking an index through every word position in the document:

positions \leftarrow all word positions in test document

$$c_{NB} = \operatorname{argmax}_{c \in C} P(c) \prod_{i \in \text{positions}} P(w_i|c)$$

el position hena asdha beha
lw 3ndk keda:
ahmed likes amira
1 2 3
fa by2olak y3ny lef 3la kol el
kalemat w e7sblhom el
equation de.

- Naive Bayes calculations are done in **log space**, to avoid underflow and increase speed.

$$c_{NB} = \operatorname{argmax}_{c \in C} \log P(c) + \sum_{i \in \text{positions}} \log P(w_i|c)$$

linear 34an el features bta3tna ba2t hya el log
fa ehna hena bn3ml summation, fa keda el space bta3na baa linear.

By considering features in **log space**:
the predicted class is **computed** as a
linear function of input features.
Naive Bayes and also Logistic
regression are **linear classifiers**.

Training the Naive Bayes Classifier

- For the class prior $P(c)$, we ask what percentage of the documents in our training set are in each class c .

✓ we have 10 docs

3->happy

4->Neutral

3->sad

$p(\text{Neutral}) = ?$

$$\hat{P}(c) = \frac{N_c}{N_{\text{doc}}}$$

N_c : the number of documents in our training data with class c
 N_{doc} : the total number of documents

- To learn the probability $P(f_i|c)$, we'll assume a feature is just the existence of a word in the document's bag of words, and so $P(w_i|c)$ is computed as the fraction of times the word w_i appears among all words in all documents of topic c .

doc1-> hhhhlma

doc2-> hhamk

doc3-> abcd -> c2

count(h) =

$$\hat{P}(w_i|c) = \frac{\text{count}(w_i, c)}{\sum_{w \in V} \text{count}(w, c)}$$

khud balak dol kol el kalemat, 7ata lw kano
duplicates, y3ny lw l2et 7agat metkrra kaman,
ha3dha brdo

mttla5batsh !!!

- Laplace smoothing is commonly used in naive Bayes text categorization:

$$\hat{P}(w_i|c) = \frac{\text{count}(w_i, c) + 1}{\sum_{w \in V} (\text{count}(w, c) + 1)} = \frac{\text{count}(w_i, c) + 1}{(\sum_{w \in V} \text{count}(w, c)) + |V|}$$

V is the total number of unique words in all the documents of the training set

el V msh bna5ud fe e3tbarna el duplicates, bna5ud el unique words elly fe kol el classes.

Naive Bayes Classifier Example

34an a3rf el

	Cat	Documents
Training	-	just plain boring entirely predictable and lacks energy no surprises and very few laughs
Test	?	very powerful the most fun film of the summer predictable with no fun

$$P(-) = \frac{3}{5} \quad P(+) = \frac{2}{5}$$

$$P(\text{"predictable"}|-) = \frac{1+1}{14+20}$$

$$P(\text{"no"}|-) = \frac{1+1}{14+20}$$

$$P(\text{"fun"}|-) = \frac{0+1}{14+20}$$

$$P(\text{"predictable"}|+) = \frac{0+1}{9+20}$$

$$P(\text{"no"}|+) = \frac{0+1}{9+20}$$

$$P(\text{"fun"}|+) = \frac{1+1}{9+20}$$

P(with) = 0, hal hnrbha fel m3ana? Iaa bn3ml neglect lel unknown words w khas.

- For the test sentence $S = \text{"predictable with no fun"}$, after removing the unknown word 'with', the chosen class is therefore computed as follows:

$$P(-|S) = P(-)P(S|-) = \frac{3}{5} \times \frac{2 \times 2 \times 1}{34^3} = 6.1 \times 10^{-5}$$

$$P(+|S) = P(+)P(S|+) = \frac{2}{5} \times \frac{1 \times 1 \times 2}{29^3} = 3.2 \times 10^{-5}$$

The model thus predicts the class **negative** for the test sentence

Optimizing for Sentiment Analysis

Binary Naive Bayes: for sentiment classification whether a word occurs or not seems to matter more than its frequency. Thus, it often improves performance to clip the word counts in each document at 1.

Four original documents:

- it was pathetic the worst part was the boxing scenes
- no plot twists or great scenes
- + and satire and great plot twists
- + great scenes great film

After per-document binarization:

- it was pathetic the worst part boxing scenes
- no plot twists or great scenes
- + and satire great plot twists
- + great scenes film

	NB Counts		Binary Counts	
	+	-	+	-
and	2	0	1	0
boxing	0	1	0	1
film	1	0	1	0
great	3	1	2	1
it	0	1	0	1
no	0	1	0	1
or	0	1	0	1
part	0	1	0	1
pathetic	0	1	0	1
plot	1	1	1	1
satire	1	0	1	0
scenes	1	2	1	2
the	0	2	0	1
twists	1	1	1	1
was	0	2	0	1
worst	0	1	0	1

Optimizing for Sentiment Analysis

Dealing with Negation:

- Consider the difference between *I really like this movie (positive)* and *I didn't like this movie (negative)*.
→ The negation expressed alters the inferences we draw from the predicate *like*.
- Similarly, negation can modify a negative word to produce a positive review (*don't dismiss this film, doesn't let us get bored*).
- During text normalization, prepend the prefix **NOT** to every word after a token of logical negation (*n't, not, no, never*) until the next punctuation mark.

didn't like this movie , but I
becomes

didn't NOT_like NOT_this NOT_movie , but I

tb3n msh de el tre2a el wa7eda
w de msh aktur haga effictive,
bs da msh m3nah en el tor2 el
simple msh btgeb accuracy.

if you are concerned more, read
more about these topics.

- Newly formed ‘words’ like *NOT like*, *NOT recommend* will thus occur more often in negative document and act as cues for negative sentiment, while words like *NOT bored*, *NOT dismiss* will acquire positive associations.

Optimizing for Sentiment Analysis

Sentiment Lexicons: lists of words that are pre-annotated with positive or negative sentiment.

- + : *admirable, beautiful, confident, dazzling, ecstatic, favor, glee, great*
- : *awful, bad, bias, catastrophe, cheat, deny, envious, foul, harsh, hate*

A common way to use lexicons in a naive Bayes classifier is to add a feature that is counted whenever a word from that lexicon occurs.

Evaluation

lazm ay 7aga fl denya te3mlha, lazm te3mlha evaluation.
w ezay te2dr enk tsl7 el model bta3k.

- A **confusion matrix** for visualizing how well a binary classification system performs against **gold standard labels**.

shof el gold fen
w 7ot odama el
recall

shof el system
output fen w 7ot
odamo el percisoin.

		gold standard labels		percision-> kolohom nfs el 7rf
		gold positive	gold negative	
system output labels	system positive	true positive	false positive	
	system negative	false negative	true negative	
recall	$\frac{tp}{tp+fn}$	$recall = \frac{tn}{tn+fp}$	$precision = \frac{tp}{tp+fp}$	$accuracy = \frac{tp+tn}{tp+fp+tn+fn}$

- F1-score:** is the most frequently used metric that incorporates aspects of both precision and recall in a single metric. (F “1” since P and R are treated as equally important)

$$F_1 = \frac{2PR}{P+R}$$

harmonic mean -> search about it.

look for F-Score it is not same as the F1-Score.

efham 34an lw 7fzt el matrix
de w sha2lebna el matrix
hanz3l kolena.

false negative -> el model
bta3ak mtl3 negative
w de false
yeb2a el s7 enha ttl3
positive.

e7na bn2ol false aw true 3la
el 7aga elly el model
mtl3ha.

Evaluation

w2ft 3nd en el accuracy lw7dha matkfeesh.

- Although accuracy might seem a natural metric, we generally don't use it for text classification tasks
→ That's because accuracy doesn't work well when the classes are unbalanced e.g., we are discovering something that is rare.

Example:

We have a million tweet, only 100 of them are discussing their love (or hatred) for our pie “about pie”, while the other 999,900 are about something completely unrelated “not about pie”.

- Imagine a simple classifier that stupidly classified every tweet as “not about pie”. This classifier would have 999,900 true negatives and only 100 false negatives. The accuracy for this model is very high, at $999,900/1,000,000 = 99.9\%$. → But of course this fabulous ‘no pie’ classifier would be completely useless, since it wouldn't find a single one of the customer comments we are looking for.
- Solution: other metrics “recall” and “precision”. This classifier, despite having a fabulous accuracy of 99.99%, has a terrible recall/precision of 0 (since there are no true positives).

Precision is a good measure, when the cost of False Positive is high.
Recall is a good measure, when the cost of False Negative is high.

Evaluation

- Confusion matrix for a **three-class** classification task.

el tre2a de htl5btk! fokak mnha

h2olak el tre2a el ana
b3mlha, w b3den enta
e3mlha bl tre2a el try7k

ersm el 2x2 matrix.
2o2af 3nd kol cell es2l
nfsk
men elly el model by2ol
enha x w hya y msln.
7elha mra w etnen htt3wd.

		gold labels		
		urgent	normal	spam
system output	urgent	8	10	1
	normal	5	60	50
	spam	3	30	200
recall		$\text{recall}_{\text{u}} = \frac{8}{8+5+3}$	$\text{recall}_{\text{n}} = \frac{60}{10+60+30}$	$\text{recall}_{\text{s}} = \frac{200}{1+50+200}$
precision		$\text{precision}_{\text{u}} = \frac{8}{8+10+1}$	$\text{precision}_{\text{n}} = \frac{60}{5+60+50}$	$\text{precision}_{\text{s}} = \frac{200}{3+30+200}$
DVR		8	11	340

recall = cell / summation of col (in this case only)
if we rotated the matrix it will be the row.... mttl5btsh arguk.

percision nfs el klam....

- In the figure, it shown how to compute a **distinct** precision and recall value for each class.
- In order to derive a **single metric** that tells us how well the **system** is doing, we can combine these values in two ways:
 - Macroaveraging**, we compute the performance for each class, and then **average over classes**.
 - Microaveraging**, we collect the **decisions** for all classes into a single confusion matrix, and then compute **precision and recall** from that table.

Evaluation

Class 1: Urgent		Class 2: Normal		Class 3: Spam		Pooled	
true	true	true	true	true	true	true	true
urgent	not	normal	not	spam	not	yes	no
system urgent	8	11	60	55	200	33	268
system not	8	340	40	212	51	83	99

precision = $\frac{8}{8+11} = .42$

precision = $\frac{60}{60+55} = .52$

precision = $\frac{200}{200+33} = .86$

macroaverage precision = $\frac{.42+.52+.86}{3} = .60$

microaverage precision = $\frac{268}{268+99} = .73$ Yes

- *Microaverage* is dominated by the *more frequent class* (in this case spam), since the counts are pooled.
- *Macroaverage* better reflects the statistics of the *smaller classes*, and so is more appropriate when performance on all the classes is *equally important*.



Thank You

0100 0100

0110 0110

1001 1001

0100 0110

0100 001

1101 0101

1010
0110

1001
0101

0100
1101

1101
0101

0110
1001

0101
0101

0100
0011

0110

1001