



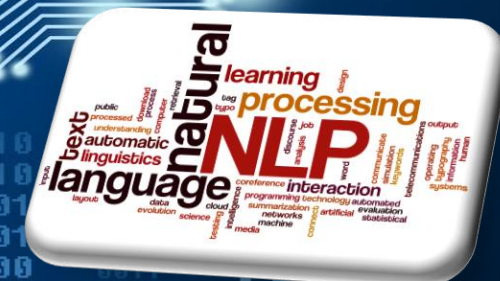
Cairo University

CAIRO UNIVERSITY

Cairo University  
Faculty of Engineering  
Computer Engineering Department

# Natural Language Processing

Dr. Sandra Wahid



# Prediction

- Predict next few words someone is going to say? What word, for example, is likely to follow

“Please turn your homework ...”



in/over but not refrigerator/the

- We will introduce models that assign a **probability** to each possible **next word**.
- The same models will also serve to assign a probability to an **entire sentence**.



# Prediction

Why predict upcoming words, or assign probabilities to sentences?

- Speech Recognition: identify words in noisy, ambiguous input
    - Ex: “*I will be back soon*” is more probable than “*I will be back tone*”
  - Spelling Correction/Grammatical Error Correction:
    - Ex: “*Their are two midterms*” *There* was mistyped as *Their*. The phrase *There are* will be much more probable than *Their are*
  - Machine Translation:
    - Ex: Suppose we are translating a Chinese source sentence to English:
    - The following set consists of potential rough English translations:

he introduced reporters to the main contents of the statement  
he briefed to reporters the main contents of the statement  
**he briefed reporters on the main contents of the statement**
- 他 向 记者 介绍了 主要 内容  
He to reporters introduced main content
- A probabilistic model of word sequences could suggest that “*briefed reporters on*” is a more probable English phrase than “*briefed to reporters*” (which has an awkward *to* after *briefed*) or “*introduced reporters to*” (which uses a verb that is less fluent English in this context), allowing us to correctly select the boldfaced sentence above.

# Language Models (LMs)

- Models that assign probabilities to sequences of words are called language models or LMs, examples: **N-Gram**, **RNN** LMs.
- An **n-gram** is a sequence of **n words**:
  - a 2-gram (**bigram**) is a *two-word* sequence of words like “please turn”, “turn your”, or “your homework”
  - a 3-gram (**trigram**) is a *three-word* sequence of words like “please turn your”, or “turn your homework”.

## Goals:

- estimate the **probability of the last word** of an n-gram given the previous words
- assign **probabilities to entire sequences**

# N-Gram

- Computing  $P(w/h)$ , the probability of a word  $w$  given some history  $h$ .
  - Example: the history  $h$  is “*its water is so transparent that*” and we want to know the probability that the next word is *the*:  $P(\text{the}|\text{its water is so transparent that})$ .
- One way to estimate this probability is from **relative frequency counts** using very large corpus.
  - “Out of the times we saw the history  $h$ , how many times was it followed by the word  $w$ ”

$$P(\text{the}|\text{its water is so transparent that}) = \frac{C(\text{its water is so transparent that the})}{C(\text{its water is so transparent that})}$$

## Problems with this scheme:

- Even the web **isn't big enough** to give us good estimates in most cases.
    - Some cases can have **zero counts**.
- This is because language is creative.



# N-Gram

- Computing probabilities of entire sequences like  $P(w_1; w_2; \dots; w_n) \rightarrow$  use **chain rule of probability**

$$\begin{aligned} P(X_1 \dots X_n) &= P(X_1)P(X_2|X_1)P(X_3|X_{1:2}) \dots P(X_n|X_{1:n-1}) \\ &= \prod_{k=1}^n P(X_k|X_{1:k-1}) \end{aligned}$$

Applying the chain rule to words, we get

$$\begin{aligned} P(w_{1:n}) &= P(w_1)P(w_2|w_1)P(w_3|w_{1:2}) \dots P(w_n|w_{1:n-1}) \\ &= \prod_{k=1}^n P(w_k|w_{1:k-1}) \end{aligned}$$

But using the chain rule doesn't really seem to help us!  
We don't know any way to compute the exact probability of a word  
given a long sequence of preceding words  
 $\rightarrow$  we can **approximate** the history by just the last few words.

# Bigram Model

- Approximates the probability of a word given all the previous words by using only the conditional probability of the **preceding word**.

$$P(w_n|w_{1:n-1}) \approx P(w_n|w_{n-1})$$

- The assumption that the probability of a word depends only on the previous word is called a **Markov assumption**.
- Generalization of the bigram:

$$P(w_n|w_{1:n-1}) \approx P(w_n|w_{n-N+1:n-1})$$


- where  $N=2 \rightarrow$  bigram,  $N=3 \rightarrow$  trigram,

# Maximum Likelihood Estimation (MLE)

Chain Rule



Markov  
Assumption



$$P(w_{1:n}) \approx \prod_{k=1}^n P(w_k | w_{k-1})$$

- How do we estimate these bigram or n-gram probabilities?
  - An intuitive way to estimate probabilities is MLE by getting counts from a corpus, and **normalizing** the counts so that they lie between 0 and 1.
  - Count of the bigram  $C(xy)$  and normalize by the sum of all the bigrams that share the same first word  $x$ :
- We can simplify this equation, since the sum of all bigram counts that start with a given word  $w_{n-1}$  must be equal to the unigram count for that word  $w_{n-1}$

$$P(w_n | w_{n-1}) = \frac{C(w_{n-1}w_n)}{\sum_w C(w_{n-1}w)}$$

$$P(w_n | w_{n-1}) = \frac{C(w_{n-1}w_n)}{C(w_{n-1})}$$

For the general case of MLE n-gram parameter estimation:

$$P(w_n | w_{n-N+1:n-1}) = \frac{C(w_{n-N+1:n-1} w_n)}{C(w_{n-N+1:n-1})}$$



# Some Practical Issues

- In practice it's more common to use trigram models, which condition on the previous two words rather than the previous word, or 4-gram or even 5-gram models, when there is sufficient training data.
- Represent and compute language model probabilities in log format as **log probabilities**.
  - Since probabilities are less than or equal to 1, the more probabilities we multiply together, the smaller the product becomes → may result in **numerical underflow**.
  - Adding in log space is equivalent to multiplying in linear space.
  - To convert back into probabilities if we need to report them at the end; then we can just take the exp of the logprob.

$$p_1 \times p_2 \times p_3 \times p_4 = \exp(\log p_1 + \log p_2 + \log p_3 + \log p_4)$$

# Evaluating Language Models

- Does our language model prefer good sentences to bad ones?  
→ Assign higher probability to “real” or “frequently observed” sentences than “ungrammatical” or “rarely observed” sentences.

## Performance evaluation of a language model

### Extrinsic Evaluation:

embed it in an application and measure how much the **application** improves → running big NLP systems end-to-end is often very expensive

### Intrinsic Evaluation:

measures the quality of a model **independent** of any application → need a **test set (held out corpora)** that are not part of the training set

# Evaluating Language Models: Perplexity

- In practice, we don't use raw probability as our metric for evaluating language models but a variant called *perplexity* (sometimes called *PP* for short).
- PP of a language model on a test set is the **inverse probability** of the test set, normalized by its number of words.

- For a test set  $W = w_1 w_2 \dots w_N$ :

$$\begin{aligned} \text{PP}(W) &= P(w_1 w_2 \dots w_N)^{-\frac{1}{N}} \\ &= \sqrt[N]{\frac{1}{P(w_1 w_2 \dots w_N)}} \end{aligned}$$

- Using chain rule:

$$\text{PP}(W) = \sqrt[N]{\prod_{i=1}^N \frac{1}{P(w_i | w_1 \dots w_{i-1})}}$$

- For bigram model:

$$\text{PP}(W) = \sqrt[N]{\prod_{i=1}^N \frac{1}{P(w_i | w_{i-1})}}$$



# Evaluating Language Models: Perplexity

**Minimizing perplexity** is equivalent to **maximizing** the test set probability.  
If a model assigns a high probability to the test set, it means that it is **not surprised** to see it (it's **not perplexed** by it).

- The entire sequence of words in some test set will cross many sentence boundaries  
→ we need to include the begin- and end-sentence markers <s> and </s> in the probability computation.
- Example of how perplexity can be used to compare different n-gram models.

	Unigram	Bigram	Trigram
Perplexity	962	170	109

- An (intrinsic) improvement in perplexity **does not guarantee** an (extrinsic) improvement in the performance of a language processing task. Perplexity is commonly used as a quick check on an algorithm. But a model's improvement in perplexity should always be confirmed by an end-to-end evaluation of a real task before concluding the evaluation of the model.

# Note

- If there is **no overlap** in the generated sentences between the training set and the test set, then the model is **useless**.

Use a training corpus that has a similar **genre** to whatever task we are trying to accomplish, e.g., to build a language model for a question-answering system, we need a training corpus of questions.

It is equally important to get training data in the appropriate **dialect**.

# Zeros

- If any corpus is limited, some perfectly acceptable English **word sequences** are subject to be missing from it  
→ we'll have many cases of “**zero probability n-grams**” that should really have some **non-zero probability**.

- Example: the words that follow the bigram *<denied the>* with their counts:

denied the allegations:	5
denied the speculation:	2
denied the rumors:	1
denied the report:	1

But suppose our test set has phrases like:

denied the offer
denied the loan

Our model will incorrectly estimate that the  $P(\text{offer} / \text{denied the})$  is **0**

## Problems:

1. underestimating the probability of all sorts of words that might occur which will hurt the **performance** of any application.
2. if the probability of any word in the test set is 0, the entire **probability of the test set is 0** and we can't compute perplexity at all, since we can't divide by 0.



# Unknown Words

The previous slide discussed the problem of words whose bigram probability is zero. But what about words we simply **have never seen before**?

- We can have a **closed vocabulary** system: where the test set can only contain words from a certain lexicon, and there will be no unknown words.
- Unknown words, or out of vocabulary (OOV) words:
  - The percentage of OOV words that appear in the test set is called the **OOV rate**.
  - An **open vocabulary** system is one in which we model these potential unknown words in the test set by adding a pseudo-word called <UNK>.

# Unknown Words

- Possible Solutions:

1. Turn the problem back into a closed vocabulary one:

1. **Choose a vocabulary** (word list) that is fixed in advance.
2. **Convert** in the training set any word that is not in this set (any OOV word) to the unknown word token <UNK> in a text normalization step.
3. **Estimate** the probabilities for <UNK> from its counts just like any other regular word in the training set.

2. Create such a vocabulary implicitly:

by replacing words in the training data by <UNK> based on their frequency. For example, we can replace by <UNK> all words that occur fewer than  $n$  times in the training set, where  $n$  is some small number.

# Smoothing

- What do we do with words that are in our vocabulary (they are **not unknown words**) but appear in a test set in an **unseen context** (for example they appear after a word they never appeared after in training)?
- Zeros problem solution: to keep a language model from assigning zero probability to these unseen events, we'll have to shave off a bit of probability mass from some more frequent events and give it to the events we've never seen.



This modification is called smoothing or discounting



# Laplace Smoothing

- Alternate name **add-one** smoothing: since it adds one to each count.
  - All the counts that used to be zero will now have a count of 1

- Unigram probabilities:

$$P_{\text{Laplace}}(w_i) = \frac{c_i + 1}{N + V}$$

N: total number of word tokens  
Need to adjust the denominator to take into account the extra V observations.  
V: number of unique words

- Bigram probabilities:

$$P_{\text{Laplace}}^*(w_n | w_{n-1}) = \frac{C(w_{n-1}w_n) + 1}{\sum_w (C(w_{n-1}w) + 1)} = \frac{C(w_{n-1}w_n) + 1}{C(w_{n-1}) + V}$$

# Add-k Smoothing

- One alternative to add-one smoothing is to move a **bit less** of the probability mass from the seen to the unseen events.
  - Instead of adding 1 to each count, we add a **fractional count k** (.5? .05? .01?).

$$P_{\text{Add-k}}^*(w_n|w_{n-1}) = \frac{C(w_{n-1}w_n) + k}{C(w_{n-1}) + kV}$$

- Requires a method for **choosing k**; this can be done, for example, by optimizing on a development set (devset).

# Backoff and Interpolation

- The **discounting** we have been discussing so far can help solve the problem of **zero frequency n-grams**. But there is an additional source of knowledge we can draw on.
- If we are trying to compute  $P(w_n/w_{n-2}w_{n-1})$  but we have no examples of a particular trigram  $w_nw_{n-2}w_{n-1}$ , we can instead estimate its probability by using the bigram probability  $P(w_n/w_{n-1})$ . Similarly, if we don't have counts to compute  $P(w_n/w_{n-1})$  we can look to the unigram  $P(w_n)$ .
- In other words, sometimes using **less context** is a good thing, helping to **generalize** more for contexts that the model hasn't learned much about.



# Backoff and Interpolation

1. **Backoff:** we use the *trigram* if the evidence is sufficient, otherwise we use the *bigram*, otherwise the *unigram*. → we only “back off” to a lower-order n-gram if we have zero evidence for a higher-order.
2. **Interpolation:** we always mix the probability estimates from all the n-gram estimators, weighing and combining the *trigram*, *bigram*, and *unigram* counts.

Example: simple linear interpolation:

$$\hat{P}(w_n|w_{n-2}w_{n-1}) = \lambda_1 P(w_n|w_{n-2}w_{n-1}) + \lambda_2 P(w_n|w_{n-1}) + \lambda_3 P(w_n)$$

$$\sum_i \lambda_i = 1$$



# Thank You