



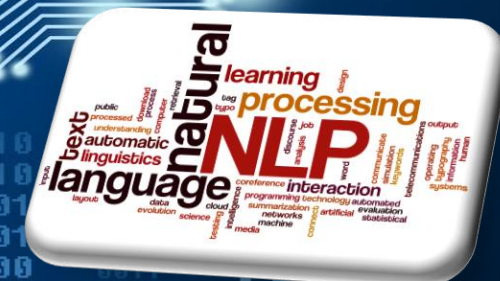
Cairo University

CAIRO UNIVERSITY

Cairo University  
Faculty of Engineering  
Computer Engineering Department

# Natural Language Processing

Dr. Sandra Wahid



# Word Embeddings

- Previously, we saw how to represent a word as a **sparse, long vector** with dimensions corresponding to words in the vocabulary or documents in a collection.
- We now introduce a more powerful word representation: **embeddings, short dense vectors**.
  - With number of dimensions  $d$  ranging from 50-1000.
  - The vectors are dense: instead of vector entries being sparse, mostly-zero counts or functions of counts, the values will be real-valued numbers that can be negative.
- It turns out that dense vectors work better in every NLP task than sparse vectors.
- One method for computing embeddings is Word2vec where embeddings are **static**: meaning that the method learns one fixed embedding for each word in the vocabulary.
- Other methods such as BERT, are for learning **dynamic contextual** embeddings in which the vector for each word is different in different contexts.



# Word2vec

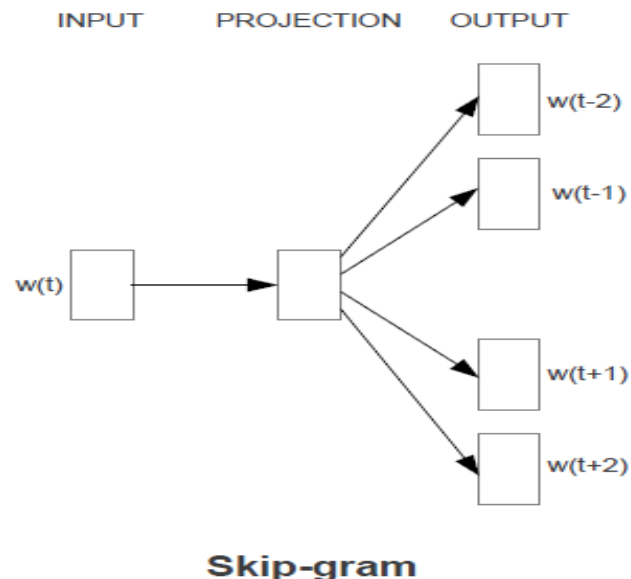
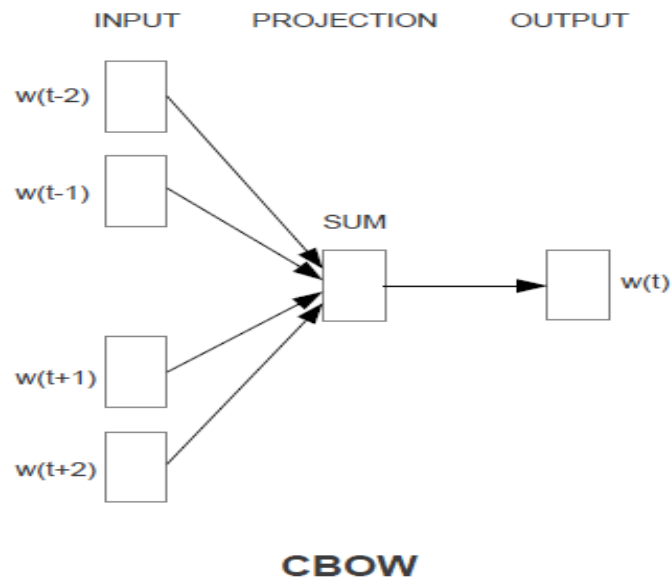
- Word2vec is developed by *Tomas Mikolov et al. at Google* where two models are proposed:

CBOW

- Continuous Bag-Of-Words
- where a word is predicted based on its context (the surrounding words).

SG

- Continuous Skip-Gram
- where the surrounding words are predicted given a current word.



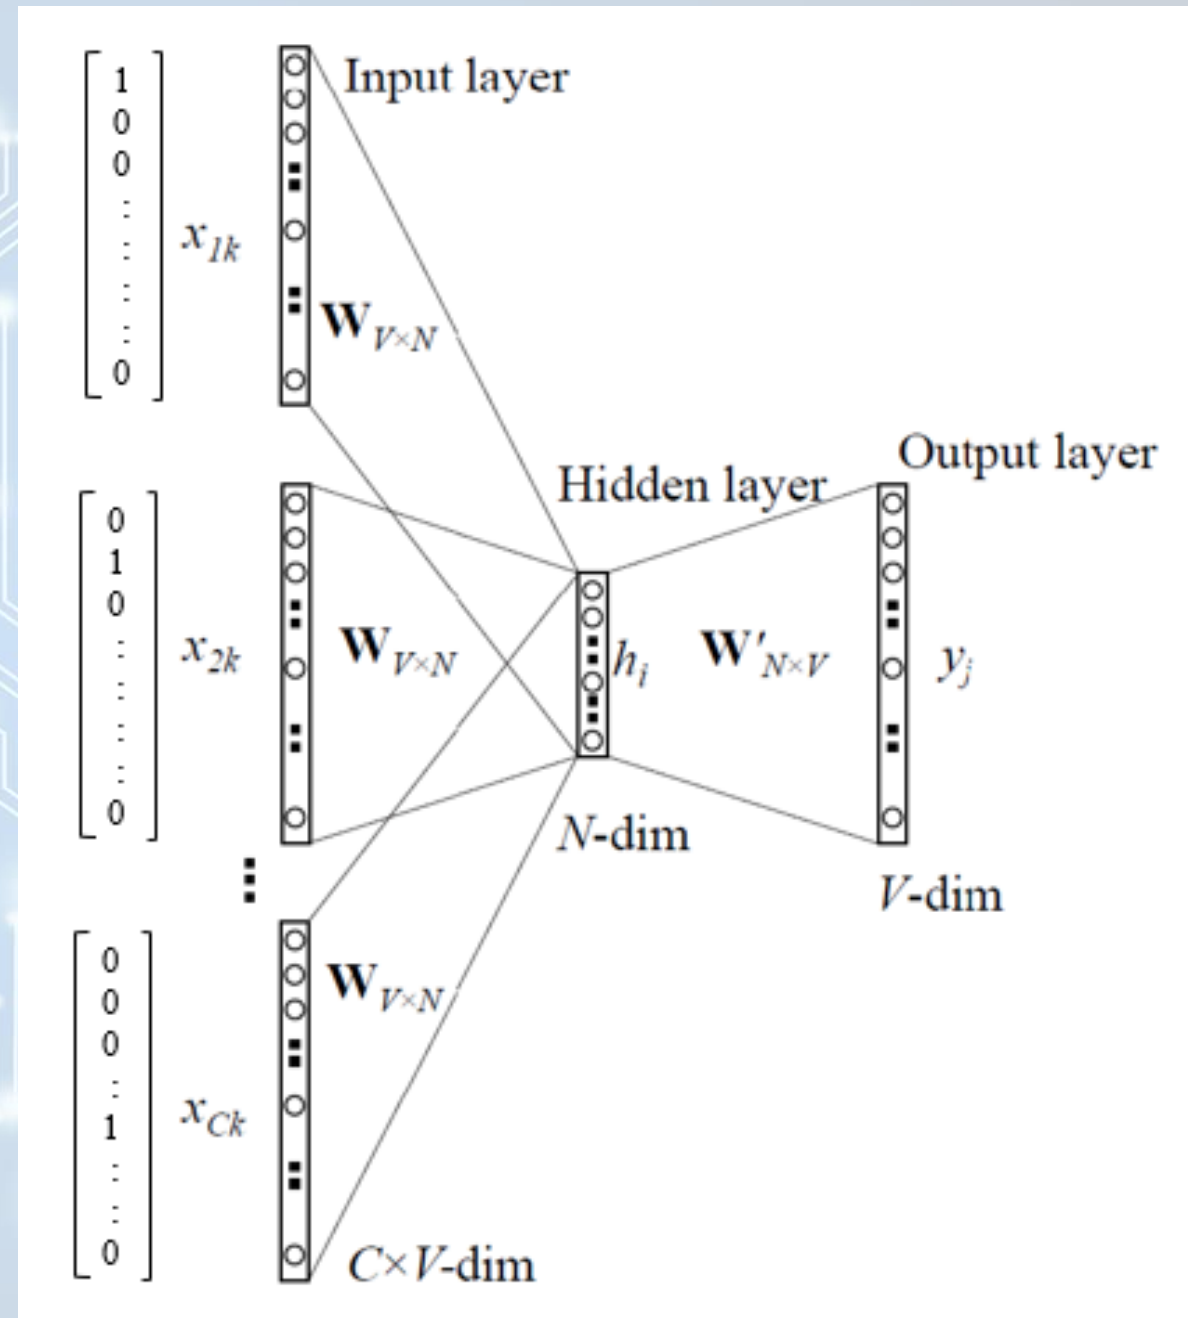
The models' architecture is a neural network with three layers: an input layer, a hidden layer and an output layer. The models are trained using huge text corpora

Though, we won't use the neural network after training!!!

# CBOW

## Input Layer:


- Its size is  $CV$ 
  - $C$ =number of context words
  - $V$ =number of words in the vocabulary
- The inputs are **C context vectors**:
  - Each context vector is one-hot encoded vector of size  $V$  (it is a vector having 1 in the position encoding the word, and 0 otherwise).



# CBOW

## Hidden Layer:


- Its size is N
  - N=number of neurons in the hidden layer  
=the dimension of embeddings  
(it is a hyper parameter)
- Fully-connected (dense) layer whose weights are the word embeddings.



# CBOW

## Output layer:

- Its size is  $V$ 
  - Where the gold output is the one-hot vector encoding the “middle/target word”.
- It is a softmax layer which is designed so that the sum of the probabilities obtained in the output layer equals 1.
  - The network is going to tell us the probability for every word in our vocabulary of being the “target word” for the given input(s) “context word(s)”

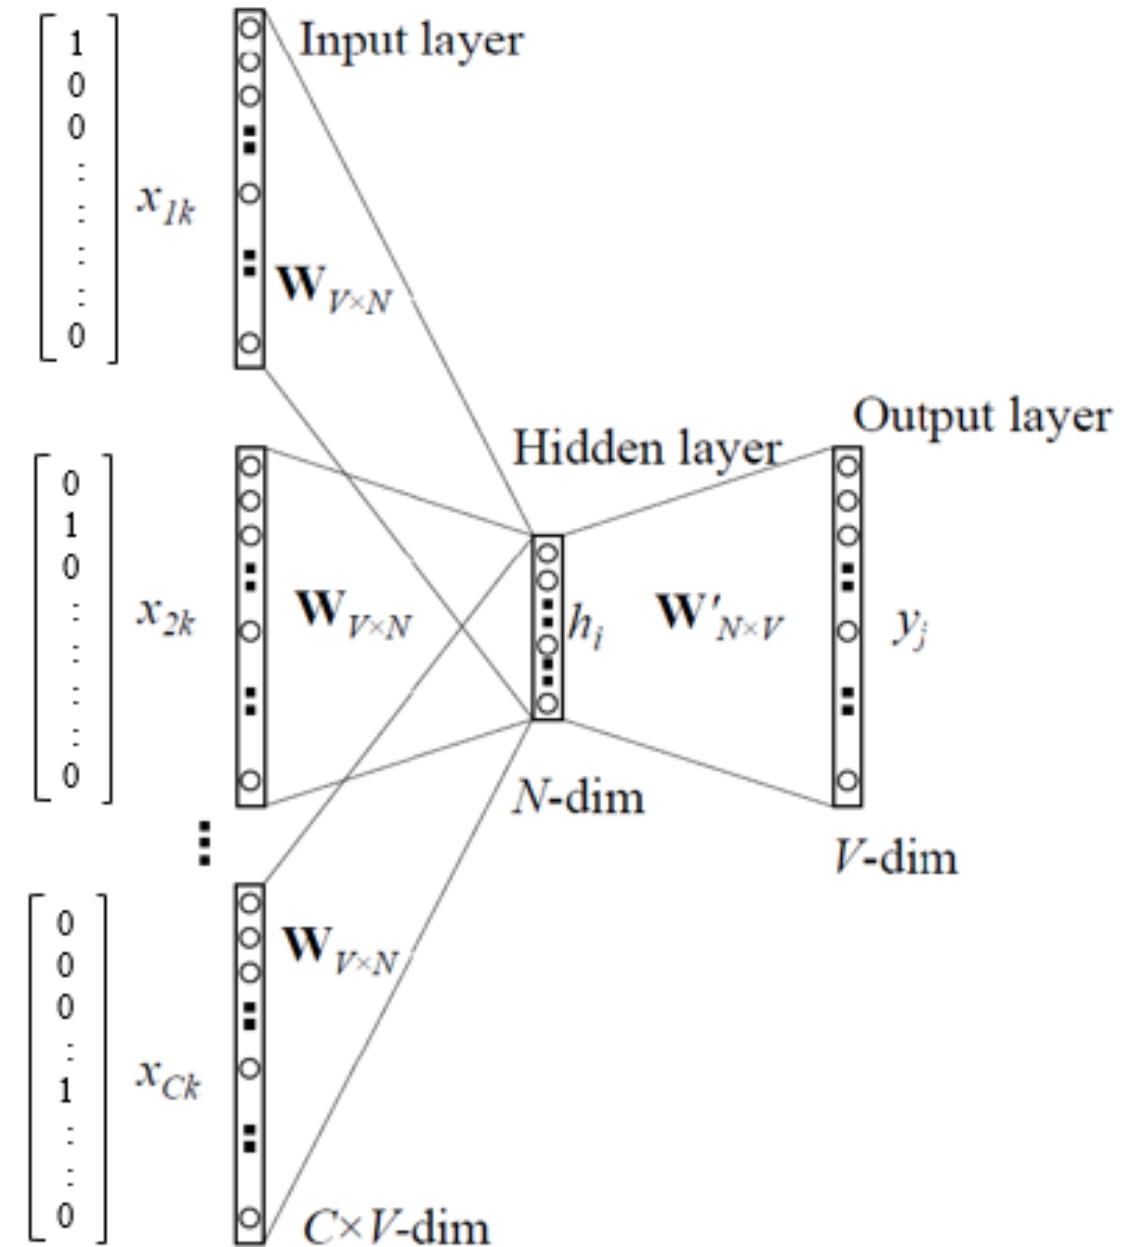




# CBOW

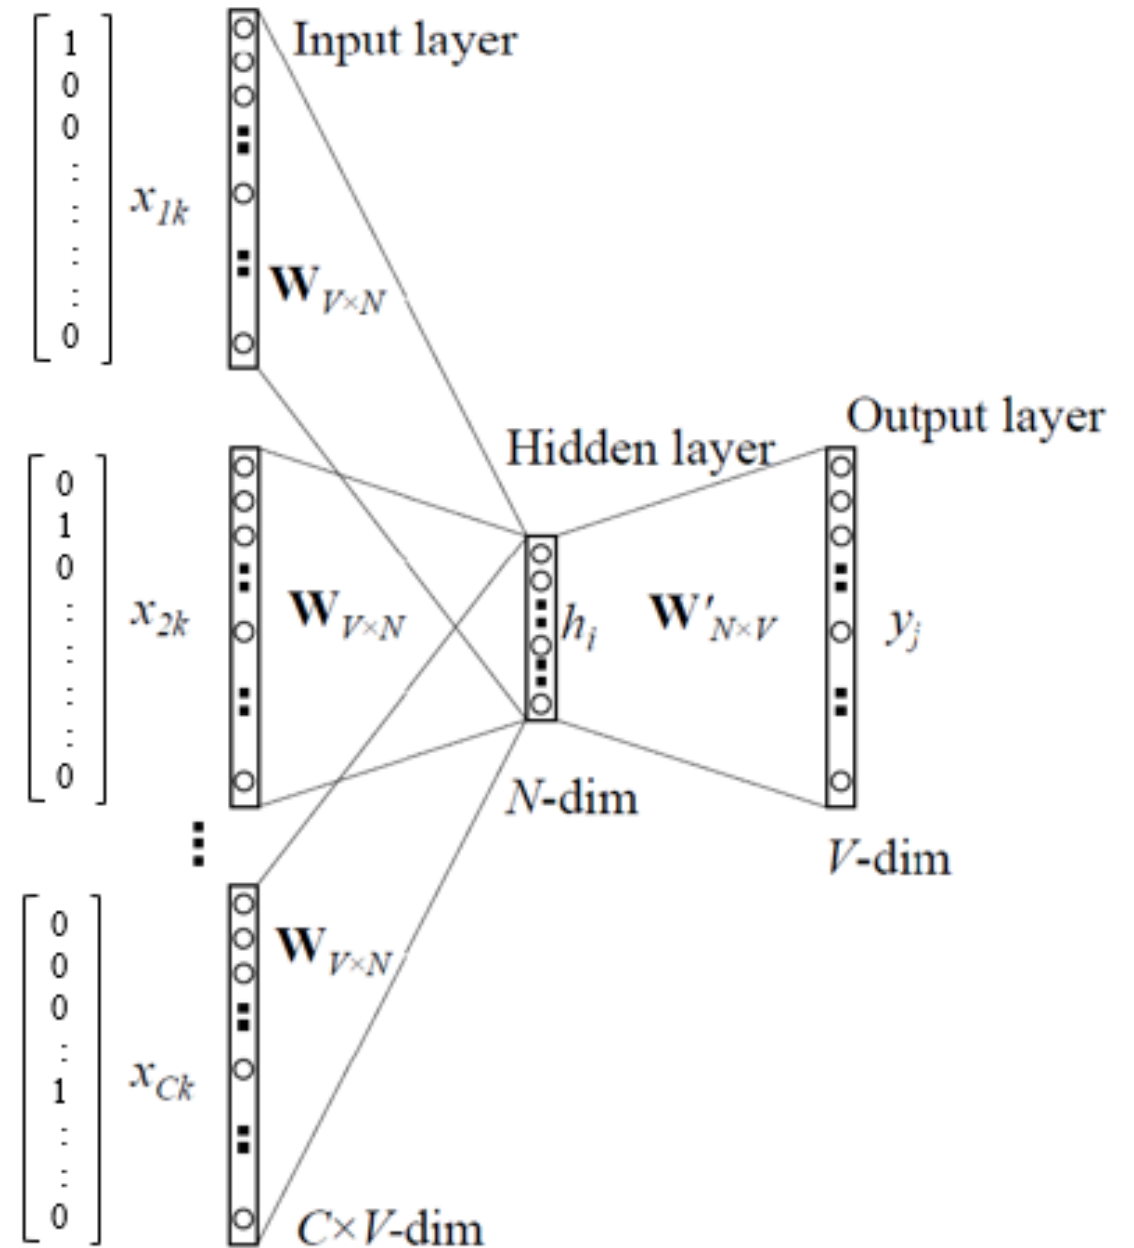
## Weights:

- There are two sets of weights:
  - one is between the input and the hidden layer:  
**Input-Hidden** layer matrix consisting of  $C$   $W[V \times N]$  matrices
  - second is between hidden and output layer:  
**Hidden-Output** layer matrix size  $= [N \times V]$



# CBOW

- There is no activation function between any layers (linear activation):
  - Linear network except for the softmax at the output layer.
- The input-to-hidden-layer weight matrix is multiplied with the input vector to produce the hidden layer output.
- The hidden input gets multiplied by hidden-output weights to produce the output then pass it to softmax.
- Error between output and target is calculated and back propagated to re-adjust the weights.
- The weights between the input layer and the hidden layer are taken as the word vectors after training.





# CBOW

- Let the available corpus be this one sentence *“word vectors are widely used”*.
- The unique words in the vocabulary are 5 words ( $V = 5$ ) as follows:  
[“word”, “vectors”, “are”, “widely”, “used”].
- Training data for  $C=1$  (context words)

training sample	context word	target word
1	word	vectors
2	vectors	are
3	are	widely
4	widely	used

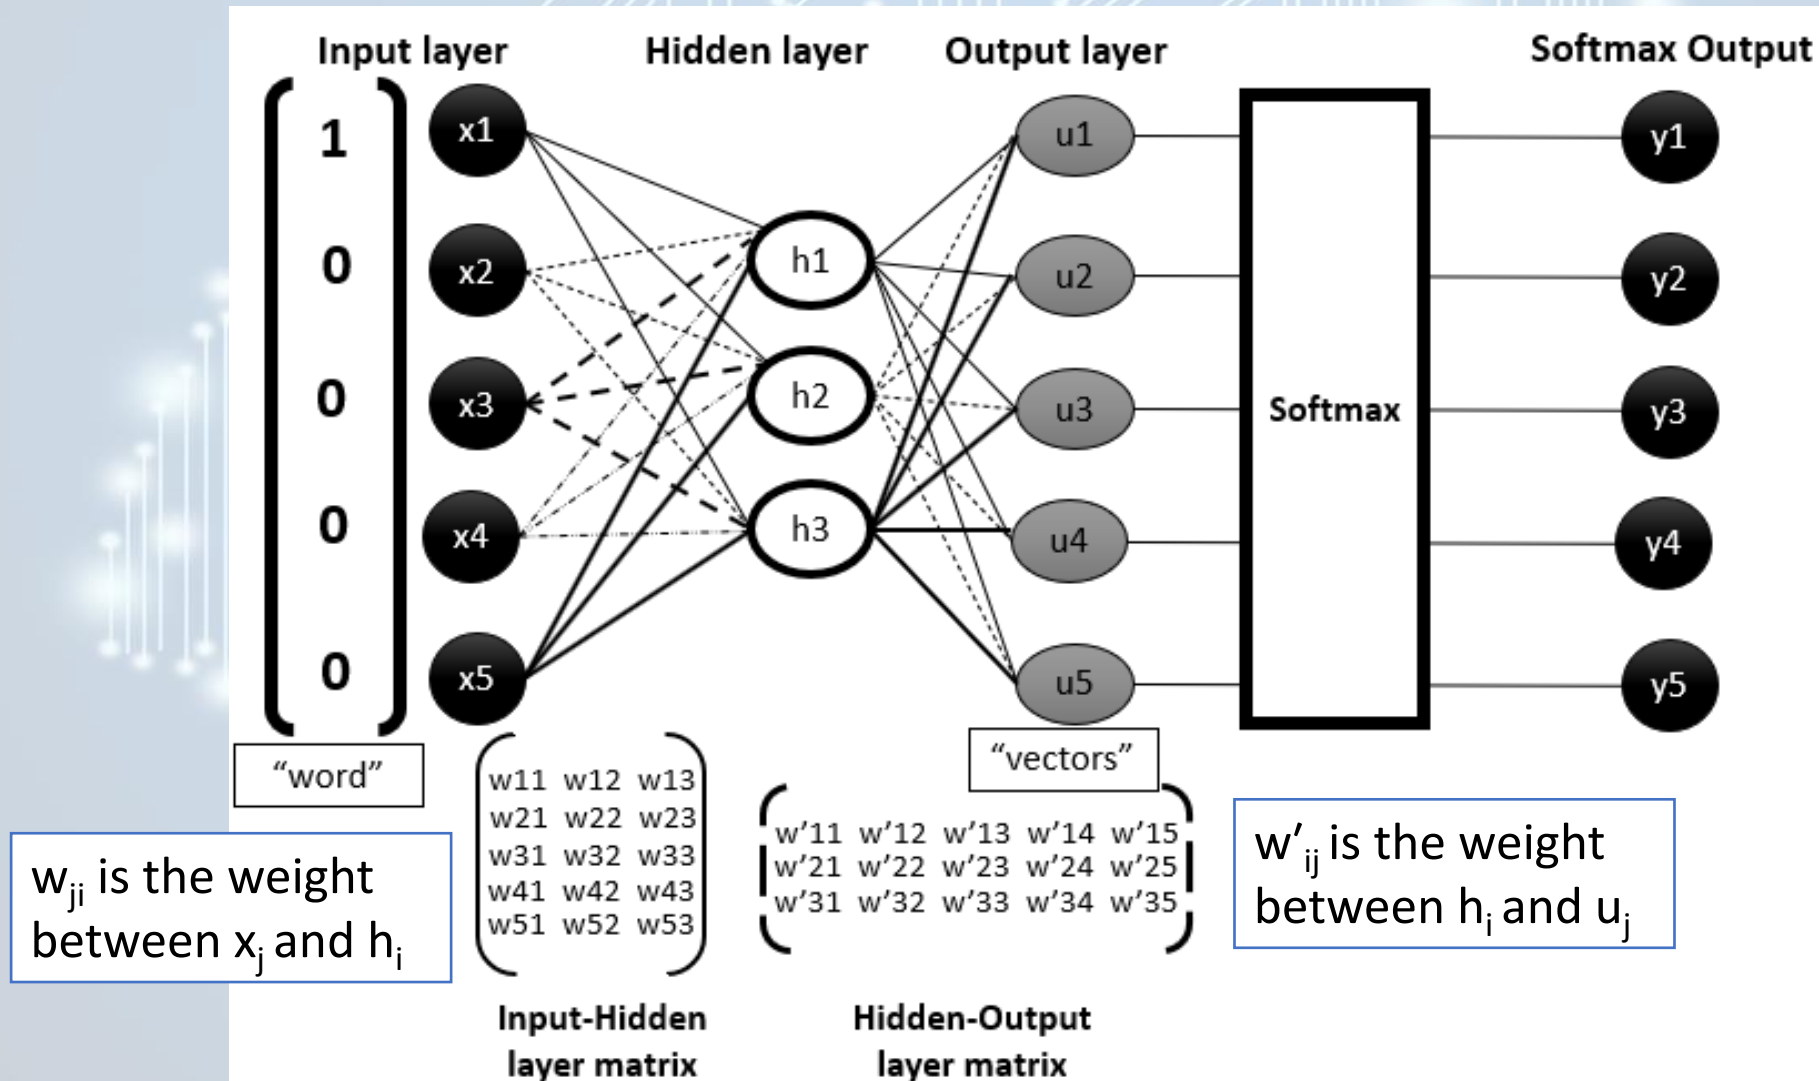
- Encoded training data:

training sample	context word	target word
1	[1,0,0,0,0]	[0,1,0,0,0]
2	[0,1,0,0,0]	[0,0,1,0,0]
3	[0,0,1,0,0]	[0,0,0,1,0]
4	[0,0,0,1,0]	[0,0,0,0,1]

# CBOW

- Let  $N = 3$ , consider the first training sample:

training sample	context word	target word
1	[1,0,0,0,0]	[0,1,0,0,0]



# CBOW


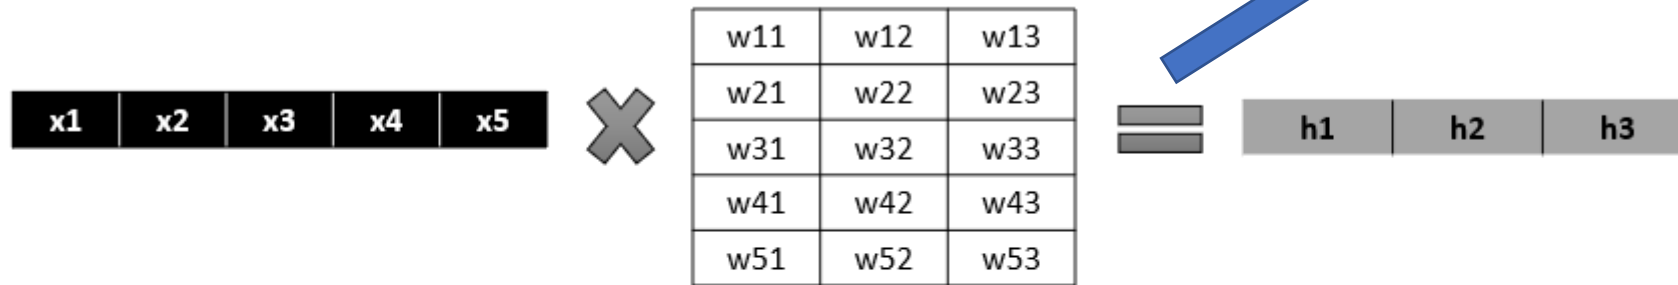
- Equations:

$$h_i = \sum_j w_{ji} * x_j \text{ where } i = 1, 2, 3 \text{ and } j = 1, 2, 3, 4, 5 \quad (1)$$

$$u_j = \sum_i w'_{ij} * h_i \text{ where } i = 1, 2, 3 \text{ and } j = 1, 2, 3, 4, 5 \quad (2)$$

$$y_j = \text{Softmax}(u_j) \text{ where } j = 1, 2, 3, 4, 5 \quad (3)$$

- Matrix Forms:



$$\begin{bmatrix} 0 & 0 & 0 & 1 & 0 \end{bmatrix} \times \begin{bmatrix} 17 & 24 & 1 \\ 23 & 5 & 7 \\ 4 & 6 & 13 \\ 10 & 12 & 19 \\ 11 & 18 & 25 \end{bmatrix} = \begin{bmatrix} 10 & 12 & 19 \end{bmatrix}$$

The output of the hidden layer is just the “word vector” for the input word.



# CBOW

- The error calculations are performed given the target vector:  $[0,1,0,0,0]$  through **back propagation** where the weight values ( $w$  and  $w'$ ) are updated.
- The training continues for **many iterations** until the weights are adjusted for all the training samples.
- Finally, the vectors of the words in the vocabulary are weights in the Input-Hidden layer matrix such that each row in the matrix is the word vector representation of the equivalent word.

The word in the vocabulary	The equivalent word vector
word	$[w_{11} \ w_{12} \ w_{13}]$
vectors	$[w_{21} \ w_{22} \ w_{23}]$
are	$[w_{31} \ w_{32} \ w_{33}]$
widely	$[w_{41} \ w_{42} \ w_{43}]$
used	$[w_{51} \ w_{52} \ w_{53}]$

# CBOW


- For  $C > 1$  (context words), the only difference will be in handling the  $C$  input vectors.
  - Each vector will be multiplied with the input-hidden layer matrix  $W[V \times N]$  returning  $C$   $[1 \times N]$  vectors.
  - Finally, to obtain a single vector  $\rightarrow$  all  $C$   $[1 \times N]$  vectors will be averaged element-wise.
- This computation is equivalent to initially computing a **mean vector** for the  $C$  input one-hot encoded vectors then proceeding as in the case of  $C = 1$ .
- For  $C=2 \rightarrow$  window size=1

training sample	context word	target word
1	(word,are)	vectors
2	(vectors,widely)	are
3	(are,used)	widely
4	widely	used

training sample	context words	mean context word	target word
1	$([1,0,0,0,0],[0,0,1,0,0])$	$[0.5,0,0.5,0,0]$	$[0,1,0,0,0]$
2	$([0,1,0,0,0],[0,0,0,1,0])$	$[0,0.5,0,0.5,0]$	$[0,0,1,0,0]$
3	$([0,0,1,0,0],[0,0,0,0,1])$	$[0,0,0.5,0,0.5]$	$[0,0,0,1,0]$
4	$[0,0,0,1,0]$	$[0,0,0,1,0]$	$[0,0,0,0,1]$

## Input Layer:

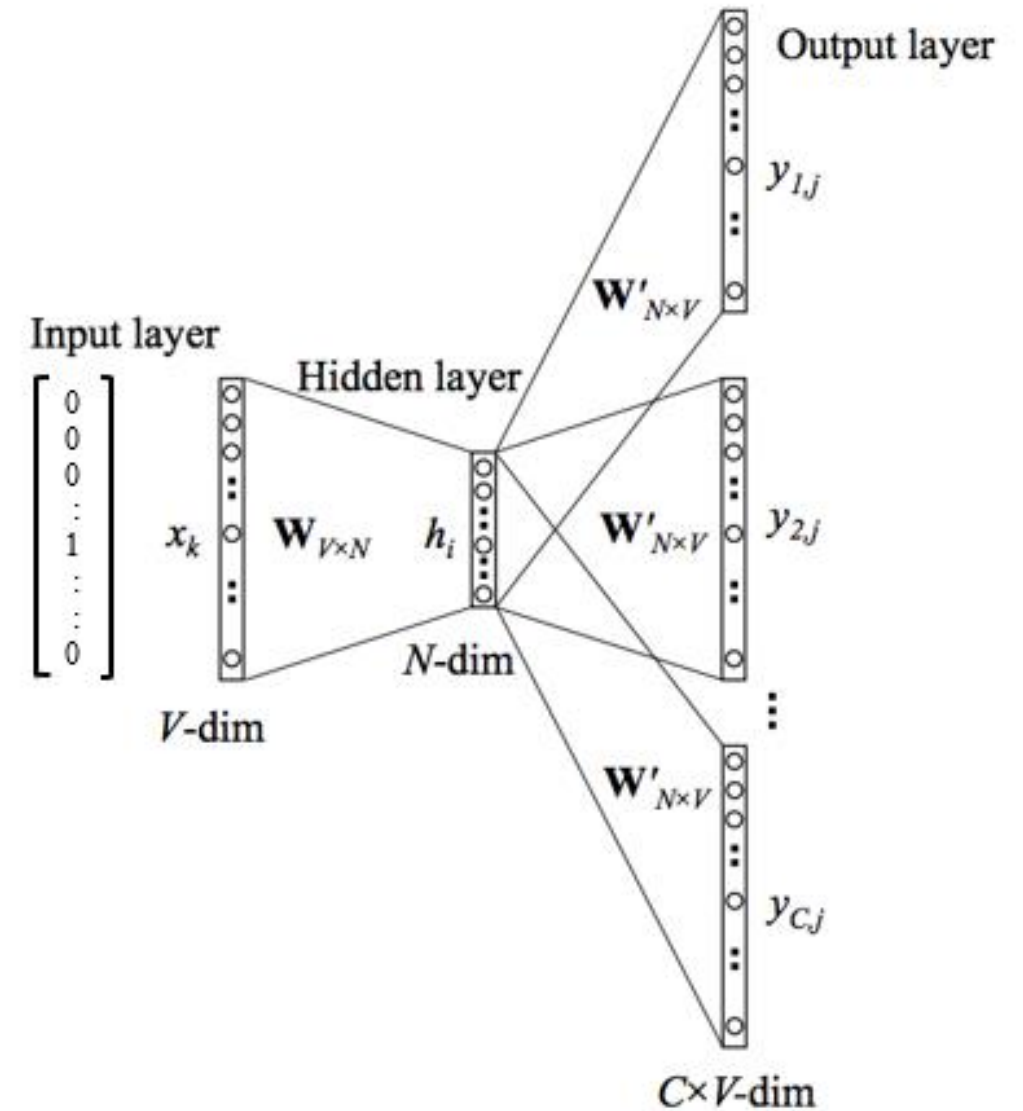
- Its size is  $V$ 
  - $V$ =number of words in the vocabulary
- The input is **1 middle/target vector**:
  - the target vector is a one-hot encoded vector (a vector having 1 in the position encoding the word, and 0 otherwise).






## Hidden Layer:

- Its size is  $N$ 
  - $N$ =number of neurons in the hidden layer  
=the dimension of embeddings  
(it is a hyper parameter)
- Fully-connected (dense) layer whose weights are the word embeddings.



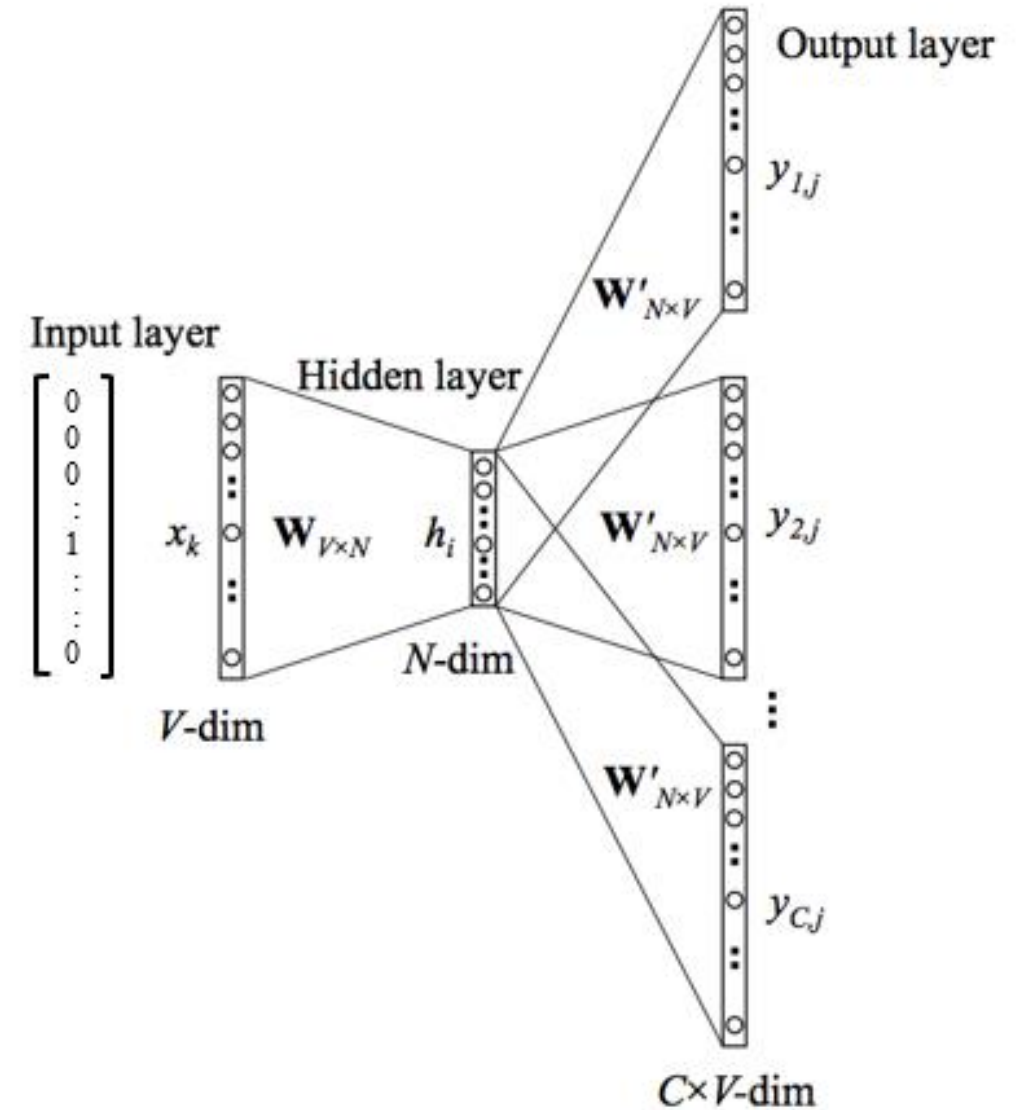
## Output Layer:

- Its size is CV
  - C=number of context words
  - V=number of words in the vocabulary
  - Where the gold outputs are C one-hot vectors each encoding one of the “context words”.
- It is a softmax layer which is designed so that the sum of the probabilities obtained in the output layer equals 1.
  - The network is going to tell us the probability for every word in our vocabulary of being the “context word” for the given input “target word”
- **C separate errors** are calculated and are added element-wise to obtain a final error vector which is back propagated to update the weights.




## Weights:

- There are two sets of weights:
  - one is between the input and the hidden layer:  
**Input-Hidden** layer matrix size  $= [V \times N]$
  - second is between hidden and output layer:  
**Hidden-Output** layer matrix consisting of  $C$   $W' [N \times V]$  matrices





- There is no activation function between any layers (linear activation):
  - Linear network except for the softmax at the output layer.
- The weights between the input layer and the hidden layer are taken as the word vectors after training.



- Let the available corpus be this one sentence “*word vectors are widely used*”.
- The unique words in the vocabulary are 5 words ( $V = 5$ ) as follows:  
[“word”, “vectors”, “are”, “widely”, “used”].
- Training data for  $C=2$  (context words)  $\rightarrow$  window size=1

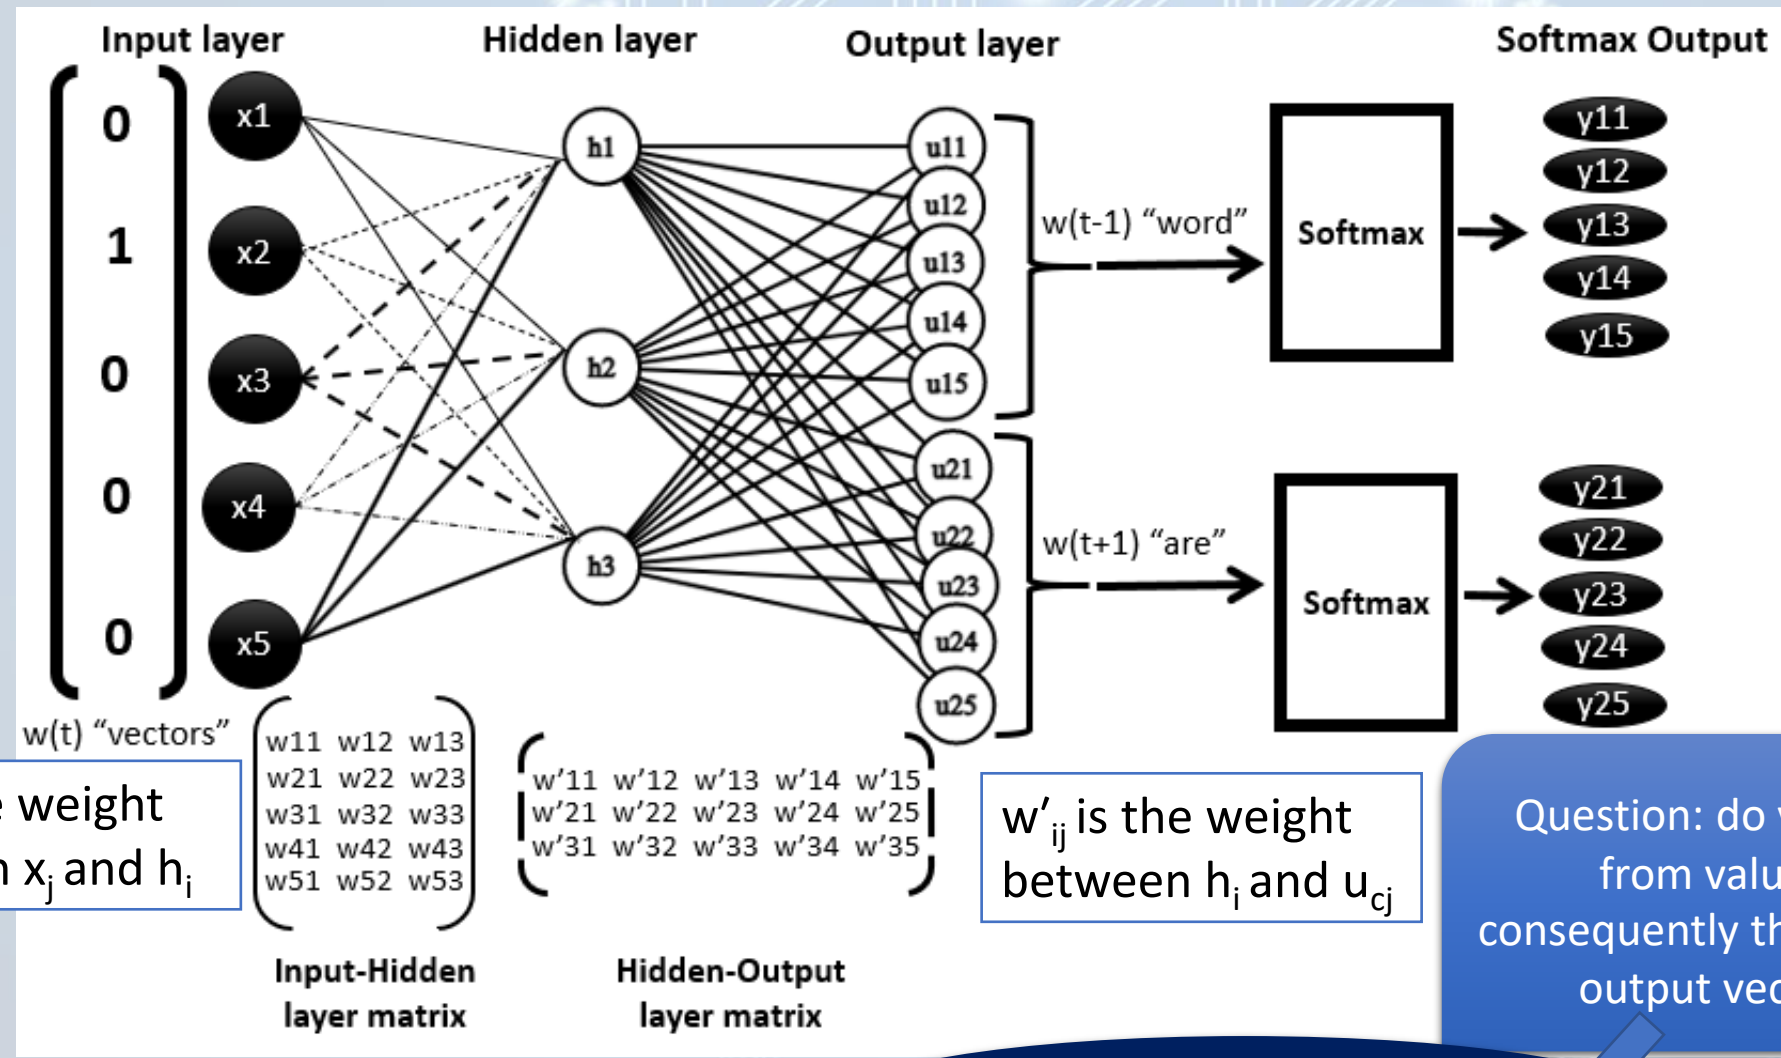
training sample	context words	target word
1	(word,are)	vectors
2	(vectors,widely)	are
3	(are,used)	widely
4	(widely)	used

- Encoded training data:

training sample	context words	target word
1	([1,0,0,0,0],[0,0,1,0,0])	[0,1,0,0,0]
2	([0,1,0,0,0],[0,0,0,1,0])	[0,0,1,0,0]
3	([0,0,1,0,0],[0,0,0,0,1])	[0,0,0,1,0]
4	([0,0,0,1,0])	[0,0,0,0,1]

- Let  $N = 3$ , consider the first training sample:

training sample	context words	target word
1	$[(1,0,0,0,0), [0,0,1,0,0)]$	$[0,1,0,0,0]$



$w_{ji}$  is the weight between  $x_j$  and  $h_i$

$w'_{ij}$  is the weight between  $h_i$  and  $u_{cj}$

Question: do values  $u_{11} \rightarrow u_{15}$  differ from values  $u_{21} \rightarrow u_{25}$  (and consequently the  $y$ ) i.e. do we have one output vector or  $C$  vectors???

Only one since  $W'$  matrix is the same



- The training process is similar to CBOW except that the error calculation differs as there are  $C$  gold outputs. Therefore, the final error is computed as the summation of the error for each context word as follows:

$$E = \sum_{c=1}^C E_c$$

- Finally, the vectors of the words in the vocabulary are weights in the Input-Hidden layer matrix such that each row in the matrix is the word vector representation of the equivalent word.

The word in the vocabulary	The equivalent word vector
word	$[w_{11} \ w_{12} \ w_{13}]$
vectors	$[w_{21} \ w_{22} \ w_{23}]$
are	$[w_{31} \ w_{32} \ w_{33}]$
widely	$[w_{41} \ w_{42} \ w_{43}]$
used	$[w_{51} \ w_{52} \ w_{53}]$

# SG Calculations

- For input  $\mathbf{x}$  with  $x_k=1$  and  $x_{k'}=0$  for all  $k' \neq k$ , the outputs of the hidden layer will be row  $k$  in  $\mathbf{W}$ :  $\mathbf{h} = \mathbf{x}^T \mathbf{W} = \mathbf{W}_{(k, \cdot)} := \mathbf{v}_{w_I}$
- The input to the  $j$ th node of the  $c$ th output word is:  $u_{c,j} = \mathbf{v}_{w_j}'^T \mathbf{h}$  →  $j$ th column in  $\mathbf{W}'$  matrix
- Since the output layers for each output word share the same weights so  $u_{c,j} = u_j$
- Computing the  $j$ th node of the  $c$ th output word via softmax function:

$$p(w_{c,j} = w_{O,c} | w_I) = y_{c,j} = \frac{\exp(u_{c,j})}{\sum_{j'=1}^V \exp(u_{j'})}$$

→ This is the probability that the output of the  $j$ th node of the  $c$ th output is equal to the actual value of the  $j$ th index of the  $c$ th output vector (one-hot encoded)

- This way the forward pass is done so next the weights  $\mathbf{W}$  and  $\mathbf{W}'$  need to be learned with backpropagation and stochastic gradient descent.

# SG Calculations

- First, a loss function is defined as:

**Negative log likelihood**

→ objective is to minimize it

$$-\sum_{c=1}^C \log$$

$$E = -\log p(w_{O,1}, w_{O,2}, \dots, w_{O,C} | w_I)$$

$$= -\log \prod_{c=1}^C \frac{\exp(u_{c,j_c^*})}{\sum_{j'=1}^V \exp(u_{j'})}$$

$$= -\sum_{c=1}^C u_{j_c^*} + C \cdot \log \sum_{j'=1}^V \exp(u_{j'})$$

The probability of the output words (context words) given the input (target/middle) word

$j_c^*$  is the index of the  $c$ th output

- Take the derivative w.r.t.  $u_{c,j}$ :

$$\frac{\partial E}{\partial u_{c,j}} = y_{c,j} - t_{c,j}$$

$t_{c,j}=1$  if the  $j$ th node of the  $c$ th true output word is equal to 1 and 0 otherwise.

How is this computed??

- Compute the derivative w.r.t.  $W'$ :

$$\begin{aligned} \frac{\partial E}{\partial w'_{ij}} &= \sum_{c=1}^C \frac{\partial E}{\partial u_{c,j}} \cdot \frac{\partial u_{c,j}}{\partial w'_{ij}} \\ &= \sum_{c=1}^C (y_{c,j} - t_{c,j}) \cdot h_i \end{aligned}$$

Using the chain rule

- The gradient descent update equation for  $W'$ :

$$w'_{ij}^{(new)} = w'_{ij}^{(old)} - \eta \cdot \sum_{c=1}^C (y_{c,j} - t_{c,j}) \cdot h_i$$



# SG Calculations

- Computing for  $W$ , begin with taking the derivative of the error w.r.t.  $h_i$ :

$$\frac{\partial E}{\partial h_i} = \sum_{j=1}^V \frac{\partial E}{\partial u_j} \cdot \frac{\partial u_j}{\partial h_i} \quad \text{Using the chain rule}$$

$$= \sum_{j=1}^V \sum_{c=1}^C (y_{c,j} - t_{c,j}) \cdot w'_{ij}$$

- Compute the derivative w.r.t.  $W$ :

$$\frac{\partial E}{\partial w_{ki}} = \frac{\partial E}{\partial h_i} \cdot \frac{\partial h_i}{\partial w_{ki}} \quad \text{Using the chain rule}$$

$$= \sum_{j=1}^V \sum_{c=1}^C (y_{c,j} - t_{c,j}) \cdot w'_{ij} \cdot x_k$$

- The gradient descent update equation for  $W$ :

$$w_{ji}^{(new)} = w_{ji}^{(old)} - \eta \cdot \sum_{j=1}^V \sum_{c=1}^C (y_{c,j} - t_{c,j}) \cdot w'_{ij} \cdot x_j$$

$$[V \times N] = [V \times N] - [V \times 1] ([N \times V] [V \times 1])^T$$

$$W_{input}^{(new)} = W_{input}^{(old)} - \eta \cdot x \cdot (W_{output} \sum_{c=1}^C e_c)^T$$

- Matrix forms:

$$[N \times V] = [N \times V] - [N \times 1] [1 \times V]$$

$$W_{output}^{(new)} = W_{output}^{(old)} - \eta \cdot h \cdot \sum_{c=1}^C e_c$$

Each gradient descent update requires a **summation over the entire vocabulary** → **Computationally expensive**. So techniques such as hierarchical softmax and negative sampling are used to make this computation more efficient.

# Other Kinds of Static Embeddings



FastText

GloVe

Sphere

# FastText

- An extension of word2vec developed by *Piotr Bojanowski et al. at Facebook*.
- Addresses a problem with word2vec “**unknown words**”—words that appear in a test corpus but were unseen in the training corpus.
- A related problem is **word sparsity**, such as in languages with rich morphology, where some of the many forms for each noun and verb may only occur rarely.
- FastText deals with these problems by using **subword models**, representing each word as **itself plus a bag of constituent n-grams**, with special boundary symbols < and > added to each word.
- For example, with  $n = 3$  the word *where* would be represented by the sequence <where> plus the character n-grams: <wh, whe, her, ere, re>
  - Note that the sequence <her>, corresponding to the word *her* is different from the tri-gram *her* from the word *where*.
- Now, the vocabulary is the union of the subwords of all words (subwords refer to the words themselves + n-grams). Then a skip-gram/cbow embedding is learned for each subword.
- The vector of a word is the sum of the whole word vector in addition to the sum of its subwords' vectors.
- Unknown words can then be presented only by the sum of the constituent n-grams.



# GloVe

- Short for Global Vectors, because the model is based on capturing global corpus statistics developed by *Jeffrey Pennington et al. at Stanford University*.
- Unlike word2vec that captures local statistics of a corpus, GloVe captures both local and global statistics of a corpus.
  - In word2vec, the vector learnt for a certain word is only affected by the surrounding context words.
  - In GloVe, the global statistics of the corpus are considered as it operates on **co-occurrence matrices** that are built across the entire corpus.
- GloVe is based on **ratios of probabilities** from the word-word co-occurrence matrix.
- GloVe is a global log-bilinear regression model.

# Sphere

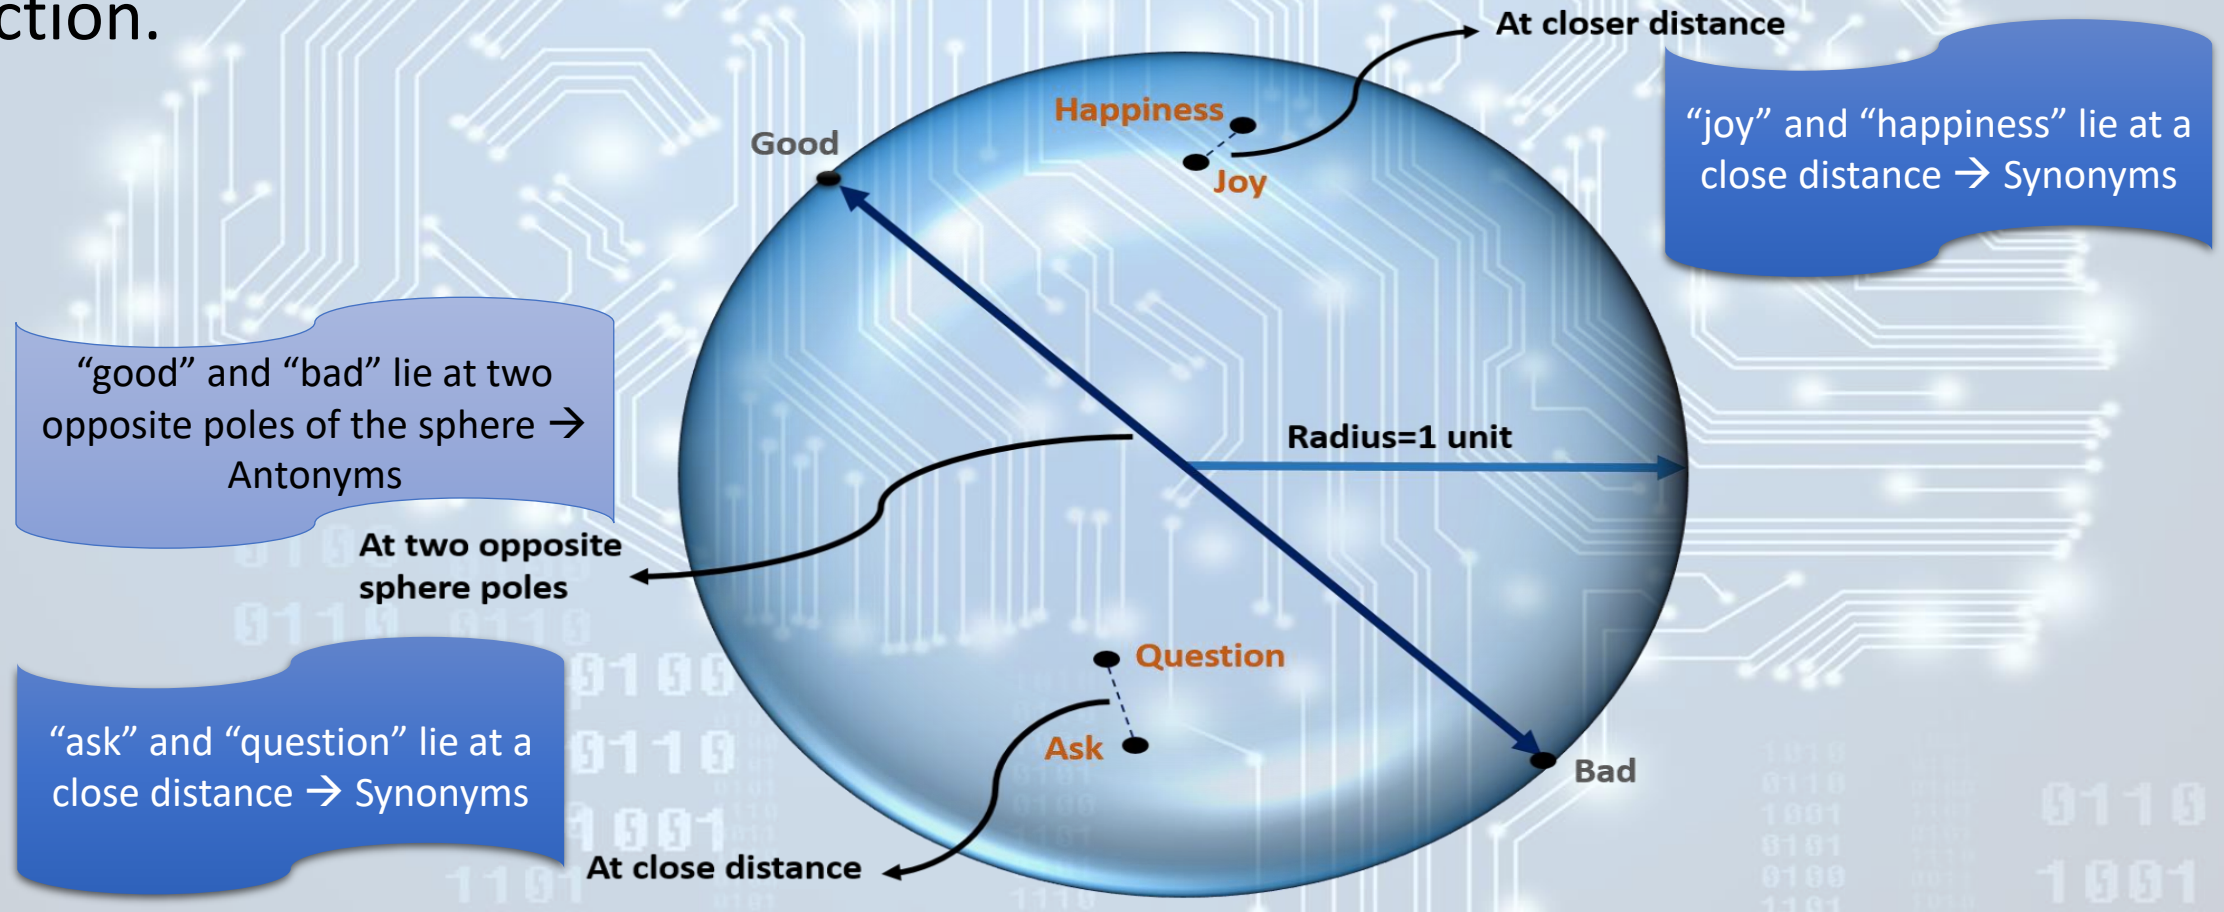
- The approaches such as word2vec, fastText and GloVe are **distributional** approaches:
  - Antonyms (*good, bad*) can lie in close proximity (**cosine similarity close to 1**)  
→ as antonyms can often occur in similar contexts.
  - Synonyms (*happiness, joy*) lie in close proximity (**cosine similarity close to 1**).
  - Unrelated words (*car, apple*) lie far away (**cosine similarity close to 0**).
  - Other antonyms can lie far away (**cosine similarity close to 0**).

There is **no clear distinction** between the different semantic relations among words such as synonyms, antonyms and unrelated words.



# Sphere

- Sphere approach developed by *Sandra Rizkallah et al. at Cairo University* solves this problem by embedding the vectors of opposite or antonyms words at opposite poles of the sphere (**cosine similarity close to -1**).
- Sphere is a semi-supervised relaxation algorithm that optimizes a devised error function.





# Evaluating Vector Models

- The most important evaluation metric for vector models is **extrinsic** evaluation on tasks.
- It is also useful to have **intrinsic** evaluations:
  1. **Similarity**: computing the correlation between an algorithm's word similarity scores and word similarity ratings assigned by humans.
  2. **Semantic Textual Similarity**: evaluates the performance of sentence-level similarity algorithms, consisting of a set of pairs of sentences, each pair with human-labeled similarity scores.
  3. **Analogy**: the system has to solve problems of the form "*a is to b as c is to \_\_\_\_*" e.g.: "*Cairo is to Egypt as Rome is to \_\_\_\_*"  $\rightarrow$  "*Italy*". Need to find *d* such that the associated word vectors *va*, *vb*, *vc*, *vd* are related to each other in the following relationship: " *$vb - va = vd - vc$* " where the *vd* is obtained to be the vector with highest cosine similarity to the vector (*vb*-*va*+*vc*).



# Thank You